

Docker Fundamentals for Linux Exercises

Please note that examples of expected output have in some places been truncated or abbreviated to better fit this document. These examples are rough guides only; if you see more information in your terminal than in the text, that is expected.

Windows Users: Please note that in all exercises we will use Unix style paths using forward slashes ('/') instead of backslashes ('\'). On Windows you can work directly with such paths by either using a **Bash** terminal or a **Powershell** terminal. Powershell can work with both Windows and Unix style paths.

Be aware that copy-pasting of commands or code snippets from this PDF may apply changes to some characters e.g. quotes, tabs, which may leads to errors. Please consider typing suggested commands and code snippets in case you encounter any issues.

Introduction to Strigo: Strigo is our lab platform, which your instructor will provide you with access credentials for. Log in right away and get familiar with the interface by following the Strigo primer: <https://bit.ly/2SikQL6>.

Contents

Exercises	5
1 Running & Inspecting Containers	5
1.1 Conclusion	8
2 Interactive Containers	8
2.1 Writing to Containers	8
2.2 Reconnecting to Containers	8
2.3 Using Container Listing Options	9
2.4 Conclusion	9
3 Detached Containers and Logging	10
3.1 Running a Container in the Background	10
3.2 Attaching to Container Output	10
3.3 Using Logging Options	11
3.4 Conclusion	11
4 Starting, Stopping, Inspecting and Deleting Containers	11
4.1 Starting and Restarting Containers	11
4.2 Inspecting a Container	12
4.3 Deleting Containers	12
4.4 Conclusion	13
5 Interactive Image Creation	13
5.1 Modifying a Container	13
5.2 Capturing Container State as an Image	14
5.3 Conclusion	14
6 Creating Images with Dockerfiles (1/2)	14
6.1 Writing and Building a Dockerfile	14
6.2 Using the Build Cache	15

6.3	Using the <code>history</code> Command	15
6.4	Conclusion	16
7	Creating Images with Dockerfiles (2/2)	16
7.1	Setting Default Commands	16
7.2	Combining Default Commands and Options	17
7.3	Conclusion	17
8	Multi-Stage Builds	17
8.1	Defining a multi-stage build	17
8.2	Building Intermediate Images	19
8.3	Optional: Building from Scratch	19
8.4	Optional: Enabling BuildKit	20
8.5	Conclusion	21
9	Managing Images	21
9.1	Making an Account on Docker's Hosted Registry	22
9.2	Tagging and Listing Images	22
9.3	Sharing Images on Docker Hub	22
9.4	Conclusion	23
10	Database Volumes	23
10.1	Launching Postgres	23
10.2	Writing to the Database	23
10.3	Conclusion	24
11	Introduction to Container Networking	24
11.1	Inspecting the Default Bridge	24
11.2	Connecting Containers to <code>docker0</code>	25
11.3	Defining Additional Bridge Networks	26
11.4	Conclusion	27
12	Container Port Mapping	27
12.1	Port Mapping at Runtime	27
12.2	Exposing Ports from the Dockerfile	27
12.3	Conclusion	28
13	Creating a Swarm	28
13.1	Starting Swarm	28
13.2	Adding Workers to the Swarm	29
13.3	Promoting Workers to Managers	29
13.4	Conclusion	29
14	Starting a Service	30
14.1	Creating an Overlay Network and Service	30
14.2	Inspecting Service Logs	31
14.3	Cleanup	31
14.4	Conclusion	31
15	Node Failure Recovery	31
15.1	Setting up a Service	32
15.2	Simulating Node Failure	32
15.3	Force Rebalancing	32
15.4	Cleanup	32
15.5	Conclusion	33
16	Swarm Scheduling	33
16.1	Restricting Resource Consumption	33

16.2	Configuring Global Scheduling	35
16.3	Scheduling via Node Constraints	36
16.4	Scheduling Topology-Aware Services	36
16.5	Conclusion	37
17	Provisioning Swarm Configuration	37
17.1	Creating a Stack	37
17.2	Defining and Using <code>.env</code> Files	38
17.3	Defining and Using Docker Configs	39
17.4	Defining and Using Docker Secrets	40
17.5	Conclusion	42
18	Routing to Services	42
18.1	Routing Cluster-Internal Traffic	42
18.2	Routing Cluster-External Traffic	44
18.3	Conclusion	45
19	Updating Applications	45
19.1	Deploying Dockercoins	45
19.2	Scaling Up an Application	46
19.3	Creating Rolling Updates	47
19.4	Parallelizing Updates	47
19.5	Auto-Rollback Failed Updates	48
19.6	Optional Challenge: Improving Dockercoins	49
19.7	Conclusion	49
20	Installing Kubernetes	49
20.1	Initializing Kubernetes	49
20.2	Conclusion	50
21	Kubernetes Orchestration	51
21.1	Creating Pods	51
21.2	Creating ReplicaSets	52
21.3	Creating Deployments	53
21.4	Conclusion	55
22	Provisioning Kube Configuration	55
22.1	Provisioning ConfigMaps	55
22.2	Provisioning Secrets	58
22.3	Conclusion	60
23	Kubernetes Networking	60
23.1	Routing Traffic with Calico	60
23.2	Routing and Load Balancing with Services	63
23.3	Optional: Deploying DockerCoins onto the Kubernetes Cluster	64
23.4	Conclusion	67
24	Cleaning up Docker Resources	67
24.1	Conclusion	68
25	Inspection Commands	68
25.1	Inspecting System Information	68
25.2	Monitoring System Events	68
25.3	Conclusion	69
26	Plugins	69
26.1	Installing a Plugin	69
26.2	Enabling and Disabling a Plugin	70

26.3 Inspecting a Plugin	70
26.4 Using the Plugin	70
26.5 Removing a Plugin	71
26.6 Conclusion	71
27 Starting a Compose App	71
27.1 Inspecting a Compose App	71
27.2 Starting the App	72
27.3 Viewing Logs	72
27.4 Conclusion	72
28 Scaling a Compose App	73
28.1 Scaling a Service	73
28.2 Investigating Bottlenecks	73
28.3 Conclusion	74
29 Appendix: A Brief Introduction to Bash	74
29.1 Managing Files and Folders	74
29.2 Some Common Tools	75
29.2.1 Inspecting Processes	75
29.2.2 The Superuser	76
29.2.3 Pinging of an address	76
29.2.4 Making HTTP Requests	76
29.2.5 Working with Commands	77
29.3 Conclusion	78
Instructor Demos	78
1 Instructor Demo: Process Isolation	78
1.1 Exploring the PID Kernel Namespace	78
1.2 Imposing Resource Limitations With Cgroups	79
1.3 Conclusion	80
2 Instructor Demo: Creating Images	81
2.1 Understanding Image Build Output	81
2.2 Managing Image Layers	82
2.3 Conclusion	83
3 Instructor Demo: Basic Volume Usage	83
3.1 Using Named Volumes	84
3.2 Mounting Host Paths	85
3.3 Conclusion	85
4 Instructor Demo: Single Host Networks	86
4.1 Following Default Docker Networking	86
4.2 Establishing Custom Docker Networks	88
4.3 Forwarding a Host Port to a Container	91
4.4 Conclusion	91
5 Instructor Demo: Self-Healing Swarm	91
5.1 Setting Up a Swarm	91
5.2 Scheduling Workload	92
5.3 Maintaining Desired State	93
5.4 Conclusion	94
6 Instructor Demo: Kubernetes Basics	94
6.1 Initializing Kubernetes	94
6.2 Exploring Kubernetes Scheduling	95

6.3	Exploring Containers in a Pod	96
6.4	Conclusion	97
7	Instructor Demo: Docker Compose	97
7.1	Exploring the Compose File	97
7.2	Communicating Between Containers	98
7.3	Conclusion	99

Exercises

1 Running & Inspecting Containers

By the end of this exercise, you should be able to:

- Start a container
 - List containers in multiple ways
 - Query the docker command line help
 - Remove containers
1. Create and start a new CentOS 7 container running ping to 8.8.8.8. Docker will download the CentOS 7 image since you do not have it available locally.

```
[centos@node-0 ~]$ docker container run centos:7 ping 8.8.8.8

Unable to find image 'centos:7' locally
7: Pulling from library/centos
a02a4930cb5d: Pull complete
Digest: sha256:184e5f35598e333bfa7de10d8fb1cebb5ee4df5bc0f970bf2b1e7c7345136426
Status: Downloaded newer image for centos:7
PING 8.8.8.8 (8.8.8.8) 56(84) bytes of data.
64 bytes from 8.8.8.8: icmp_seq=1 ttl=108 time=7.07 ms
64 bytes from 8.8.8.8: icmp_seq=2 ttl=108 time=7.11 ms
64 bytes from 8.8.8.8: icmp_seq=3 ttl=108 time=7.03 ms
64 bytes from 8.8.8.8: icmp_seq=4 ttl=108 time=7.09 ms
64 bytes from 8.8.8.8: icmp_seq=5 ttl=108 time=7.01 ms
64 bytes from 8.8.8.8: icmp_seq=6 ttl=108 time=7.00 ms
^C
--- 8.8.8.8 ping statistics ---
6 packets transmitted, 6 received, 0% packet loss, time 5006ms
rtt min/avg/max/mdev = 7.008/7.056/7.110/0.039 ms
....
```

Press CTRL+C after a few pings. This stops and exits the container.

2. This first container sent its STDOUT to your terminal. Create a second container, this time in *detached mode*:

```
[centos@node-0 ~]$ docker container run --detach centos:7 ping 8.8.4.4

8aef3d0d411c7b02532292ec3267a54f9258eaafb71d3d73a8ad41e702bd35a2
```

Instead of seeing the executed command (ping 8.8.4.4), Docker engine displays a long hexadecimal number, which is the full *container ID* of your new container. The container is running detached, which means the container is running as a background process, rather than printing its STDOUT to your terminal.

3. List the running Docker containers using the `docker container ls` command. You will see only one container running.

```
[centos@node-0 ~]$ docker container ls
```

CONTAINER ID	IMAGE	COMMAND	CREATED	STATUS	PORTS	NAMES
8aef3d0d411c	centos:7	"ping 8.8.4.4"	3 minutes ago	Up 3 minutes		zen_jang

4. Now you know that the `docker container ls` command only shows running containers. You can show all containers that exist (running or stopped) by using `docker container ls --all`. Your container ID and name will vary. Note that you will see two containers: a stopped container and a running container.

```
[centos@node-0 ~]$ docker container ls --all
```

CONTAINER ID	IMAGE	COMMAND	CREATED	STATUS	NAMES
8aef3d0d411c	centos:7	"ping 8.8.4.4"	2 minutes ago	Up 2 minutes	zen_jang
00f763b9308d	centos:7	"ping 8.8.8.8"	4 minutes ago	Exited (0)...	inspiring_cheb

Where did those names come from? All containers have names, which in most Docker CLI commands can be substituted for the container ID as we'll see in later exercises. By default, containers get a randomly generated name of the form <adjective>_<scientist / technologist>, but you can choose a name explicitly with the `--name` flag in `docker container run`.

5. Start up another detached container, this time giving it a name "opendnsping".

```
[centos@node-0 ~]$ docker container run --detach --name opendnsping \
centos:7 ping 208.67.222.222
```

```
3bdc61a95e76fdfe2597ef18aa00321a53dc9c36b2db97f738f8a623ecad
```

6. List all your containers again. You can see all of the containers, including your new one with your customized name.

```
[centos@node-0 ~]$ docker container ls --all
```

CONTAINER ID	IMAGE	COMMAND	CREATED	STATUS	NAMES
3bdc61a95e76	centos:7	"ping 208.67.222.222"	23 seconds ago	Up 22 seconds	opendnsping
8aef3d0d411c	centos:7	"ping 8.8.4.4"	4 minutes ago	Up 14 minutes	zen_jang
00f763b9308d	centos:7	"ping 8.8.8.8"	9 minutes ago	Exited (0)...	inspirin...

7. Next, remove the exited container. To do this, use `docker container rm <container ID>`. In the example above, the Docker container ID is 00f763b9308d. You only need as many characters as to uniquely identify this container from all the others.

```
[centos@node-0 ~]$ docker container rm <container ID>
```

```
00f763b9308d
```

8. Now try to remove one of the other Docker containers using the same command. It does not work. Why?

```
[centos@node-0 ~]$ docker container rm <container ID>
```

```
Error response from daemon: You cannot remove a running container
3bdc61a95e76fdfe2597ef18aa00321a53dc9c36b2db97f738f8a623ecad.
Stop the container before attempting removal or force remove
```

9. You can see that running containers are not removed. You'll have to look for an option to remove a running container. In order to find out the option you need to do a force remove, check the command line help. To do this with the `docker container rm` command, use the `--help` option:

```
[centos@node-0 ~]$ docker container rm --help
```

```
Usage: docker container rm [OPTIONS] CONTAINER [CONTAINER...]
```

Remove one or more containers

Options:

-f, --force	Force the removal of a running container (uses SIGKILL)
-l, --link	Remove the specified link
-v, --volumes	Remove the volumes associated with the container

Help works with all Docker commands Not only can you use `--help` with `docker container rm`, but it works on all levels of docker commands. For example, `docker --help` provides you with all available docker commands, and `docker container --help` provides you with all available container commands.

10. Now, run a force remove on the running container you tried to remove in the two previous steps. This time it works.

```
[centos@node-0 ~]$ docker container rm --force <container ID>
8aef3d0d411c
```

11. Start another detached container pinging 8.8.8.8, with the name `pinggoogledns`.

```
[centos@node-0 ~]$ docker container run --detach --name pinggoogledns \
    centos:7 ping 8.8.8.8
38e121e629611daa0726a21d634bc5189400377d82882cc6fd8a3870dc9943a0
```

12. Now that you've finished your testing, you need to remove your containers. In order to remove all of them at once, you want to get only the container IDs. Look at `docker container ls --help` to get the information you need:

```
[centos@node-0 ~]$ docker container ls --help

Usage:  docker container ls [OPTIONS]

List containers

Aliases:
  ls, ps, list

Options:
  -a, --all                Show all containers (default shows just running)
  -f, --filter filter      Filter output based on conditions provided
  --format string          Pretty-print containers using a Go template
  -n, --last int           Show n last created containers (includes all states)
  -l, --latest             Show the latest created container (includes all states)
  --no-trunc              Don't truncate output
  -q, --quiet              Only display numeric IDs
  -s, --size              Display total file sizes
```

13. To get only the container IDs, use the `--quiet` option. If you want to use only the container IDs of all existing containers to perform an action on, you can use `--quiet` with the `--all` option.

```
[centos@node-0 ~]$ docker container ls --all --quiet
3bdc61a95e76
38e121e62961
```

14. Since we are done running pings on the public DNS servers, kill the containers. To do this, use the syntax `docker container rm --force <containerID>`. However, this only kills one container at a time. We want to kill all the containers, no matter what state the containers are in. To get this information, you will

need to use the output from `docker container ls --quiet --all`. To capture this output within the command, use `$(...)` to nest the listing command inside the `docker container rm` command.

```
[centos@node-0 ~]$ docker container rm --force \
    $(docker container ls --quiet --all)

3bdc61a95e76
38e121e62961
```

1.1 Conclusion

This exercise taught you how to start, list, and kill containers. In this exercise you ran your first containers using `docker container run`, and how they are running commands inside the containers. You also learned to how to list your containers, and how to kill the containers using the command `docker container rm`. If you run into trouble, you've learned that the `--help` option can provide you with information that could help get you answers.

2 Interactive Containers

By the end of this exercise, you should be able to:

- Launch an interactive shell in a new or existing container
- Run a child process inside a running container
- List containers using more options and filters

2.1 Writing to Containers

1. Create a container using the `centos:7` image. Connect to its bash shell in interactive mode using the `-i` flag and request a TTY connection using the `-t` flag:

```
[centos@node-0 ~]$ docker container run -it centos:7 bash
```

2. Explore your container's filesystem with `ls`, and then create a new file. Use `'ls'` again to confirm you have successfully created your file. Use the `'-l'` option with `'ls'` to list the files and directories in long list format.

```
[root@2b8de2ffdf85 /]# ls -l
[root@2b8de2ffdf85 /]# echo 'Hello there...' > test.txt
[root@2b8de2ffdf85 /]# ls -l
```

3. Exit the connection to the container:

```
[root@2b8de2ffdf85 /]# exit
```

4. Run the same command as before to start a container using the `'centos:7'` image:

```
[centos@node-0 ~]$ docker container run -it centos:7 bash
```

5. Use `'ls'` to explore your container. You will see that your previously created `test.txt` is nowhere to be found in your new container. Exit this container in the same way you did above.

2.2 Reconnecting to Containers

1. We wish to recover `'test.txt'` written to our container in the first example, but starting a new container didn't get us there. We need to restart and reconnect to our original container. List all your stopped containers:

```
[centos@node-0 ~]$ docker container ls -a
```

CONTAINER ID	IMAGE	COMMAND	CREATED	STATUS
cc19f7e9aa91	centos:7	"bash"	About a minute ago	Exited (0) About a minute ago


```
2b8de2ffdf85  centos:7  "bash"  2 minutes ago  Exited (0) About a minute ago
...
```

2. We can restart a container via the container ID listed in the first column. Use the container ID for the first centos:7 container you created with bash as its command (see the CREATED column above to make sure you're choosing the *first* bash container you ran):

```
[centos@node-0 ~]$ docker container start <container ID>
[centos@node-0 ~]$ docker container ls
```

```
CONTAINER ID  IMAGE      COMMAND      CREATED          STATUS          ...
2b8de2ffdf85  centos:7   "bash"       5 minutes ago    Up 21 seconds   ...
```

Your container status has changed from Exited to Up, via `docker container start`.

3. Run `ps -ef` inside the container you just restarted using Docker's `exec` command (`exec` runs the specified process as a child of the PID 1 process inside the container):

```
[centos@node-0 ~]$ docker container exec <container ID> ps -ef
```

What process is PID 1 inside the container? Find the PID of that process on the host machine by using:

```
[centos@node-0 ~]$ docker container top <container ID>
```

4. Launch a bash shell in your running container with `docker container exec`:

```
[centos@node-0 ~]$ docker container exec -it <container ID> bash
```

5. List the contents of the container's filesystem with `ls -l`; your `test.txt` should be where you left it. Exit the container again by typing `exit`.

2.3 Using Container Listing Options

1. In the last step, we saw how to get the short container ID of all our containers using `docker container ls -a`. Try adding the `--no-trunc` flag to see the entire container ID:

```
[centos@node-0 ~]$ docker container ls -a --no-trunc
```

This long ID is the same as the string that is returned after starting a container with `docker container run`.

2. List only the container ID using the `-q` flag:

```
[centos@node-0 ~]$ docker container ls -a -q
```

3. List the last container to have been created using the `-l` flag:

```
[centos@node-0 ~]$ docker container ls -l
```

4. Finally, you can also filter results with the `--filter` flag; for example, try filtering by exit code:

```
[centos@node-0 ~]$ docker container ls -a --filter "exited=0"
```

The output of this command will list the containers that have exited successfully.

5. Clean up with:

```
[centos@node-0 ~]$ docker container rm -f $(docker container ls -aq)
```

2.4 Conclusion

In this demo, you saw that files added to a container's filesystem do not get added to all containers created from the same image. Changes to a container's filesystem are local to itself, and exist only in that particular container. You also learned how to restart a stopped Docker container using `docker container start`, how to run a command

in a running container using `docker container exec`, and also saw some more options for listing containers via `docker container ls`.

3 Detached Containers and Logging

By the end of this exercise, you should be able to:

- Run a container detached from the terminal
- Fetch the logs of a container
- Attach a terminal to the STDOUT of a running container

3.1 Running a Container in the Background

1. First try running a container as usual; the STDOUT and STDERR streams from whatever is PID 1 inside the container are directed to the terminal:

```
[centos@node-0 ~]$ docker container run centos:7 ping 127.0.0.1 -c 2

PING 127.0.0.1 (127.0.0.1) 56(84) bytes of data.
64 bytes from 127.0.0.1: icmp_seq=1 ttl=64 time=0.021 ms
64 bytes from 127.0.0.1: icmp_seq=2 ttl=64 time=0.029 ms

--- 127.0.0.1 ping statistics ---
2 packets transmitted, 2 received, 0% packet loss, time 1019ms
rtt min/avg/max/mdev = 0.021/0.025/0.029/0.004 ms
```

2. The same process can be run in the background with the `-d` flag:

```
[centos@node-0 ~]$ docker container run -d centos:7 ping 127.0.0.1

d5ef517cc113f36738005295066b271ae604e9552ce4070caffbacdc3893ae04
```

This time, we only see the container's ID; its STDOUT isn't being sent to the terminal.

3. Use this second container's ID to inspect the logs it generated:

```
[centos@node-0 ~]$ docker container logs <container ID>
```

These logs correspond to STDOUT and STDERR from the container's PID 1. Also note when using container IDs: you don't need to specify the entire ID. Just enough characters from the start of the ID to uniquely identify it, often just 2 or 3, is sufficient.

3.2 Attaching to Container Output

1. We can attach a terminal to a container's PID 1 output with the `attach` command; try it with the last container you made in the previous step:

```
[centos@node-0 ~]$ docker container attach <container ID>
```

2. We can leave attached mode by then pressing `CTRL+C`. After doing so, list your running containers; you should see that the container you attached to has been killed, since the `CTRL+C` issued killed PID 1 in the container, and therefore the container itself.
3. Try running the same thing in detached interactive mode:

```
[centos@node-0 ~]$ docker container run -d -it centos:7 ping 127.0.0.1
```

4. Attach to this container like you did the first one, but this time detach with `CTRL+P CTRL+Q` (sequential, not simultaneous), and list your running containers. In this case, the container should still be happily running in the background after detaching from it.

3.3 Using Logging Options

1. We saw previously how to read the entire log of a container's PID 1; we can also use a couple of flags to control what logs are displayed. `--tail n` limits the display to the last `n` lines; try it with the container that should be running from the last step:

```
[centos@node-0 ~]$ docker container logs --tail 5 <container ID>
```

You should see the last 5 pings from this container.

2. We can also follow the logs as they are generated with `-f`:

```
[centos@node-0 ~]$ docker container logs -f <container ID>
```

The container's logs get piped in real time to the terminal (CTRL+C to break out of following mode - note this doesn't kill the process like when we attached to it, since now we're tailing the logs, not attaching to the process).

3. Finally, try combining the tail and follow flags to begin following the logs from 10 lines back in history.

3.4 Conclusion

In this exercise, we saw our first detached containers. Almost all containers you ever run will be running in detached mode; you can use `container attach` to interact with their PID 1 processes, as well as `container logs` to fetch their logs. Note that both `attach` and `logs` interact with the PID 1 process only - if you launch child processes inside a container, it's up to you to manage their STDOUT and STDERR streams. Also, be careful when killing processes after attaching to a container; as we saw, it's easy to attach to a container and then kill it, by issuing a CTRL+C to the PID 1 process you've attached to.

4 Starting, Stopping, Inspecting and Deleting Containers

By the end of this exercise, you should be able to:

- Restart containers which have exited
- Distinguish between stopping and killing a container
- Fetch container metadata using `docker container inspect`
- Delete containers

4.1 Starting and Restarting Containers

1. Start by running a container in the background, and check that it's really running:

```
[centos@node-0 ~]$ docker container run -d centos:7 ping 8.8.8.8  
[centos@node-0 ~]$ docker container ls
```

2. Stop the container using `docker container stop`, and check that the container is indeed stopped:

```
[centos@node-0 ~]$ docker container stop <container ID>  
[centos@node-0 ~]$ docker container ls -a
```

Note that the `stop` command takes a few seconds to complete. `docker container stop` first sends a SIGTERM to the PID 1 process inside a container, asking it to shut down nicely; it then waits 10 seconds before sending a SIGKILL to kill it off, ready or not. The exit code you see (137 in this case) is the exit code returned by the PID 1 process (ping) upon being killed by one of these signals.

3. Start the container again with `docker container start`, and attach to it at the same time with the `-a` flag:

```
[centos@node-0 ~]$ docker container start -a <container ID>
```

As you saw previously, this brings the container from the `Exited` to the `Up` state; in this case, we're also attaching to the `PID 1` process.

4. Detach and stop the container with `CTRL+C`, then restart the container without attaching and follow the logs starting from 10 lines previous.
5. Finally, stop the container with `docker container kill`:

```
[centos@node-0 ~]$ docker container kill <container ID>
```

Unlike `docker container stop`, `container kill` just sends the `SIGKILL` right away - no grace period.

4.2 Inspecting a Container

1. Start your ping container again, then inspect the container details using `docker container inspect`:

```
[centos@node-0 ~]$ docker container start <container ID>
[centos@node-0 ~]$ docker container inspect <container ID>
```

You get a JSON object describing the container's config, metadata and state.

2. Find the container's IP and long ID in the JSON output of `inspect`. If you know the key name of the property you're looking for, try piping to `grep`:

```
[centos@node-0 ~]$ docker container inspect <container ID> | grep IPAddress
```

The output should look similar to this:

```
"SecondaryIPAddresses": null,
"IPAddress": "<Your IP Address>"
```

3. Now try grepping for `Cmd`, the `PID 1` command being run by this container. `grep`'s simple text search doesn't always return helpful results:

```
[centos@node-0 ~]$ docker container inspect <container ID> | grep Cmd

"Cmd": [
```

4. A more powerful way to filter this JSON is with the `--format` flag. Syntax follows Go's text/template package: <http://golang.org/pkg/text/template/>. For example, to find the `Cmd` value we tried to `grep` for above, instead try:

```
[centos@node-0 ~]$ docker container inspect --format='{{.Config.Cmd}}' <container ID>

[ping 8.8.8.8]
```

This time, we get a the value of the `Config.Cmd` key from the `inspect` JSON.

5. Keys nested in the JSON returned by `docker container inspect` can be chained together in this fashion. Try modifying this example to return the IP address you grepped for previously.
6. Finally, we can extract all the key/value pairs for a given object using the `json` function:

```
[centos@node-0 ~]$ docker container inspect --format='{{json .Config}}' <container ID>
```

Try adding `| jq` to this command to get the same output a little bit easier to read.

4.3 Deleting Containers

1. Start three containers in background mode, then stop the first one.
2. List only exited containers using the `--filter` flag we learned earlier, and the option `status=exited`.
3. Delete the container you stopped above with `docker container rm`, and do the same listing operation as above to confirm that it has been removed:

```
[centos@node-0 ~]$ docker container rm <container ID>
[centos@node-0 ~]$ docker container ls ...
```

4. Now do the same to one of the containers that's still running; notice `docker container rm` won't delete a container that's still running, unless we pass it the force flag `-f`. Delete the second container you started above:

```
[centos@node-0 ~]$ docker container rm -f <container ID>
```

5. Try using the `docker container ls` flags we learned previously to remove the last container that was run, or all stopped containers. Recall that you can pass the output of one shell command `cmd-A` into a variable of another command `cmd-B` with syntax like `cmd-B $(cmd-A)`.
6. When done, clean up any containers you may still have:

```
[centos@node-0 ~]$ docker container rm -f $(docker container ls -aq)
```

4.4 Conclusion

In this exercise, you explored the lifecycle of a container, particularly in terms of stopping and restarting containers. Keep in mind the behavior of `docker container stop`, which sends a `SIGTERM`, waits a grace period, and then sends a `SIGKILL` before forcing a container to stop; this two step process is designed to give your containers a chance to shut down 'nicely': dump their state to a log, finish a database transaction, or do whatever your application needs them to do in order to exit without causing additional problems. Make sure you bear this in mind when designing containerized software.

Also keep in mind the `docker container inspect` command we saw, for examining container metadata, state and config; this is often the first place to look when trying to troubleshoot a failed container.

5 Interactive Image Creation

By the end of this exercise, you should be able to:

- Capture a container's filesystem state as a new docker image
- Read and understand the output of `docker container diff`

5.1 Modifying a Container

1. Start a bash terminal in a CentOS container:

```
[centos@node-0 ~]$ docker container run -it centos:7 bash
```

2. Install a couple pieces of software in this container - there's nothing special about `wget`, any changes to the filesystem will do. Afterwards, exit the container:

```
[root@dfe86ed42be9 /]# yum install -y which wget
[root@dfe86ed42be9 /]# exit
```

3. Finally, try `docker container diff` to see what's changed about a container relative to its image; you'll need to get the container ID via `docker container ls -a` first:

```
[centos@node-0 ~]$ docker container ls -a
[centos@node-0 ~]$ docker container diff <container ID>

C /root
A /root/.bash_history
C /usr
C /usr/bin
```

```
A /usr/bin/gsoelim
...
```

Those Cs at the beginning of each line stand for files Changed, and A for Added; lines that start with D indicate Deletions.

5.2 Capturing Container State as an Image

1. Installing `which` and `wget` in the last step wrote information to the container's read/write layer; now let's save that read/write layer as a new read-only image layer in order to create a new image that reflects our additions, via the `docker container commit`:

```
[centos@node-0 ~]$ docker container commit <container ID> myapp:1.0
```

2. Check that you can see your new image by listing all your images:

```
[centos@node-0 ~]$ docker image ls
```

REPOSITORY	TAG	IMAGE ID	CREATED	SIZE
myapp	1.0	34f97e0b087b	8 seconds ago	300MB
centos	7	5182e96772bf	44 hours ago	200MB

3. Create a container running `bash` using your new image, and check that `which` and `wget` are installed:

```
[centos@node-0 ~]$ docker container run -it myapp:1.0 bash
[root@2ecb80c76853 /]# which wget
```

The `which` commands should show the path to the specified executable, indicating they have been installed in the image. Exit your container when done by typing `exit`.

5.3 Conclusion

In this exercise, you saw how to inspect the contents of a container's read / write layer with `docker container diff`, and commit those changes to a new image layer with `docker container commit`. Committing a container as an image in this fashion can be useful when developing an environment inside a container, when you want to capture that environment for reproduction elsewhere.

6 Creating Images with Dockerfiles (1/2)

By the end of this exercise, you should be able to:

- Write a Dockerfile using the `FROM` and `RUN` commands
- Build an image from a Dockerfile
- Anticipate which image layers will be fetched from the cache at build time
- Fetch build history for an image

6.1 Writing and Building a Dockerfile

1. Create a folder called `myimage`, and a text file called `Dockerfile` within that folder. In `Dockerfile`, include the following instructions:

```
FROM centos:7

RUN yum update -y
RUN yum install -y wget
```

This serves as a recipe for an image based on `centos:7`, that has all its default packages updated and `wget` installed on top.

2. Build your image with the build command. Don't miss the `.` at the end; that's the path to your Dockerfile. Since we're currently in the directory `myimage` which contains it, the path is just `.` (here).

```
[centos@node-0 myimage]$ docker image build -t myimage .
```

You'll see a long build output - we'll go through the meaning of this output in a demo later. For now, your image creation was successful if the output ends with `Successfully tagged myimage:latest`.

3. Verify that your new image exists with `docker image ls`, then use your new image to run a container and `wget` something from within that container, just to confirm that everything worked as expected:

```
[centos@node-0 myimage]$ docker container run -it myimage bash
[root@1d86d4093cce /]# wget example.com
[root@1d86d4093cce /]# cat index.html
[root@1d86d4093cce /]# exit
```

You should see the HTML from `example.com`, downloaded by `wget` from within your container.

4. It's also possible to pipe a Dockerfile in from STDIN; try rebuilding your image with the following:

```
[centos@node-0 myimage]$ cat Dockerfile | docker image build -t myimage -f - .
```

(This is useful when reading a Dockerfile from a remote location with `curl`, for example).

6.2 Using the Build Cache

In the previous step, the second time you built your image should have completed immediately, with each step except the first reporting using `cache`. Cached build steps will be used until a change in the Dockerfile is found by the builder.

1. Open your Dockerfile and add another `RUN` step at the end to install `vim`:

```
FROM centos:7

RUN yum update -y
RUN yum install -y wget
RUN yum install -y vim
```

2. Build the image again as before; which steps is the cache used for?
3. Build the image again; which steps use the cache this time?
4. Swap the order of the two `RUN` commands for installing `wget` and `vim` in the Dockerfile:

```
FROM centos:7

RUN yum update -y
RUN yum install -y vim
RUN yum install -y wget
```

Build one last time. Which steps are cached this time?

6.3 Using the history Command

1. The `docker image history` command allows us to inspect the build cache history of an image. Try it with your new image:

```
[centos@node-0 myimage]$ docker image history myimage:latest
```

IMAGE	CREATED	CREATED BY	SIZE
-------	---------	------------	------

f2e85c162453	8 seconds ago	/bin/sh -c yum install -y wget	87.2MB
93385ea67464	12 seconds ago	/bin/sh -c yum install -y vim	142MB
27ad488e6b79	3 minutes ago	/bin/sh -c yum update -y	86.5MB
5182e96772bf	44 hours ago	/bin/sh -c <i> #(nop) CMD ["/bin/bash"]</i>	<i>0B</i>
<missing>	44 hours ago	/bin/sh -c <i> #(nop) LABEL org.label-schema....</i>	<i>0B</i>
<missing>	44 hours ago	/bin/sh -c <i> #(nop) ADD file:6340c690b08865d...</i>	<i>200MB</i>

Note the image id of the layer built for the `yum update` command.

2. Replace the two `RUN` commands that installed `wget` and `vim` with a single command:

```
...
RUN yum install -y wget vim
```

3. Build the image again, and run `docker image history` on this new image. How has the history changed?

6.4 Conclusion

In this exercise, we've seen how to write a basic Dockerfile using `FROM` and `RUN` commands, some basics of how image caching works, and seen the `docker image history` command. Using the build cache effectively is crucial for images that involve lengthy compile or download steps. In general, moving commands that change frequently as late as possible in the Dockerfile will minimize build times. We'll see some more specific advice on this later in this lesson.

7 Creating Images with Dockerfiles (2/2)

By the end of this exercise, you should be able to:

- Define a default process for an image to containerize by using the `ENTRYPOINT` or `CMD` Dockerfile commands
- Understand the differences and interactions between `ENTRYPOINT` and `CMD`

7.1 Setting Default Commands

1. Add the following line to the bottom of your Dockerfile from the last exercise:

```
...
CMD ["ping", "127.0.0.1", "-c", "5"]
```

This sets `ping` as the default command to run in a container created from this image, and also sets some parameters for that command.

2. Rebuild your image:

```
[centos@node-0 myimage]$ docker image build -t myimage .
```

3. Run a container from your new image with no command provided:

```
[centos@node-0 myimage]$ docker container run myimage
```

You should see the command provided by the `CMD` parameter in the Dockerfile running.

4. Try explicitly providing a command when running a container:

```
[centos@node-0 myimage]$ docker container run myimage echo "hello world"
```

Providing a command in `docker container run` overrides the command defined by `CMD`.

5. Replace the `CMD` instruction in your Dockerfile with an `ENTRYPOINT`:

```
...
ENTRYPOINT ["ping"]
```


6. Build the image and use it to run a container with no process arguments:

```
[centos@node-0 myimage]$ docker image build -t myimage .  
[centos@node-0 myimage]$ docker container run myimage
```

You'll get an error. What went wrong?

7. Try running with an argument after the image name:

```
[centos@node-0 myimage]$ docker container run myimage 127.0.0.1
```

You should see a successful ping output. Tokens provided after an image name are sent as arguments to the command specified by ENTRYPOINT.

7.2 Combining Default Commands and Options

1. Open your Dockerfile and modify the ENTRYPOINT instruction to include 2 arguments for the ping command:

```
...  
ENTRYPOINT ["ping", "-c", "3"]
```

2. If CMD and ENTRYPOINT are both specified in a Dockerfile, tokens listed in CMD are used as default parameters for the ENTRYPOINT command. Add a CMD with a default IP to ping:

```
...  
CMD ["127.0.0.1"]
```

3. Build the image and run a container with the defaults:

```
[centos@node-0 myimage]$ docker image build -t myimage .  
[centos@node-0 myimage]$ docker container run myimage
```

You should see it pinging the default IP, 127.0.0.1.

4. Run another container with a custom IP argument:

```
[centos@node-0 myimage]$ docker container run myimage 8.8.8.8
```

This time, you should see a ping to 8.8.8.8. Explain the difference in behavior between these two last containers.

7.3 Conclusion

In this exercise, we encountered the Dockerfile commands CMD and ENTRYPOINT. These are useful for defining the default process to run as PID 1 inside the container right in the Dockerfile, making our containers more like executables and adding clarity to exactly what process was meant to run in a given image's containers.

8 Multi-Stage Builds

By the end of this exercise, you should be able to:

- Write a Dockerfile that describes multiple images, which can copy files from one image to the next.
- Enable BuildKit for faster build times

8.1 Defining a multi-stage build

1. Make a new folder named 'multi' to do this exercise in, and cd into it.
2. Add a file hello.c to the multi folder containing **Hello World** in C:

```
#include <stdio.h>

int main (void)
{
    printf ("Hello, world!\n");
    return 0;
}
```

3. Try compiling and running this right on the host OS:

```
[centos@node-0 multi]$ gcc -Wall hello.c -o hello
[centos@node-0 multi]$ ./hello
```

4. Now let's Dockerize our hello world application. Add a Dockerfile to the multi folder with this content:

```
FROM alpine:3.5
RUN apk update && \
    apk add --update alpine-sdk
RUN mkdir /app
WORKDIR /app
COPY hello.c /app
RUN mkdir bin
RUN gcc -Wall hello.c -o bin/hello
CMD /app/bin/hello
```

5. Build the image and note its size:

```
[centos@node-0 multi]$ docker image build -t my-app-large .
[centos@node-0 multi]$ docker image ls | grep my-app-large
```

REPOSITORY	TAG	IMAGE ID	CREATED	SIZE
my-app-large	latest	a7d0c6fe0849	3 seconds ago	189MB

6. Test the image to confirm it was built successfully:

```
[centos@node-0 multi]$ docker container run my-app-large
```

It should print "hello world" in the console.

7. Update your Dockerfile to use an AS clause on the first line, and add a second stanza describing a second build stage:

```
FROM alpine:3.5 AS build
RUN apk update && \
    apk add --update alpine-sdk
RUN mkdir /app
WORKDIR /app
COPY hello.c /app
RUN mkdir bin
RUN gcc -Wall hello.c -o bin/hello

FROM alpine:3.5
COPY --from=build /app/bin/hello /app/hello
CMD /app/hello
```

8. Build the image again and compare the size with the previous version:

```
[centos@node-0 multi]$ docker image build -t my-app-small .
[centos@node-0 multi]$ docker image ls | grep 'my-app-'
```

REPOSITORY	TAG	IMAGE ID	CREATED	SIZE
------------	-----	----------	---------	------

my-app-small	latest	f49ec3971aa6	6 seconds ago	4.01MB
my-app-large	latest	a7d0c6fe0849	About a minute ago	189MB

As expected, the size of the multi-stage build is much smaller than the large one since it does not contain the Alpine SDK.

9. Finally, make sure the app works:

```
[centos@node-0 multi]$ docker container run --rm my-app-small
```

You should get the expected 'Hello, World!' output from the container with just the required executable.

8.2 Building Intermediate Images

In the previous step, we took our compiled executable from the first build stage, but that image wasn't tagged as a regular image we can use to start containers with; only the final FROM statement generated a tagged image. In this step, we'll see how to persist whichever build stage we like.

1. Build an image from the build stage in your Dockerfile using the `--target` flag:

```
[centos@node-0 multi]$ docker image build -t my-build-stage --target build .
```

Notice all its layers are pulled from the cache; even though the build stage wasn't tagged originally, its layers are nevertheless persisted in the cache.

2. Run a container from this image and make sure it yields the expected result:

```
[centos@node-0 multi]$ docker container run -it --rm my-build-stage /app/bin/hello
```

3. List your images again to see the size of `my-build-stage` compared to the small version of the app.

8.3 Optional: Building from Scratch

So far, every image we've built has been based on a pre-existing image, referenced in the FROM command. But what if we want to start from nothing, and build a completely original image? For this, we can build FROM `scratch`.

1. In a new directory `~/scratch`, create a file named `sleep.c` that just launches a sleeping process for an hour:

```
#include <stdio.h>
#include <unistd.h>
int main()
{
    int delay = 3600; //sleep for 1 hour
    printf ("Sleeping for %d second(s)...\n", delay);
    sleep(delay);
    return 0;
}
```

2. Create a file named Dockerfile to build this sleep program in a build stage, and then copy it to a scratch-based image:

```
FROM alpine:3.8 AS build
RUN ["apk", "update"]
RUN ["apk", "add", "--update", "alpine-sdk"]
COPY sleep.c /
RUN ["gcc", "-static", "sleep.c", "-o", "sleep"]

FROM scratch
COPY --from=build /sleep /sleep
CMD ["/sleep"]
```

This image will contain nothing but our executable and the bare minimum file structure Docker needs to stand up a container filesystem. Note we're statically linking the `sleep.c` binary, so it will have everything it needs bundled along with it, not relying on the rest of the container's filesystem for anything.

3. Build your image:

```
[centos@node-0 scratch]$ docker image build -t sleep:scratch .
```

4. List your images, and search for the one you just built:

```
[centos@node-0 scratch]$ docker image ls | grep scratch
```

REPOSITORY	TAG	IMAGE ID	CREATED	SIZE
sleep	scratch	1b68b20a85a8	9 minutes ago	128kB

This image is only 128 kB, as tiny as possible.

5. Run your image, and check out its filesystem; we can't list directly inside the container, since `ls` isn't installed in this ultra-minimal image, so we have to find where this container's filesystem is mounted on the host. Start by finding the PID of your sleep process after its running:

```
[centos@node-0 scratch]$ docker container run --name sleeper -d sleep:scratch
[centos@node-0 scratch]$ docker container top sleeper
```

UID	PID	PPID	C	STIME	TTY	TIME	CMD
root	1190	1174	0	15:21	?	00:00:00	/sleep

In this example, the PID for sleep is 1190.

6. List your container's filesystem from the host using this PID:

```
[centos@node-0 scratch]$ sudo ls /proc/<PID>/root
```

```
dev  etc  proc  sleep  sys
```

We see not only our binary `sleep` but a bunch of other folders and files. Where does these come from? `runC`, the tool for spawning and running containers, requires a json config of the container and a root file system. At runtime, Docker Engine adds these minimum requirements to form the most minimal container filesystem possible.

7. Clean up by deleting your container:

```
[centos@node-0 scratch]$ docker container rm -f sleeper
```

8.4 Optional: Enabling BuildKit

In addition to the default builder, BuildKit can be enabled to take advantages of some optimizations of the build process.

1. Back in the `~/multi` directory, turn on BuildKit:

```
[centos@node-0 multi]$ export DOCKER_BUILDKIT=1
```

2. Add an AS label to the final stage of your Dockerfile (this is not strictly necessary, but will make the output in the next step easier to understand):

```
...

FROM alpine:3.5 AS prod
RUN apk update
COPY --from=build /app/bin/hello /app/hello
CMD /app/hello
```

3. Re-build `my-app-small`, without the cache:

```
[centos@node-0 multi]$ docker image build --no-cache -t my-app-small-bk .

[+] Building 15.5s (14/14) FINISHED
=> [internal] load Dockerfile
=> => transferring dockerfile: 97B
=> [internal] load .dockerignore
=> => transferring context: 2B
=> [internal] load metadata for docker.io/library/alpine:3.5
=> CACHED [prod 1/3] FROM docker.io/library/alpine:3.5
=> [internal] load build context
=> => transferring context: 87B
=> CACHED [internal] helper image for file operations
=> [build 2/6] RUN apk update && apk add --update alpine-sdk
=> [prod 2/3] RUN apk update
=> [build 3/6] RUN mkdir /app
=> [build 4/6] COPY hello.c /app
=> [build 5/6] RUN mkdir bin
=> [build 6/6] RUN gcc -Wall hello.c -o bin/hello
=> [prod 3/3] COPY --from=build /app/bin/hello /app/hello
=> exporting to image
=> => exporting layers
=> => writing image sha256:22de288...
=> => naming to docker.io/library/my-app-small-bk
```

Notice the lines marked like [prod 2/3] and [build 4/6]: prod and build in this context are the AS labels you applied to the FROM lines in each stage of your build in the Dockerfile; from the above output, you can see that the build stages were built in parallel. Every step of the final image was completed while the build environment image was being created; the prod environment image creation was only blocked at the COPY instruction since it required a file from the completed build image.

4. Comment out the COPY instruction in the prod image definition in your Dockerfile, and rebuild; the build image is skipped. BuildKit recognized that the build stage was not necessary for the image being built, and skipped it.
5. Turn off BuildKit:

```
[centos@node-0 multi]$ export DOCKER_BUILDKIT=0
```

8.5 Conclusion

In this exercise, you created a Dockerfile defining multiple build stages. Being able to take artifacts like compiled binaries from one image and insert them into another allows you to create very lightweight images that do not include developer tools or other unnecessary components in your production-ready images, just like how you currently probably have separate build and run environments for your software. This will result in containers that start faster, and are less vulnerable to attack.

9 Managing Images

By the end of this exercise, you should be able to:

- Rename and retag an image
- Push and pull images from the public registry
- Delete image tags and image layers, and understand the difference between the two operations

9.1 Making an Account on Docker's Hosted Registry

1. If you don't have one already, head over to <https://hub.docker.com> and make an account.
For the rest of this workshop, <Docker ID> refers to the username you chose for this account.

9.2 Tagging and Listing Images

1. Download the centos:7 image from Docker Hub:

```
[centos@node-0 ~]$ docker image pull centos:7
```

2. Make a new tag of this image:

```
[centos@node-0 ~]$ docker image tag centos:7 my-centos:dev
```

Note no new image has been created; my-centos:dev is just a pointer pointing to the same image as centos:7.

3. List your images:

```
[centos@node-0 ~]$ docker image ls
```

You should have centos:7 and my-centos:dev both listed, but they ought to have the same hash under image ID, since they're actually the same image.

9.3 Sharing Images on Docker Hub

1. Push your image to Docker Hub:

```
[centos@node-0 ~]$ docker image push my-centos:dev
```

You should get a denied: requested access to the resource is denied error.

2. Login by doing `docker login`, and try pushing again. The push fails again because we haven't namespaced our image correctly for distribution on Docker Hub; all images you want to share on Docker Hub must be named like <Docker ID>/<repo name>[:<optional tag>].
3. Retag your image to be namespaced properly, and push again:

```
[centos@node-0 ~]$ docker image tag my-centos:dev <Docker ID>/my-centos:dev  
[centos@node-0 ~]$ docker image push <Docker ID>/my-centos:dev
```

4. Search Docker Hub for your new <Docker ID>/my-centos repo, and confirm that you can see the :dev tag therein.
5. Next, write a Dockerfile that uses <Docker ID>/my-centos:dev as its base image, and installs any application you like on top of that. Build the image, and simultaneously tag it as :1.0:

```
[centos@node-0 ~]$ docker image build -t <Docker ID>/my-centos:1.0 .
```

6. Push your :1.0 tag to Docker Hub, and confirm you can see it in the appropriate repository.
7. Finally, list the images currently on your node with `docker image ls`. You should still have the version of your image that wasn't namespaced with your Docker Hub user name; delete this using `docker image rm`:

```
[centos@node-0 ~]$ docker image rm my-centos:dev
```

Only the tag gets deleted, not the actual image. The image layers are still referenced by another tag.

9.4 Conclusion

In this exercise, we practiced tagging images and exchanging them on the public registry. The namespacing rules for images on registries are *mandatory*: user-generated images to be exchanged on the public registry must be named like `<Docker ID>/<repo name>[:<optional tag>]`; official images in the Docker registry just have the repo name and tag.

Also note that as we saw when building images, image names and tags are just pointers; deleting an image with `docker image rm` just deletes that pointer if the corresponding image layers are still being referenced by another such pointer. Only when the last pointer is deleted are the image layers actually destroyed by `docker image rm`.

10 Database Volumes

By the end of this exercise, you should be able to:

- Provide a docker volume as a database backing to Postgres
- Recover a Postgres database from volume contents after destroying the original Postgres container

10.1 Launching Postgres

1. Download a postgres image, and look at its history to determine its default volume usage:

```
[centos@node-0 ~]$ docker image pull postgres:9-alpine
[centos@node-0 ~]$ docker image inspect postgres:9-alpine

...
"Volumes": {
    "/var/lib/postgresql/data": {}
},
...
```

You should see a `Volumes` block like the above, indicating that those paths in the container filesystem will get volumes automatically mounted to them when a container is started based on this image.

2. Set up a running instance of this postgres container:

```
[centos@node-0 ~]$ docker container run --name some-postgres \
    -v db_backing:/var/lib/postgresql/data \
    -d postgres:9-alpine
```

Notice the explicit volume mount, `-v db_backing:/var/lib/postgresql/data`; if we hadn't done this, a randomly named volume would have been mounted to the container's `/var/lib/postgresql/data`. Naming the volume explicitly is a best practice that will become useful when we start mounting this volume in multiple containers.

10.2 Writing to the Database

1. The `psql` command line interface to postgres comes packaged with the postgres image; spawn it as a child process in your postgres container interactively, to create a postgres terminal:

```
[centos@node-0 ~]$ docker container exec \
    -it some-postgres psql -U postgres
```

2. Create an arbitrary table in the database:

```
postgres=# CREATE TABLE PRODUCTS(PRICE FLOAT, NAME TEXT);
postgres=# INSERT INTO PRODUCTS VALUES('18.95', 'widget');
postgres=# INSERT INTO PRODUCTS VALUES('1.45', 'sprocket');
```

Double check you created the table you expected, and then quit this container:

```
postgres=# SELECT * FROM PRODUCTS;

 price | name 
-----+-----
  18.95 | widget
   1.45 | sprocket
(2 rows)

postgres=# \q
```

3. Delete the postgres container:

```
[centos@node-0 ~]$ docker container rm -f some-postgres
```

4. Create a new postgres container, mounting the db_backing volume just like last time:

```
[centos@node-0 ~]$ docker container run \
  --name some-postgres \
  -v db_backing:/var/lib/postgresql/data \
  -d postgres:9-alpine
```

5. Reconnect a psql interface to your database, also like before:

```
[centos@node-0 ~]$ docker container exec \
  -it some-postgres psql -U postgres
```

6. List the contents of the PRODUCTS table:

```
postgres=# SELECT * FROM PRODUCTS;
```

The contents of the database have survived the deletion and recreation of the database container; this would not have been true if the database was keeping its data in the writable container layer. As above, use \q to quit from the postgres prompt.

10.3 Conclusion

Whenever data needs to live longer than the lifecycle of a container, it should be pushed out to a volume outside the container's filesystem; numerous popular databases are containerized using this pattern.

11 Introduction to Container Networking

By the end of this exercise, you should be able to:

- Create docker bridge networks and attach containers to them
- Design networks of containers that can successfully resolve each other via DNS and reach each other across a Docker software defined network.

11.1 Inspecting the Default Bridge

1. In the dropdown menu at the top of the Strigo webpage, click into node-1. See what networks are present on your host:

```
[centos@node-1 ~]$ docker network ls
```

You should have entries for host, none, and bridge.

2. Find some metadata about the default bridge network:


```
[centos@node-1 ~]$ docker network inspect bridge
```

Note especially the private subnet assigned by Docker's IPAM driver to this network. The first IP in this range is used as the network's gateway, and the rest will be assigned to containers as they join the network.

3. See similar info from common networking tools:

```
[centos@node-1 ~]$ ip addr
```

Note the bridge network's gateway corresponds to the IP of the docker0 device in this list. docker0 is the linux bridge itself, while bridge is the name of the default Docker network that uses that bridge.

4. Use brctl to see connections to the docker0 bridge:

```
[centos@node-1 ~]$ brctl show docker0
```

```
bridge name bridge id          STP enabled interfaces
docker0      8000.02427f12c30b    no
```

At the moment, there are no connections to docker0.

11.2 Connecting Containers to docker0

1. Start a container and reexamine the network; the container is listed as connected to the network, with an IP assigned to it from the bridge network's subnet:

```
[centos@node-1 ~]$ docker container run --name u1 -dt centos:7
[centos@node-1 ~]$ docker network inspect bridge
```

```
...
  "Containers": {
    "11da9b7db065f971f78aebf14b706b0b85f07ec10dbf6f0773b1603f48697961": {
      "Name": "u1",
      "EndpointID": "670c495...",
      "MacAddress": "02:42:ac:11:00:02",
      "IPv4Address": "172.17.0.2/16",
      "IPv6Address": ""
    }
  },
  ...
```

2. Inspect the network interfaces with ip and brctl again, now that you have a container running:

```
[centos@node-1 ~]$ ip addr
```

```
...
5: veth6f244c3@if4: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc noqueue
   master docker0 state UP
   link/ether aa:71:82:6c:f3:88 brd ff:ff:ff:ff:ff:ff link-netnsid 0
   inet6 fe80::a871:82ff:fe6c:f388/64 scope link
       valid_lft forever preferred_lft forever
```

```
[centos@node-1 ~]$ brctl show docker0
```

```
bridge name bridge id          STP enabled interfaces
docker0      8000.02427f12c30b    no      veth6f244c3
```

ip addr indicates a veth endpoint has been created and plugged into the docker0 bridge, as indicated by master docker0, and that it is connected to device index 4 in this case (indicated by the @if4 suffix to the veth device name above). Similarly, brctl now shows this veth connection on docker0 (notice that the ID for the veth connection matches in both utilities).

3. Launch a bash shell in your container, and look for the eth0 device therein:

```
[centos@node-1 ~]$ docker container exec -it u1 bash
[root@11da9b7db065 /]# yum install -y iproute
[root@11da9b7db065 /]# ip addr

...
4: eth0@if5: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc noqueue
    state UP group default
    link/ether 02:42:ac:11:00:02 brd ff:ff:ff:ff:ff:ff link-netnsid 0
    inet 172.17.0.2/16 scope global eth0
        valid_lft forever preferred_lft forever
```

We see that the eth0 device in this namespace is in fact the device that the veth connection in the host namespace indicated it was attached to, and vice versa - eth0@if5 indicates it is plugged into networking interface number 5, which we saw above was the other end of the veth connection. Docker has created a veth connection with one end in the host's docker0 bridge, and the other providing the eth0 device in the container.

11.3 Defining Additional Bridge Networks

In the last step, we investigated the default bridge network; now let's try making our own. User defined bridge networks work exactly the same as the default one, but provide DNS lookup by container name, and are firewalled from other networks by default.

1. Create a bridge network by using the bridge driver with `docker network create`:

```
[centos@node-1 ~]$ docker network create --driver bridge my_bridge
```

2. Launch a container connected to your new network via the `--network` flag:

```
[centos@node-1 ~]$ docker container run --name=u2 --network=my_bridge -dt centos:7
```

3. Use the `inspect` command to investigate the network settings of this container:

```
[centos@node-1 ~]$ docker container inspect u2
```

my_bridge should be listed under the Networks key.

4. Launch another container, this time interactively:

```
[centos@node-1 ~]$ docker container run --name=u3 --network=my_bridge -it centos:7
```

5. From inside container u3, ping u2 by name: `ping u2`. The ping succeeds, since Docker is able to resolve container names when they are attached to a custom network.
6. Try starting a container on the default network, and pinging u1 by name:

```
[centos@node-1 ~]$ docker container run centos:7 ping u1
```

```
ping: u1: Name or service not known
```

The ping fails; even though the containers are both attached to the bridge network, Docker does not provide name lookup on this default network. Try the same command again, but using u1's IP instead of name, and you should be successful.

7. Finally, try pinging u1 by IP, this time from container u2:

```
[centos@node-1 ~]$ docker container exec u2 ping <u1 IP>
```

The ping fails, since the containers reside on different networks; all Docker networks are firewalled from each other by default.

8. Clean up your containers and networks:

```
[centos@node-1 ~]$ docker container rm -f $(docker container ls -aq)
[centos@node-1 ~]$ docker network rm my_bridge
```

11.4 Conclusion

In this exercise, you explored the fundamentals of container networking. The key take away is that *containers on separate networks are firewalled from each other by default*. This should be leveraged as much as possible to harden your applications; if two containers don't need to talk to each other, put them on separate networks.

You also explored a number of API objects:

- `docker network ls` lists all networks on the host
- `docker network inspect <network name>` gives more detailed info about the named network
- `docker network create --driver <driver> <network name>` creates a new network using the specified driver; so far, we've only seen the bridge driver, for creating a linux bridge based network.
- `docker network connect <network name> <container name or id>` connects the specified container to the specified network after the container is running; the `--network` flag in `docker container run` achieves the same result at container launch.
- `docker container inspect <container name or id>` yields, among other things, information about the networks the specified container is connected to.

12 Container Port Mapping

By the end of this exercise, you should be able to:

- Forward traffic from a port on the docker host to a port inside a container's network namespace
- Define ports to automatically expose in a Dockerfile

12.1 Port Mapping at Runtime

1. Run an nginx container with no special port mappings:

```
[centos@node-1 ~]$ docker container run -d nginx
```

nginx stands up a landing page at `<ip>:80`. If you try to visit this at your host or container's IP it won't be visible; no external traffic can make it past the linux bridge's firewall to the nginx container.

2. Now run an nginx container and map port 80 on the container to port 5000 on your host using the `-p` flag:

```
[centos@node-1 ~]$ docker container run -d -p 5000:80 nginx
```

Note that the syntax is: `-p [host-port]:[container-port]`.

3. Verify the port mappings with the `docker container port` command

```
[centos@node-1 ~]$ docker container port <container id>
```

```
80/tcp -> 0.0.0.0:5000
```

4. Visit your nginx landing page at `<host ip>:5000`, e.g. using `curl -4 localhost:5000`, just to confirm it's working as expected.

12.2 Exposing Ports from the Dockerfile

1. In addition to manual port mapping, we can expose some ports for automatic port mapping on container startup using a Dockerfile. In a new directory `~/port`, create a Dockerfile:

```
FROM nginx
```

```
EXPOSE 80
```

2. Build your image as `my_nginx`:

```
[centos@node-1 port]$ docker image build -t my_nginx .
```

3. Use the `-P` flag when running to map all ports mentioned in the `EXPOSE` directive:

```
[centos@node-1 port]$ docker container run -d -P my_nginx
```

4. Use `docker container ls` or `docker container port` to find out which host ports were used, then visit your nginx landing page in a browser at `<node-1 public IP>:<port>`.

5. Clean up your containers:

```
[centos@node-1 port]$ docker container rm -f $(docker container ls -aq)
```

12.3 Conclusion

In this exercise, we saw how to explicitly map ports from our container's network stack onto ports of our host at runtime with the `-p` option to `docker container run`, or more flexibly in our Dockerfile with `EXPOSE`, which will result in the listed ports inside our container being mapped to random available ports on our host. In both cases, Docker is writing iptables rules to forward traffic from the host to the appropriate port in the container's network namespace.

13 Creating a Swarm

By the end of this exercise, you should be able to:

- Create a swarm in high availability mode
- Set default address pools
- Check necessary connectivity between swarm nodes
- Configure the swarm's TLS certificate rotation

13.1 Starting Swarm

1. Switch back to `node-0` by using the dropdown menu. Then, initialize swarm and create a cluster with a default address pool for a discontinuous address range of `10.85.0.0/16` and `10.91.0.0/16` with a default subnet size of 128 addresses. This will be your first manager node:

```
[centos@node-0 ~]$ docker swarm init \
  --default-addr-pool 10.85.0.0/16 \
  --default-addr-pool 10.91.0.0/16 \
  --default-addr-pool-mask-length 25
```

2. Confirm that Swarm Mode is active and that the default address pool configuration has been registered by inspecting the output of:

```
[centos@node-0 ~]$ docker system info

...
Swarm: active
...
  Default Address Pool: 10.85.0.0/16 10.91.0.0/16
  SubnetSize: 25
...
```

3. See all nodes currently in your swarm by doing:

```
[centos@node-0 ~]$ docker node ls
```

A single node is reported in the cluster.

4. Change the certificate rotation period from the default of 90 days to one week, and rotate the certificate now:

```
[centos@node-0 ~]$ docker swarm ca --rotate --cert-expiry 168h
```

Note that the `docker swarm ca [options]` command *must* receive the `--rotate` flag, or all other flags will be ignored.

5. Display UDP and TCP activity on your manager:

```
[centos@node-0 ~]$ sudo netstat -plunt | grep -E "2377|7946|4789"
```

You should see (at least) TCP+UDP 7946, UDP 4789, and TCP 2377. What are each of these ports for?

13.2 Adding Workers to the Swarm

A single node swarm is not a particularly interesting swarm; let's add some workers to really see Swarm Mode in action.

1. On your manager node (node-0), get the swarm 'join token' you'll use to add worker nodes to your swarm:

```
[centos@node-0 ~]$ docker swarm join-token worker
```

2. Switch to node-1 from the dropdown menu.
3. Paste in the join token you found in the first step above. node-1 will join the swarm as a worker.
4. Inspect the network on node-1 with `sudo netstat -plunt` like you did for the manager node. Are the same ports open? Why or why not?
5. Do `docker node ls` after switching back to the manager node-0 again, and you should see both your nodes and their status. Note that `docker node ls` won't work on a worker node, as the cluster status is maintained only by the manager nodes.
6. Finally, use the same join token to add two more workers (node-2 and node-3) to your swarm. When you're done, confirm that `docker node ls` on your one manager node reports 4 nodes in the cluster - one manager, and three workers.

13.3 Promoting Workers to Managers

At this point, our swarm has a single manager, node-0. If this node goes down, we'll lose the ability to maintain and schedule workloads on our swarm. In a real deployment, this is unacceptable; we need some redundancy to our system, and Swarm achieves this by allowing a raft consensus of multiple managers to preserve swarm state.

1. Promote two of your workers to manager status by executing, on the current manager node:

```
[centos@node-0 ~]$ docker node promote node-1 node-2
```

2. Finally, do a `docker node ls` to check and see that you now have three managers. Note that manager nodes also count as worker nodes - tasks can still be scheduled on them as normal.

13.4 Conclusion

In this exercise, you set up a basic high-availability swarm. In practice, it is crucial to have at least 3 (and always an odd number) of managers in order to ensure high availability of your cluster, and to ensure that the management, control, and data plane communications a swarm maintains can proceed unimpeded between all nodes.

14 Starting a Service

By the end of this exercise, you should be able to:

- Schedule a docker service across a swarm
- Predict and understand the scoping behavior of docker overlay networks
- Scale a service on swarm up or down

14.1 Creating an Overlay Network and Service

1. Create a multi-host overlay network which you can connect your service to:

```
[centos@node-0 ~]$ docker network create --driver overlay my_overlay
```

2. Verify that the network subnet was taken from the address pool defined when creating your swarm:

```
[centos@node-0 ~]$ docker network inspect my_overlay
...
"Subnet": "10.85.0.0/25",
"Gateway": "10.85.0.1"
...
```

The overlay network has been assigned a subnet from the address pool we specified when creating our swarm.

3. Create a service featuring an alpine container pinging Google resolvers, plugged into your overlay network:

```
[centos@node-0 ~]$ docker service create --name pinger \
--network my_overlay alpine ping 8.8.8.8
```

Note the syntax is a lot like `docker container run`; an image (`alpine`) is specified, followed by the PID 1 process for that container (`ping 8.8.8.8`).

4. Get some information about the currently running services:

```
[centos@node-0 ~]$ docker service ls
```

ID	NAME	MODE	REPLICAS	IMAGE
bmthsr0m6xvr	pinger	replicated	1/1	alpine:latest

5. Check which node the container was created on:

```
[centos@node-0 ~]$ docker service ps pinger
```

ID	NAME	IMAGE	NODE	DESIRED STATE	CURRENT STATE
lmm...	pinger.1	alpine:latest	node-0	Running	Running 42 seconds ago

In my case, the one container we started for this service was scheduled on `node-0`.

6. Scale up the number of concurrent tasks that our pinger service is running to 3:

```
[centos@node-0 ~]$ docker service update pinger --replicas=3
```

```
pinger
overall progress: 3 out of 3 tasks
1/3: running [=====>]
2/3: running [=====>]
3/3: running [=====>]
verify: Service converged
```

7. Now run `docker service ps pinger` to inspect the service. How were tasks distributed across your swarm?

- Run `docker network inspect my_overlay` on any node that has a `pinger` task scheduled on it. Look for the `Containers` key in the output; it indicates the containers on this node attached to the `my_overlay` network. Also, look for the `Peers` list; mine looks like:

```
"Peers": [
  {
    "Name": "967da1a0349f",
    "IP": "10.10.12.136"
  },
  {
    "Name": "d1cd9f4a25bb",
    "IP": "10.10.35.36"
  },
  {
    "Name": "d7e00d4376ca",
    "IP": "10.10.57.19"
  }
]
```

Challenge: Looking at your own `Peers` list, what do the IPs correspond to?

14.2 Inspecting Service Logs

- Manager nodes can assemble all logs for all tasks of a given service:

```
[centos@node-0 ~]$ docker service logs pinger
```

The ping logs for all 3 pinging containers will be displayed.

- If instead you'd like to see the logs of a single task, on a manager node run `docker service ps pinger`, choose any task ID, and run `docker service logs <task ID>`.

14.3 Cleanup

- Remove all existing services, in preparation for future exercises:

```
[centos@node-0 ~]$ docker service rm $(docker service ls -q)
```

14.4 Conclusion

In this exercise, we saw the basics of creating, scheduling and updating services. A common mistake people make is thinking that a service is just the containers scheduled by the service; in fact, a Docker service is the definition of *desired state* for those containers. Changing a service definition does not in general change containers directly; it causes them to get rescheduled by Swarm in order to match their new desired state.

15 Node Failure Recovery

By the end of this exercise, you should be able to:

- Anticipate swarm scheduling decisions when nodes fail and recover
- Force swarm to reallocate workload across a swarm

15.1 Setting up a Service

1. Set up a myProxy service with four replicas on one of your manager nodes:

```
[centos@node-0 ~]$ docker service create --replicas 4 --name myProxy nginx
```

2. Now watch the output of `docker service ps` on the same node:

```
[centos@node-0 ~]$ watch docker service ps myProxy
```

This should be stable for now, but will let us monitor scheduling updates as we interfere with the rest of our swarm.

15.2 Simulating Node Failure

1. Switch into node-3, and simulate a node failure by rebooting it:

```
[centos@node-3 ~]$ sudo reboot now
```

2. Back on your manager node, watch the updates to `docker service ps`; what happens to the task running on the rebooted node? Look at its desired state, any other tasks that get scheduled with the same name, and keep watching until node-3 comes back online.

15.3 Force Rebalancing

By default, if a node fails and rejoins a swarm it *will not* get its old workload back; if we want to redistribute workload across a swarm after new nodes join (or old nodes rejoin), we need to force-rebalance our tasks

1. Back on the manager node, exit the watch mode with CTRL+C.
2. Force rebalance the tasks:

```
[centos@node-0 ~]$ docker service update --force myProxy
```

3. After the service converges, check which nodes the service tasks are scheduled on:

```
[centos@node-0 ~]$ docker service ps myProxy
... NAME                NODE        DESIRED STATE    CURRENT STATE
... myProxy.1           node-0      Running          Running about a minute ago
... \_ myProxy.1        node-0      Shutdown         Shutdown about a minute ago
... myProxy.2           node-3      Running          Running about a minute ago
... \_ myProxy.2        node-1      Shutdown         Shutdown about a minute ago
... myProxy.3           node-1      Running          Running about a minute ago
... \_ myProxy.3        node-2      Shutdown         Shutdown about a minute ago
... myProxy.4           node-2      Running          Running about a minute ago
... \_ myProxy.4        node-0      Shutdown         Shutdown about a minute ago
... \_ myProxy.4        node-3      Shutdown         Shutdown 2 minutes ago
```

The `_` shape indicate *ancestor* tasks which have been shut down and replaced by a new task, typically after reconfiguring the service or rebalancing like we've done here. Once the rebalance is complete, the current tasks for the myProxy service should be evenly distributed across your swarm.

15.4 Cleanup

1. On your manager node, remove all existing services, in preparation for future exercises:

```
[centos@node-0 ~]$ docker service rm $(docker service ls -q)
```


15.5 Conclusion

In this exercise, you saw swarm's scheduler in action - when a node is lost from the swarm, tasks are automatically rescheduled to restore the state of our services. Note that nodes joining or rejoining the swarm do not get workload automatically reallocated from existing nodes to them; rescheduling only happens when tasks crash, services are first scheduled, or you force a reschedule as above.

16 Swarm Scheduling

By default, the Swarm scheduling algorithm tries to spread workload out roughly evenly across your Swarm. In many cases, we want to exert more nuanced control over what containers get scheduled where, in order to respect resource availability, hardware requirements, application high availability and other concerns. By the end of this exercise, you should be able to:

- Impose CPU and memory resource reservations and limitations
- Schedule services in global or replicated mode, and define appropriate use cases for each
- Schedule tasks on a subset of nodes via label constraints
- Schedule topology-aware services

16.1 Restricting Resource Consumption

By default, containers can consume as much CPU and memory as they want; in practice, unconstrained CPU usage leads to noisy-neighbor problems where one container can starve all other co-located containers of CPU time, and unconstrained memory usage leads to processes getting killed with out-of-memory errors. When scheduling services, we must prevent containers from overconsuming compute resources, and make sure we're scheduling only as many containers on a host as that host can realistically support.

1. Create a service based on the `training/stress:3.0` container, with a few flags to make it consume two full CPUs and allocate a gigabyte of memory. We'll also introduce the `--detach` flag, which sends the service startup to the background:

```
[centos@node-0 ~]$ docker service create --name compute-stress \
  --replicas 4 \
  --detach \
  training/stress:3.0 --vm 2 --vm-bytes 1024M
```

2. Check the resource consumption of your containers on one of your hosts:

```
[centos@node-0 ~]$ docker stats
```

CONTAINER ID	NAME	CPU %	MEM USAGE / LIMIT	MEM %
b97b645e3a4f	compute-stress.4.o7v...	199.38%	1.128GiB / 3.7GiB	30.48%

CTRL+C to escape from the stats view once you've seen it. The one container on this host is consuming two full CPUs and over a gigabyte of memory; anything else scheduled on this node is going to be starved of resources.

3. Remove this service, recreate it with a limitation on how much CPU its containers are allowed to consume, and check the stats again:

```
[centos@node-0 ~]$ docker service rm compute-stress
[centos@node-0 ~]$ docker service create --name compute-stress \
  --replicas 4 \
  --limit-cpu 1 \
  --detach \
  training/stress:3.0 --vm 2 --vm-bytes 1024M
[centos@node-0 ~]$ docker stats
```

CONTAINER ID	NAME	CPU %	MEM USAGE / LIMIT	MEM %
d311d88debd9	compute-stress.2.6w0...	100.52%	1.158GiB / 3.7GiB	31.29%

The `--limit-cpu 1` flag imposes control group limits on our containers, preventing them from consuming more than one core's worth of cycles.

4. We've throttled our CPU consumption above, but that one container is still hogging around a gigabyte of memory; we can similarly limit memory consumption to prevent random out-of-memory process kills from taking down the host. Remove and recreate your service, this time with a memory limit:

```
[centos@node-0 ~]$ docker service rm compute-stress
[centos@node-0 ~]$ docker service create --name compute-stress \
  --replicas 4 \
  --limit-cpu 1 \
  --limit-memory 512M \
  --detach \
  training/stress:3.0 --vm 2 --vm-bytes 1024M
```

5. List your containers with `docker container ls -a`; you should see a list of exited containers. Inspect one of them, and look for its out-of-memory status:

```
[centos@node-0 ~]$ docker container ls -a
CONTAINER ID   IMAGE                CREATED          STATUS
3eb69ace1d66   training/stress:3.0 ... 4 seconds ago   Created
eac5f4c35142   training/stress:3.0 ... 5 seconds ago   Exited (1) 5 seconds ago

[centos@node-0 ~]$ docker container inspect eac5f4c35142 | grep OOM

"OOMKilled": true,
```

`--limit-memory` works a little differently than `--limit-cpu`; while the CPU limit throttled a running container, the memory limit kills the container with an Out Of Memory exception if it tries to exceed its memory limit; this way, we avoid the random out-of-memory process kill that the kernel usually imposes, and which can take down a worker by potentially killing the Docker daemon itself.

6. So far, we've limited the amount of resources a container can consume once scheduled, but we still haven't prevented the scheduler from *overprovisioning* containers on a node; we would like to prevent Docker from scheduling more containers on a node than that node can support. Delete and recreate your service one more time, this time without exceeding your memory limit and also imposing a scheduling reservation with `--reserve-memory`:

```
[centos@node-0 ~]$ docker service rm compute-stress
[centos@node-0 ~]$ docker service create --name compute-stress \
  --replicas 4 \
  --limit-cpu 1 \
  --limit-memory 512M \
  --reserve-memory 512M \
  --detach \
  training/stress:3.0 --vm 2 --vm-bytes 128M

[centos@node-0 ~]$ docker service ps compute-stress
```

You should see your four tasks running happily.

7. Now scale up your service to more replicas than your current cluster can support:

```
[centos@node-0 ~]$ docker service update compute-stress --replicas=40 --detach
[centos@node-0 ~]$ docker service ps compute-stress
```

ID	NAME	CURRENT STATE	ERROR
...			

```

z6t...   compute-stress.31   Pending 2 seconds ago   "no suitable node (insufficien..."
xf3...   compute-stress.32   Pending 2 seconds ago   "no suitable node (insufficien..."
...

```

Many of your tasks should be stuck in CURRENT STATE: Pending. Inspect one of them:

```

[centos@node-0 ~]$ docker inspect xf3xigecdfw8

...
  "Status": {
    "Timestamp": "2019-01-29T16:31:45.5859483Z",
    "State": "pending",
    "Message": "pending task scheduling",
    "Err": "no suitable node (insufficient resources on 4 nodes)",
    "PortStatus": {}
  }
...

```

From the Status block in the task info, we can see that this task isn't getting scheduled because there isn't sufficient resources available in the cluster to support it.

Always limit resource consumption using *both* limits and reservations for CPU and memory. Failing to do so is a very common mistake that can lead to a widespread cluster outage; if your cluster is running at near-capacity and one node fails, its workload will get rescheduled to other nodes, potentially causing them to also fail, initiating a cascading cluster outage.

8. Clean up by removing your service:

```

[centos@node-0 ~]$ docker service rm compute-stress

```

16.2 Configuring Global Scheduling

So far, all the services we've created have been run in the default *replicated mode*; as we've seen, Swarm spreads containers out across the cluster, potentially respecting resource reservations. Sometimes, we want to run services that create *exactly one* container on each host in our cluster; this is typically used for deploying daemon, like logging, monitoring, or node management tools which need to run locally on every node.

1. Create a globally scheduled service:

```

[centos@node-0 ~]$ docker service create --mode global \
  --name my-global \
  centos:7 ping 8.8.8.8

```

2. Check what nodes your service tasks were scheduled on:

```

[centos@node-0 ~]$ docker service ps my-global

```

ID	NAME	IMAGE	NODE
rjl91n1i5o4g	competent_taussig.j9hmr8ne6s8ysyb3k3y0wtp9	centos:7	node-3
pzkiv2kpsu26	competent_taussig.c3va5z6je8zpkozgn0cm5kllt	centos:7	node-2
k767q7i1f73t	competent_taussig.afoyo4r860dbwve1h4dm3dsrq	centos:7	node-1
wem26fmzq2k5	competent_taussig.lnki68wnzp0r7456zo6xc46s2	centos:7	node-0

One task is scheduled on every node in the swarm; as you add or remove nodes from the swarm, the global service will be rescaled appropriately.

3. Remove your service with `docker service rm my-global`.

16.3 Scheduling via Node Constraints

Sometimes, we want to confine our containers to specific nodes; for this, we can use *constraints* and *node properties*.

1. Add a label `datacenter` with value `east` to two nodes of your swarm:

```
[centos@node-0 ~]$ docker node update --label-add datacenter=east node-0
[centos@node-0 ~]$ docker node update --label-add datacenter=east node-1
```

2. Add a label `datacenter` with value `west` to the other two nodes:

```
[centos@node-0 ~]$ docker node update --label-add datacenter=west node-2
[centos@node-0 ~]$ docker node update --label-add datacenter=west node-3
```

Note these labels are user-defined; `datacenter` and its values `east` and `west` can be anything you like.

3. Schedule a service constrained to run on nodes labeled as `datacenter=east`:

```
[centos@node-0 ~]$ docker service create --replicas 4 \
  --constraint node.labels.datacenter==east \
  --name east-deploy \
  centos:7 ping 8.8.8.8
```

4. Check what nodes your tasks were scheduled on as above; they should all be on nodes bearing the `datacenter==east` label (node-0 and node-1).
5. Remove your service, and schedule another, this time constrained to run only on worker nodes:

```
[centos@node-0 ~]$ docker service rm east-deploy
[centos@node-0 ~]$ docker service create --replicas 4 \
  --constraint node.role==worker \
  --name worker-only \
  centos:7 ping 8.8.8.8
```

Once again, check where your `worker-only` service tasks got scheduled; they'll all be on node-3, your only worker.

Keep workload off of managers using these selectors, especially in production. If something goes badly wrong with a workload container and causes its host to crash, we don't want to take down a manager and possibly lose our raft consensus. As we've seen, Swarm can recover from losing workers automatically; a lost manager consensus can be much harder to recover from.

6. Clean up by removing this service: `docker service rm worker-only`.

16.4 Scheduling Topology-Aware Services

Oftentimes, we want to schedule workload to be tolerant of faults in our datacenters; we wouldn't want every replica for a service on one power zone or one rack which can go down all at once, for example.

1. Create a service using the `--placement-pref` flag to spread replicas across our `datacenter` label:

```
[centos@node-0 ~]$ docker service create --name my_proxy \
  --replicas=2 --publish 8000:80 \
  --placement-pref spread=node.labels.datacenter \
  nginx
```

There should be `nginx` containers present on nodes with every possible value of the `node.labels.datacenter` label, one in `datacenter=east` nodes, and one in `datacenter=west` nodes.

2. Use `docker service ps my_proxy` as above to check that replicas got spread across the `datacenter` labels.
3. Clean up: `docker service rm my_proxy`.

16.5 Conclusion

In this exercise, we saw how to use resource allocations, global scheduling, labels and node properties to influence scheduling. A few best practices:

- **Always apply memory and CPU limits and reservations**, and in essentially all cases, the limit should be less than or equal to the reservation to ensure nodes are never overprovisioned.
- **Keep workload off of manager** as mentioned above, to prevent application logic bugs from taking down your manager consensus
- **Don't overconstrain your scheduler**: it can be tempting to exert strict control over exactly what gets scheduled exactly where; don't. If a service is constrained to run on a very small set of nodes, and those nodes go down, the service will become unschedulable and suffer an outage. Let your orchestrator's scheduler make decisions as independently as possible in order to maximize workload resilience.

17 Provisioning Swarm Configuration

When deploying an application, especially one meant to be migrated across different environments, it's helpful to be able to provision configuration like environment variables and config files to your services in a modular, pluggable fashion. By the end of this exercise, you should be able to:

- Assemble application components together as a Docker stack
- Provision insecure configuration to service containers using `.env` files and Docker configs
- Provision secure configuration to service containers using Docker secrets

17.1 Creating a Stack

So far, we've run individual services with `docker service create`. As we build more complex applications consisting of multiple components, we'd like a way to capture them all in a single file we can version control and recreate; for this, we can use *stack files*.

1. Create a file called `mystack.yaml` with the following content:

```
version: "3.7"

services:
  db:
    image: postgres:9.6
```

This simple stack file will create a single service named `db`, based on the `postgres:9.6` image.

Docker stack file syntax is based on Docker Compose; we'll see numerous examples of this syntax in this workshop, but if you'd like the full reference, see the docs at <https://dockr.ly/2iHUpeX>.

2. Deploy your stack:

```
[centos@node-0 ~]$ docker stack deploy -c mystack.yaml dbdemo
```

Your service is created, along with a default network for the stack (more on service networking in a future exercise).

3. List your stacks and, see its services:

```
[centos@node-0 ~]$ docker stack ls

NAME                SERVICES            ORCHESTRATOR
dbdemo              1                   Swarm

[centos@node-0 ~]$ docker service ls
```

ID	NAME	MODE	REPLICAS	IMAGE
xb7c19heahku	dbdemo_db	replicated	1/1	postgres:9.6

By default, your service gets named as the stack (dbdemo), concatenated with the key you labeled your service with in your stack file (db).

4. Delete your stack:

```
[centos@node-0 ~]$ docker stack rm dbdemo
```

17.2 Defining and Using .env Files

Many configurations don't have strong security needs, and can be stored and transmitted unencrypted. For these, we can use *Docker config* objects.

1. Create a file called `myvars.env` listing environment variables to define inside your container:

```
POSTGRES_USER=moby
POSTGRES_DB=mydb
```

If defined at postgres startup, these environment variables will set the default username and database for postgres.

2. Modify your stack file so your db service consumes this `.env` file:

```
version: "3.7"

services:
  db:
    image: postgres:9.6
    env_file:
      - myvars.env
```

3. Redeploy your stack:

```
[centos@node-0 ~]$ docker stack deploy -c mystack.yaml dbdemo
```

4. We'd like to confirm that the environment variables got set and had the desired effect; to do so, list all the tasks running for every service in your stack:

```
[centos@node-0 ~]$ docker stack ps dbdemo
```

ID	NAME	IMAGE	NODE	DESIRED STATE	CURRENT STATE
uil...	dbdemo_db.1	postgres:9.6	node-0	Running	Running about a minute ago

As expected, we have one service with one task, which got scheduled in my case on node-0.

5. Find the container corresponding to the single task started for your db service:

```
[centos@node-0 ~]$ docker inspect <task ID> | grep ContainerID

"ContainerID": "b2ffe30...",
```

6. Confirm the environment variables you provisioned actually got set (note you'll have to run this on the node listed in the NODE column in the output of `docker stack ps dbdemo` above, node-0 for me):

```
[centos@node-0 ~]$ docker container inspect <container ID> | grep POSTGRES

"POSTGRES_DB=mydb",
"POSTGRES_USER=moby",
```

- Also on the node running the postgres container, run a command line interface (psql) inside this container to confirm your config was used to correctly set up the default user and database:

```
[centos@node-0 ~]$ docker container exec -it <container ID> psql -U moby -d mydb

psql (9.6.11)
Type "help" for help.

moby=# \du
                                List of roles
Role name |                               Attributes                               | Member of
-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
moby      | Superuser, Create role, Create DB, Replication, Bypass RLS | {}
moby=# \q
```

We can see that the user moby and default database mydb were created as expected.

17.3 Defining and Using Docker Configs

The config we've seen so far is centered around defining environment variables in our containers, but oftentimes we need entire configuration files or scripts to be available within our containerized environments. We can provision these flexibly in our stack definitions using *docker configs*.

- Create a database initialization script db-init.sh:

```
#!/bin/bash
set -e

psql -v ON_ERROR_STOP=1 --username "$POSTGRES_USER" --dbname "$POSTGRES_DB" <<-EOSQL
CREATE TABLE PRODUCTS(PRICE FLOAT, NAME TEXT);
INSERT INTO PRODUCTS VALUES('18.95', 'widget');
INSERT INTO PRODUCTS VALUES('1.45', 'sprocket');
EOSQL
```

- On startup, the postgres container will automatically run any file *.sh found in the directory /docker-entrypoint-initdb.d. Modify your stack file to look like this:

```
version: "3.7"

services:
  db:
    image: postgres:9.6
    env_file:
      - myvars.env
    configs:
      - source: initscript
        target: /docker-entrypoint-initdb.d/init.sh

configs:
  initscript:
    file: ./db-init.sh
```

Here we see our first concrete example of composing two objects together in a stack file: our original service, and a new top-level key, configs:, which lists all the config objects we can provision to our service objects.

- Update your stack:

```
[centos@node-0 ~]$ docker stack deploy -c mystack.yaml dbdemo
```

Notice we didn't actually delete the old version of our stack first; recreating a stack with the same name and stack file will apply updates to a running stack.

4. List and inspect your config objects:

```
[centos@node-0 ~]$ docker config ls
```

ID	NAME	CREATED	UPDATED
hjrbeqqpe8l25r7u70sulung4	dbdemo_initscript	3 minutes ago	3 minutes ago

```
[centos@node-0 ~]$ docker config inspect --pretty <config ID>
```

```
ID:          hjrbeqqpe8l25r7u70sulung4
Name:        dbdemo_initscript
Labels:
  - com.docker.stack.namespace=dbdemo
Created at:  2019-01-30 15:36:31.953234447 +0000 utc
Updated at:  2019-01-30 15:36:31.953234447 +0000 utc
Data:
  #!/bin/bash
  set -e

  psql -v ON_ERROR_STOP=1 --username "$POSTGRES_USER" --dbname "$POSTGRES_DB" <<-EOSQL
    CREATE TABLE PRODUCTS(PRICE FLOAT, NAME TEXT);
    INSERT INTO PRODUCTS VALUES('18.95', 'widget');
    INSERT INTO PRODUCTS VALUES('1.45', 'sprocket');
  EOSQL
```

We can recover the plain-text contents of any config option in this manner.

5. Reconnect to your postgres database, and confirm the data got loaded correctly (remember to do this from whichever node is running your postgres container):

```
[centos@node-0 ~]$ docker container exec -it <container ID> psql -U moby -d mydb
psql (9.6.11)
Type "help" for help.

mydb=# SELECT * FROM products;
```

price	name
18.95	widget
1.45	sprocket

```
(2 rows)

mydb=# \q
```

17.4 Defining and Using Docker Secrets

In everything we've seen so far, our configurations are stored unencrypted and are recoverable directly from their definition. In some cases, this isn't good enough; when we want to store and distribute secure information like passwords or access tokens, we want this information to be encrypted by default. For this, we use *Docker secrets*.

Postgres will set the password for remote login based on the contents of the file with path specified in the `POSTGRES_PASSWORD_FILE` environment variable on startup; we'll use a secret to set this environment variable securely.

1. On node-0, place your postgres password 12345678 in a file called `mypassword`.

2. Turn the contents of mypassword into a Docker secret:

```
[centos@node-0 ~]$ docker secret create password ./mypassword
[centos@node-0 ~]$ rm mypassword
```

Note we immediately remove the plaintext mypassword - of course we don't want it sitting around in plain text for someone to find later.

3. Inspect your secret:

```
[centos@node-0 ~]$ docker secret inspect password

[
  {
    "ID": "agxqp9v4zdch2igeh59zt1qyb",
    "Version": {
      "Index": 5548
    },
    "CreatedAt": "2019-01-30T15:50:45.035216925Z",
    "UpdatedAt": "2019-01-30T15:50:45.035216925Z",
    "Spec": {
      "Name": "password",
      "Labels": {}
    }
  }
]
```

Unlike configs, Docker won't return the value of a secret at the command line once encrypted in the raft datastore. Only containers authorized to use this secret will be able to recover it in plain text.

4. By default, secrets are provisioned in containers as plaintext files at the path `/run/secrets/<secretname>`. Modify your stack file to consume your secret, and point to it with the `POSTGRES_PASSWORD_FILE` environment variable:

```
version: "3.7"

services:
  db:
    image: postgres:9.6
    env_file:
      - myvars.env
    configs:
      - source: initscript
        target: /docker-entrypoint-initdb.d/init.sh
    secrets:
      - password
    environment:
      - POSTGRES_PASSWORD_FILE=/run/secrets/password

configs:
  initscript:
    file: ./db-init.sh

secrets:
  password:
    external: true
```

Here we're adding a third top-level object, `secrets:`, to our stack; the `external: true` key indicates that we defined this object outside of our stack and are just using it here, which is a typical pattern for secrets so we can avoid having them sitting around in plain text at any time.

- Update your stack, confirm the environment variables are set correctly, and check that the password is available at `/run/secrets/password` as expected:

```
[centos@node-0 ~]$ docker stack deploy -c mystack.yaml dbdemo

[centos@node-0 ~]$ docker stack ps dbdemo
ID                NAME                IMAGE                NODE
n0c9efwva2ri      dbdemo_db.1         postgres:9.6         node-0

[centos@node-0 ~]$ docker inspect <task ID> | grep ContainerID
    "ContainerID": "c1c7cef...",

[centos@node-0 ~]$ docker container inspect <container ID> | grep POSTGRES
    "POSTGRES_DB=mydb",
    "POSTGRES_PASSWORD_FILE=/run/secrets/password",
    "POSTGRES_USER=moby",

[centos@node-0 ~]$ docker container exec <container ID> cat /run/secrets/password
1234568
```

(Remember to do the `docker container ...` commands on the node the task is actually running on). With this secret configuration, our postgres password is available in plaintext only inside the container filesystem that it has been explicitly provisioned to in our stack file.

- Clean up by removing your stack:

```
[centos@node-0 ~]$ docker stack rm dbdemo
```

17.5 Conclusion

In this exercise, we saw several different methods for defining and provisioning configurations, as well as our first example of a complete stack file for defining and composing all the elements of our application. Deciding what information to provision via configurations is an important architectural choice; in general, anything that's going to change when moving from environment to environment is a good candidate for a config, since env files, docker configs, and docker secrets are all modular and defined separately from the service definition itself; by separating configs in this way, we can just swap the config out when changing environments, without redefining our services. The (usually worse) alternative to provisioning by config is to include this information directly in the image; this is a good choice for information that is the same in all environments you plan on running that image in, but can lead to image management complexity and loss of portability if environment-specific information is hard-coded into the image. Of course, secure information like passwords should *never* be hard-coded into images; they should strictly be provisioned as Docker secrets, and consumed only from the temporary filesystem inside the container to which they are mounted.

18 Routing to Services

By the end of this exercise, you should be able to:

- Route traffic originating either inside or outside your cluster to stateless Swarm services.

18.1 Routing Cluster-Internal Traffic

- By *cluster-internal traffic*, we mean traffic from originating from a container running on your swarm, sending a request to another container running on the same swarm. Let's create a stack with two such services, attached to a custom overlay network; create a file `net-demo.yaml` with the following content:

```

version: "3.7"

services:
  destination:
    image: training/whoami:latest
    deploy:
      replicas: 3
    networks:
      - demonet

  origin:
    image: centos:7
    command: ["sleep", "100000"]
    networks:
      - demonet

networks:
  demonet:
    driver: overlay

```

Here our destination service is using the `deploy:replicas` key to ask for three replica containers based on the `training/whoami` image; this image serves a simple web page on port 8000 that reports the ID of the container its running in. Our origin service is using the `command` key to define what process to run in a container based on the `centos:7` image.

2. Deploy your stack, and find out what node your origin container is running on:

```

[centos@node-0 ~]$ docker stack deploy -c net-demo.yaml netstack

[centos@node-0 ~]$ docker stack ps netstack

```

ID	NAME	IMAGE	NODE
tqsurlytrruu	netstack_destination.1	training/whoami:latest	node-0
lp3hwp3c4nih	netstack_origin.1	centos:7	node-3
g94q730kdto5	netstack_destination.2	training/whoami:latest	node-1
n5297tisleso	netstack_destination.3	training/whoami:latest	node-2

3. Switch to the node running your origin container (node-3 in the example above), and create a bash shell inside that container:

```

[centos@node-3 ~]$ docker container ls

```

CONTAINER ID	IMAGE	COMMAND
3047dbb47a7a	centos:7	"sleep 100000"

```

[centos@node-3 ~]$ docker container exec -it <container ID> bash

[root@3047dbb47a7a /]#

```

4. From within your origin container, attempt to curl your destination container by service name several times:

```

[root@3047dbb47a7a /]# curl destination:8000
I'm a16b6e9741db
[root@3047dbb47a7a /]# curl destination:8000
I'm d3d800059d7b
[root@3047dbb47a7a /]# curl destination:8000
I'm bace80287419
[root@3047dbb47a7a /]# curl destination:8000
I'm a16b6e9741db

```

The service name defined in your stack file is DNS resolvable, and load balances traffic in round robin fashion across all the containers corresponding to that service. In this way, your application logic (simply `curl` in this example) doesn't need to do any explicit service discovery or routing to contact another service; all of that is handled by Docker's networking layer.

5. (Optional): Still inside your `centos` container, install `nslookup`, and use it to see what destination is actually resolving to:

```
[root@3047dbb47a7a /]# yum install -y bind-utils

[root@3047dbb47a7a /]# nslookup destination
Server:      127.0.0.11
Address:     127.0.0.11#53

Non-authoritative answer:
Name:   destination
Address: 10.85.5.133
```

That IP (10.85.5.133) is the *virtual IP* of your destination service. Docker automatically configures an IP virtual server on every host in your Swarm to route traffic from this VIP to the corresponding backend containers as we saw above.

6. Exit your container, and clean up by removing your stack:

```
[root@3047dbb47a7a /]# exit
[centos@node-0 ~]$ docker stack rm netstack
```

18.2 Routing Cluster-External Traffic

In the last section, we routed traffic from one Docker service to another, all within the same swarm. If we want to route ingress traffic from an external network to a service, we have to expose it on the *routing mesh*.

1. Create a new stack file `mesh.yaml`:

```
version: "3.7"

services:
  destination:
    image: training/whoami:latest
    deploy:
      replicas: 3
    ports:
      - 8080:8000
```

Here the `ports:` key is specifying that traffic arriving at port 8080 on *any node in the swarm* should be forwarded to port 8000 of our `whoami` container.

2. Deploy your stack, and `curl` the public IP of any node in your swarm on port 8080 a few times:

```
[centos@node-0 ~]$ docker stack deploy -c mesh.yaml mesh

[centos@node-0 ~]$ curl 52.55.221.133:8080
I'm 17f79606a85b
[centos@node-0 ~]$ curl 52.55.221.133:8080
I'm a8f60ac476b9
[centos@node-0 ~]$ curl 52.55.221.133:8080
I'm b285c54f5fb0
[centos@node-0 ~]$ curl 52.55.221.133:8080
I'm 17f79606a85b
```

Traffic is forwarded from the host's network namespace's port 8080 to the `whoami` containers' network namespaces in round robin fashion, similarly to above, but this time from *outside* our swarm; try visiting this IP and port in your browser to convince yourself this is externally reachable (but note that many browsers' caching behavior will hide the round robin load balancing; if you want to see the load balancing in the browser, try using a 'history-less' browser mode, closing and reopening the browser or completely purging history between each refresh).

3. Try the same `curl` again, but with the public IP for a different node in your swarm; it'll work the exact same way, since the routing mesh forwards traffic inbound to port 8080 on *any* node in the swarm on to the containers scheduled for your destination service.

Cluster internal versus cluster external routing is often misunderstood by new Swarm users; note that the port mapping 8080:8000 we created to make our service available externally on the mesh net was *not necessary* for making our service reachable by other services internally to our swarm. In general, you should not expose services on the mesh net unless they need to be reachable on the external network.

4. Clean up by deleting your mesh stack.

18.3 Conclusion

In this exercise, you saw the basic service discovery and routing that Swarm enables via DNS lookup of service names, VIP routing and port forwarding across network namespaces. In general, using this networking plane to do our service discovery helps make our inter-service communication more robust against container failure. If we were communicating with a container directly by its IP, we would have to constantly monitor whether that container was still actually reachable at that IP; Swarm effectively does that for us when we communicate via DNS resolution of service names to VIPs, since a failed container will get pulled out of the virtual IP server's round-robin routing automatically. In the case of using the mesh net to route ingress traffic from outside our cluster, mapping the external port on *every* host in the swarm onto the service's VIP means our external load balancers (which is realistically the point of ingress for external traffic to our swarm), doesn't need to know anything about where our service is scheduled; it can simply send requests to any node in the swarm, and Docker handles the rest.

19 Updating Applications

Once we have defined an application as a Docker stack, we will periodically want to update its scale, configuration, and source code. By the end of this exercise, you should be able to:

- Scale up microservice components of a stack to improve application performance
- Define and trigger a rolling update of a service
- Define and trigger an automatic rollback of a failed service update

19.1 Deploying Dockercoins

For this exercise, we'll work with a toy application called *Dockercoins*. This application is a toy 'dockercoin' miner, consisting of five microservices interacting in the following workflow:

1. A **worker** container requests a random number from a random number generator **rng**
2. After receiving a random number, the worker pushes it to a **hasher** container, which computes a hash of this number.
3. If the hash of the random number starts with 0, we accept this as a Dockercoin, and forward it to a **redis** container.
4. Meanwhile, a **webui** container monitors the rate of coins being sent to redis, and visualizes this as a chart on a web page.
5. Download the Dockercoins app from Github and change directory to `~/orchestration-workshop/dockercoins`:

```
[centos@node-0 ~]$ git clone -b ee3.0 \
  https://github.com/docker-training/orchestration-workshop.git

[centos@node-0 ~]$ cd ~/orchestration-workshop/dockercoins
```

6. The Dockercoins application is defined in `docker-compose.yml`; have a look at this file, and make sure you understand what every key is doing. Once you're satisfied with this, deploy the stack:

```
[centos@node-0 dockercoins]$ docker stack deploy -c docker-compose.yml dockercoins
```

Visit the Dockercoins web frontend at `<public IP>:8000`, for any public IP in your swarm. You should see Dockercoins getting mined at a rate of a few per second.

19.2 Scaling Up an Application

If we've written our services to be stateless, we might hope for linear performance scaling in the number of replicas of that service. For example, our `worker` service requests a random number from the `rng` service and hands it off to the `hasher` service; the faster we make those requests, the higher our throughput of Dockercoins should be, as long as there are no other confounding bottlenecks.

1. Modify the `worker` service definition in `docker-compose.yml` to set the number of replicas to create using the `deploy` and `replicas` keys:

```
worker:
  image: training/dockercoins-worker:1.0
  networks:
    - dockercoins
  deploy:
    replicas: 2
```

2. Update your app by running the same command you used to launch it in the first place:

```
[centos@node-0 dockercoins]$ docker stack deploy -c docker-compose.yml dockercoins
```

Check the web frontend; after a few seconds, you should see about double the number of hashes per second, as expected.

3. Scale up even more by changing the `worker` replicas to 10. A small improvement should be visible, but certainly not an additional factor of 5. Something else is bottlenecking Dockercoins; let's investigate the two services `worker` is interacting with: `rng` and `hasher`.
4. The `rng` and `hasher` services are exposed on host ports 8001 and 8002, so we can use `httping` to probe their latency:

```
[centos@node-0 dockercoins]$ httping -c 5 localhost:8001
[centos@node-0 dockercoins]$ httping -c 5 localhost:8002
```

`rng` is much slower to respond, suggesting that it might be the bottleneck. If this random number generator is based on an entropy collector (random voltage microfluctuations in the machine's power supply, for example), it won't be able to generate random numbers beyond a physically limited rate; we need more machines collecting more entropy in order to scale this up. This is a case where it makes sense to run exactly one copy of this service per machine, via `global` scheduling (as opposed to potentially many copies on one machine, or whatever the scheduler decides as in the default `replicated` scheduling).

5. Modify the definition of our `rng` service in `docker-compose.yml` to be globally scheduled:

```
rng:
  image: training/dockercoins-rng:1.0
  networks:
    - dockercoins
  ports:
```

```
- "8001:80"
deploy:
  mode: global
```

6. Scheduling can't be changed on the fly, so we need to stop our app and restart it:

```
[centos@node-0 dockercoins]$ docker stack rm dockercoins
[centos@node-0 dockercoins]$ docker stack deploy -c=docker-compose.yml dockercoins
```

7. Check the web frontend again; the overall factor of 10 improvement (from ~3 to ~35 hashes per second) should now be visible.

19.3 Creating Rolling Updates

Beyond scaling up an existing application, we'll periodically want to update the underlying source code of one or more of our components; Swarm provides mechanisms for rolling out updates in a controlled fashion that minimizes downtime.

1. First, let's change one of our services a bit: open `orchestration-workshop/dockercoins/worker/worker.py` in your favorite text editor, and find the following section:

```
def work_once():
    log.debug("Doing one unit of work")
    time.sleep(0.1)
```

Change the 0.1 to a 0.01. Save the file, exit the text editor.

2. Rebuild the worker image with a tag of `<Docker ID>/dockercoins-worker:1.1`, and push it to Docker Hub.
3. Change the `image:` value for the worker service in your `docker-compose.yml` file to that of the image you just pushed.
4. Start the update:

```
[centos@node-0 ~]$ docker stack deploy -c='docker-compose.yml' dockercoins
```

Use `docker stack ps dockercoins` to watch tasks get updated to our new 1.1 image one at a time.

19.4 Parallelizing Updates

1. We can also set our updates to run in batches by configuring some options associated with each service. Change the update parallelism to 2 and the delay to 5 seconds on the worker service by editing its definition in the `docker-compose.yml`:

```
worker:
  image: training/dockercoins-worker:1.0
  networks:
    - dockercoins
  deploy:
    replicas: 10
    update_config:
      parallelism: 2
      delay: 5s
```

2. Roll back the worker service to 1.0:

```
[centos@node-0 ~]$ docker stack deploy -c=docker-compose.yml dockercoins
```

3. On node-1, watch your updates:

```
[centos@node-1 ~]$ watch -n1 "docker service ps dockercoins_worker \
| grep -v Shutdown.*Shutdown"
```

You should see two tasks get shutdown and restarted with the 1.0 image every five seconds.

19.5 Auto-Rollback Failed Updates

In the event of an application or container failure on deployment, we'd like to automatically roll the update back to the previous version.

1. Update the worker service with some parameters to define rollback:

```
worker:
  image: training/dockercoins-worker:1.0
  networks:
  - dockercoins
  deploy:
    replicas: 10
    update_config:
      parallelism: 2
      delay: 5s
      failure_action: rollback
      max_failure_ratio: 0.2
      monitor: 20s
```

These parameters will trigger a rollback if more than 20% of services tasks fail in the first 20 seconds after an update.

2. Update your stack to make sure the rollback parameters are in place before you attempt to update your image:

```
[centos@node-0 ~]$ docker stack deploy -c=docker-compose.yml dockercoins
```

3. Make a broken version of the worker service to trigger a rollback with; try removing all the import commands at the top of worker.py, for example. Then rebuild the worker image with a tag <Docker ID>/dockercoins-worker:bugged, push it to Docker Hub, and attempt to update your service:

```
[centos@node-0 worker]$ docker image build -t <Docker ID>/dockercoins-worker:bugged .
[centos@node-0 worker]$ docker image push <Docker ID>/dockercoins-worker:bugged
```

4. Update your docker-compose.yml file to use the :bugged tag for the worker service, and redeploy your stack as above.
5. The connection to node-1 running watch should show the :bugged tag getting deployed, failing, and rolling back to :1.0 automatically over the next minute or two:

NAME	IMAGE	CURRENT STATE
dockercoins_worker.1	dockercoins-worker:1.0	Running 2 minutes ago
dockercoins_worker.2	dockercoins-worker:1.0	Running 2 minutes ago
dockercoins_worker.3	dockercoins-worker:1.0	Running about a minute ago
_ dockercoins_worker.3	dockercoins-worker:bugged	Failed about a minute ago
_ dockercoins_worker.3	dockercoins-worker:bugged	Failed about a minute ago
_ dockercoins_worker.3	dockercoins-worker:bugged	Failed about a minute ago
dockercoins_worker.4	dockercoins-worker:1.0	Running about a minute ago
_ dockercoins_worker.4	dockercoins-worker:bugged	Failed about a minute ago
dockercoins_worker.5	dockercoins-worker:1.0	Running about a minute ago
_ dockercoins_worker.5	dockercoins-worker:bugged	Failed about a minute ago
_ dockercoins_worker.5	dockercoins-worker:bugged	Failed about a minute ago
_ dockercoins_worker.5	dockercoins-worker:bugged	Failed about a minute ago
dockercoins_worker.6	dockercoins-worker:1.0	Running 2 minutes ago
dockercoins_worker.7	dockercoins-worker:1.0	Running 2 minutes ago

dockercoins_worker.8	dockercoins-worker:1.0	Running 2 minutes ago
dockercoins_worker.9	dockercoins-worker:1.0	Running about a minute ago
_ dockercoins_worker.9	dockercoins-worker:bugged	Failed about a minute ago
dockercoins_worker.10	dockercoins-worker:1.0	Running 2 minutes ago

For example, this table indicates that tasks 3, 4, 5 and 9 all attempted to update to the :bugged tag, failed, and successfully rolled back to the :1.0 tag (the _ symbol is meant to indicate failed ancestors for an individual task; so dockercoins_worker.3 above made three failed attempts to run the :bugged image before rolling back).

Use CTRL+C to exit this watch view when done.

6. Clean up by removing your stack:

```
[centos@node-0 dockercoins]$ docker stack rm dockercoins
```

19.6 Optional Challenge: Improving Dockercoins

Dockercoins' stack file is very rudimentary, and not at all suitable for production. If you have time, try modifying Dockercoins' stack file with some of the best practices you've learned in this workshop; think about things like security, operational stability, latency and scheduling. This activity is best done in groups! Partner up with someone else and discuss what improvements you can make, then try them out and make sure they work as you expected.

19.7 Conclusion

In this exercise, we explored deploying and redeploying an application as stacks and services. Note that relaunching a running stack updates all the objects it manages in the most non-disruptive way possible; there is usually no need to remove a stack before updating it. In production, rollback contingencies should always be used to cautiously upgrade images, cutting off potential damage before an entire service is taken down.

20 Installing Kubernetes

By the end of this exercise, you should be able to:

- Set up a Kubernetes cluster with one master and two nodes

20.1 Initializing Kubernetes

1. On node-0, initialize the cluster with kubeadm:

```
[centos@node-0 ~]$ sudo kubeadm init --pod-network-cidr=192.168.0.0/16 \
--ignore-preflight-errors=SystemVerification
```

If successful, the output will end with a join command:

...

You can now join any number of machines by running the following on each node as root:

```
kubeadm join 10.10.29.54:6443 --token xxx --discovery-token-ca-cert-hash sha256:yyy
```

2. To start using your cluster, you need to run:

```
[centos@node-0 ~]$ mkdir -p $HOME/.kube
[centos@node-0 ~]$ sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
[centos@node-0 ~]$ sudo chown $(id -u):$(id -g) $HOME/.kube/config
```

3. List all your nodes in the cluster:

```
[centos@node-0 ~]$ kubectl get nodes
```

Which should output something like:

NAME	STATUS	ROLES	AGE	VERSION
node-0	NotReady	master	2h	v1.16.3

The NotReady status indicates that we must install a network for our cluster.

4. Let's install the Calico network driver:

```
[centos@node-0 ~]$ kubectl apply -f https://bit.ly/2KHCEg0
```

5. After a moment, if we list nodes again, ours should be ready:

```
[centos@node-0 ~]$ kubectl get nodes -w
```

NAME	STATUS	ROLES	AGE	VERSION
node-0	NotReady	master	1m	v1.16.3
node-0	NotReady	master	1m	v1.16.3
node-0	NotReady	master	1m	v1.16.3
node-0	Ready	master	2m	v1.16.3
node-0	Ready	master	2m	v1.16.3

6. Execute the join command you found above when initializing Kubernetes on node-1 and node-2 (you'll need to add sudo to the start, and --ignore-preflight-errors=SystemVerification to the end), and then check the status back on node-0:

```
[centos@node-1 ~]$ sudo kubeadm join ... --ignore-preflight-errors=SystemVerification
[centos@node-2 ~]$ sudo kubeadm join ... --ignore-preflight-errors=SystemVerification
[centos@node-0 ~]$ kubectl get nodes
```

After a few moments, there should be three nodes listed - all with the Ready status.

7. Let's see what system pods are running on our cluster:

```
[centos@node-0 ~]$ kubectl get pods -n kube-system
```

which results in something similar to this:

NAME	READY	STATUS	RESTARTS	AGE
calico-kube-controllers-6b64bcd855-lnrw7	1/1	Running	0	2m13s
calico-node-btgn6	1/1	Running	0	54s
calico-node-rdhrb	0/1	Running	0	20s
calico-node-s7dnc	1/1	Running	0	2m13s
coredns-5644d7b6d9-c5j5c	1/1	Running	0	7m29s
coredns-5644d7b6d9-gm4rl	1/1	Running	0	7m29s
etcd-node-0	1/1	Running	0	6m39s
kube-apiserver-node-0	1/1	Running	0	6m41s
kube-controller-manager-node-0	1/1	Running	0	6m26s
kube-proxy-jc6h7	1/1	Running	0	20s
kube-proxy-nczr9	1/1	Running	0	7m29s
kube-proxy-sf9n6	1/1	Running	0	54s
kube-scheduler-node-0	1/1	Running	0	6m51s

We can see the pods running on the master: etcd, api-server, controller manager and scheduler, as well as calico and DNS infrastructure pods deployed when we installed calico.

20.2 Conclusion

At this point, we have a Kubernetes cluster with one master and two workers ready to accept workloads.

21 Kubernetes Orchestration

By the end of this exercise, you should be able to:

- Define and launch basic pods, replicaSets and deployments using `kubectl`
- Get metadata, configuration and state information about a kubernetes object using `kubectl describe`
- Update an image for a pod in a running kubernetes deployment

21.1 Creating Pods

1. On your master node, create a yaml file `pod.yaml` to describe a simple pod with the following content:

```
apiVersion: v1
kind: Pod
metadata:
  name: demo
spec:
  containers:
  - name: nginx
    image: nginx:1.7.9
```

2. Deploy your pod:

```
[centos@node-0 ~]$ kubectl create -f pod.yaml
```

3. Confirm your pod is running:

```
[centos@node-0 ~]$ kubectl get pod demo
```

4. Get some metadata about your pod:

```
[centos@node-0 ~]$ kubectl describe pod demo
```

5. Delete your pod:

```
[centos@node-0 ~]$ kubectl delete pod demo
```

6. Modify `pod.yaml` to create a second container inside your pod:

```
apiVersion: v1
kind: Pod
metadata:
  name: demo
spec:
  containers:
  - name: nginx
    image: nginx:1.7.9
  - name: sidecar
    image: centos:7
    command: ["ping"]
    args: ["8.8.8.8"]
```

7. Deploy this new pod, and create a bash shell inside the container named `sidecar`:

```
[centos@node-0 ~]$ kubectl create -f pod.yaml
[centos@node-0 ~]$ kubectl exec -c=sidecar -it demo -- /bin/bash
```

8. From within the `sidecar` container, fetch the `nginx` landing page on the default port 80 using `localhost`:

```
[root@demo /]# curl localhost:80
```

You should see the html of the `nginx` landing page. Note **these containers can reach each other on localhost**, meaning they are sharing a network namespace. Now list the processes in your `sidecar` container:

```
[root@demo /]# ps -aux
```

You should see the ping process we containerized, the shell we created to explore this container using `kubectl exec`, and the `ps` process itself - but no `nginx`. While a network namespace is shared between the containers, they still have their own PID namespace (for example).

9. Finally, remember to exit out of this pod, and delete it:

```
[root@demo /]# exit
[centos@node-0 ~]$ kubectl delete pod demo
```

21.2 Creating ReplicaSets

1. On your master node-0, create a yaml file `replicaset.yaml` to describe a simple replicaSet with the following content:

```
apiVersion: apps/v1
kind: ReplicaSet
metadata:
  name: rs-demo
spec:
  replicas: 3
  selector:
    matchLabels:
      component: reverse-proxy
  template:
    metadata:
      labels:
        component: reverse-proxy
    spec:
      containers:
      - name: nginx
        image: nginx:1.7.9
```

Notice especially the `replicas` key, which defines how many copies of this pod to create, and the `template` section; this defines the pod to replicate, and is described almost exactly like the first pod definition we created above. The difference here is the required presence of the `labels` key in the pod's metadata, which must match the selector -> `matchLabels` item in the specification of the replicaSet.

2. Deploy your replicaSet, and get some state information about it:

```
[centos@node-0 ~]$ kubectl create -f replicaset.yaml
[centos@node-0 ~]$ kubectl describe replicaset rs-demo
```

After a few moments, you should see something like

```
Name:          rs-demo
Namespace:     default
Selector:      component=reverse-proxy
Labels:        component=reverse-proxy
Annotations:   <none>
Replicas:      3 current / 3 desired
Pods Status:   3 Running / 0 Waiting / 0 Succeeded / 0 Failed
Pod Template:
  Labels:  component=reverse-proxy
  Containers:
    nginx:
      Image:   nginx:1.7.9
      Port:    <none>
```

```

Host Port:    <none>
Environment:  <none>
Mounts:       <none>
Volumes:      <none>
Events:
  Type        Reason              Age   From                      Message
  ----        -
  Normal      SuccessfulCreate     35s   replicaset-controller     Created pod: rs-demo-jxmjj
  Normal      SuccessfulCreate     35s   replicaset-controller     Created pod: rs-demo-dmdtf
  Normal      SuccessfulCreate     35s   replicaset-controller     Created pod: rs-demo-j62fx

```

Note the replicaSet has created three pods as requested, and will reschedule them if they exit.

3. Try killing off one of your pods, and reexamining the output of the above describe command. The <pod name> comes from the last three lines in the output above, such as rs-demo-jxmjj:

```

[centos@node-0 ~]$ kubectl delete pod <pod name>
[centos@node-0 ~]$ kubectl describe replicaset rs-demo

```

The dead pod gets rescheduled by the replicaSet, similar to a failed task in Docker Swarm.

4. Delete your replicaSet:

```

[centos@node-0 ~]$ kubectl delete replicaset rs-demo

```

21.3 Creating Deployments

1. On your master node-0, create a yaml file deployment.yaml to describe a simple deployment with the following content:

```

apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
spec:
  replicas: 3
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
      - name: nginx
        image: nginx:1.7.9

```

Notice this is the exact same structure as your replicaSet yaml above, but this time the kind is Deployment. Deployments create a replicaSet of pods, but add some deployment management functionality on top of them, such as rolling updates and rollback.

2. Spin up your deployment, and get some state information:

```

[centos@node-0 ~]$ kubectl create -f deployment.yaml
[centos@node-0 ~]$ kubectl describe deployment nginx-deployment

```

The describe command should return something like:

```

Name:          nginx-deployment
Namespace:     default
CreationTimestamp: Thu, 24 May 2018 04:29:18 +0000
Labels:       <none>
Annotations:  deployment.kubernetes.io/revision=1
Selector:     app=nginx
Replicas:     3 desired | 3 updated | 3 total | 3 available | 0 unavailable
StrategyType: RollingUpdate
MinReadySeconds: 0
RollingUpdateStrategy: 25% max unavailable, 25% max surge
Pod Template:
  Labels:  app=nginx
  Containers:
    nginx:
      Image:        nginx:1.7.9
      Port:         <none>
      Host Port:    <none>
      Environment:  <none>
      Mounts:       <none>
      Volumes:      <none>
  Conditions:
    Type           Status  Reason
    ----           -
    Available       True    MinimumReplicasAvailable
    Progressing     True    NewReplicaSetAvailable
OldReplicaSets:  <none>
NewReplicaSet:   nginx-deployment-85f7784776 (3/3 replicas created)
Events:
  Type           Reason             Age   From                      Message
  ----           -
  Normal        ScalingReplicaSet   10s   deployment-controller     Scaled up replica set
                                                                    nginx-deployment-xxx to 3

```

Note the very last line, indicating this deployment actually created a replicaSet which it used to scale up to three pods.

3. List your replicaSets and pods:

```

[centos@node-0 ~]$ kubectl get replicaSet

NAME                                DESIRED   CURRENT   READY   AGE
nginx-deployment-79d4c74645         3         3         3       17s

[centos@node-0 ~]$ kubectl get pod

NAME                                READY   STATUS    RESTARTS   AGE
nginx-deployment-79d4c74645-9mtzm   1/1     Running   0          22s
nginx-deployment-79d4c74645-k7wml   1/1     Running   0          22s
nginx-deployment-79d4c74645-rrfrf   1/1     Running   0          22s

```

You should see one replicaSet and three pods created by your deployment, similar to the above.

4. Upgrade the nginx image from 1.7.9 to 1.9.1:

```

[centos@node-0 ~]$ kubectl set image deployment/nginx-deployment nginx=nginx:1.9.1

```

5. After a few seconds, kubectl describe your deployment as above again. You should see that the image has been updated, and that the old replicaSet has been scaled down to 0 replicas, while a new replicaSet (with your updated image) has been scaled up to 3 pods. List your replicaSets one more time:

```
[centos@node-0 ~]$ kubectl get replicaSets
```

You should see something like

NAME	DESIRED	CURRENT	READY	AGE
nginx-deployment-69df9ccbf8	3	3	3	4m
nginx-deployment-85f7784776	0	0	0	9m

Do a `kubectl describe replicaSet <replicaSet scaled down to 0>`; you should see that while no pods are running for this replicaSet, the old replicaSet's definition is still around so we can easily roll back to this version of the app if we need to.

6. Clean up your cluster:

```
[centos@node-0 ~]$ kubectl delete deployment nginx-deployment
```

21.4 Conclusion

In this exercise, you explored the basic scheduling objects of pods, replicaSets, and deployments. Each object is responsible for a different part of the orchestration stack; pods are the basic unit of scheduling, replicaSets do keep-alive and scaling, and deployments provide update and rollback functionality. In a sense, these objects all 'nest' one inside the next; by creating a deployment, you implicitly created a replicaSet which in turn created the corresponding pods. In most cases, you're better off creating deployments rather than replicaSets or pods directly; this way, you get all the orchestrating scheduling features you would expect in analogy to a Docker Swarm service.

22 Provisioning Kube Configuration

By the end of this exercise, you should be able to:

- write kube yaml describing secrets and config maps, and associate them with pods and deployments.

22.1 Provisioning ConfigMaps

1. Create a file on node-0 called env-config with the following content:

```
user=moby
db=mydb
```

2. Create a configMap that parses each line in this file as a separate key/value pair:

```
[centos@node-0 ~]$ kubectl create configmap dbconfig --from-env-file=env-config
```

3. Ask kubectl to repeat this configMap back to us, in yaml:

```
[centos@node-0 ~]$ kubectl get configmap dbconfig -o yaml

apiVersion: v1
data:
  db: mydb
  user: moby
kind: ConfigMap
metadata:
  creationTimestamp: "2019-01-08T15:41:31Z"
  name: dbconfig
  namespace: default
  resourceVersion: "709821"
  selfLink: /api/v1/namespaces/default/configmaps/dbconfig
  uid: df14e7a7-135b-11e9-87ee-0242ac11000a
```

The key/value pairs parsed from `env-config` are visible under the `data` key in this file; we could have created the same configMap declaratively via `kubectl create -f <yaml filename>` like we've been doing so far, using this yaml.

Imperative vs. Declarative kubectl commands: `kubectl` supports two syntaxes for most actions: *imperative*, which specifies the action to take (like `kubectl create pod ...`, `kubectl describe deployment ...` etc), and *declarative*, which specifies objects in a yaml file and creates them with `kubectl create -f <yaml filename>`. Which to use is a matter of preference; in general, I recommend *declarative* (ie file based) commands for *any action that changes the state of the system* (ie create / update / destroy operations) so that these changes can be based off of easily tracked and version controlled config files, and *imperative* commands only for gathering information without changing anything (`kubectl get ...`, `kubectl describe ...`). These are not strict rules (we just saw a convenient example of the imperative `kubectl create configmap` above, useful since an `env-file` specification is so much easier to write than the corresponding yaml), but lend themselves well to good record-keeping, automation and version control for your workloads.

4. Describe a postgres database in a pod with the following `postgres.yaml`:

```
apiVersion: v1
kind: Pod
metadata:
  name: dbdemo
  namespace: default
spec:
  containers:
  - name: pg
    image: postgres:9.6
    env:
      - name: POSTGRES_USER
        valueFrom:
          configMapKeyRef:
            name: dbconfig
            key: user
      - name: POSTGRES_DB
        valueFrom:
          configMapKeyRef:
            name: dbconfig
            key: db
```

Here we're populating the environment variables `POSTGRES_USER` and `POSTGRES_DB` from our configMap, under the `containers:env` specification. Notice that the pod definition itself makes no reference to the literal values of these environment variables; we can reconfigure our database (say for deployment in a different environment) by swapping out our `dbconfig` configMap, and leaving our pod definition untouched.

Deploy your pod as usual with `kubectl create -f postgres.yaml`.

5. Describe your `dbdemo` pod:

```
[centos@node-0 ~]$ kubectl describe pod dbdemo
...
Environment:
  POSTGRES_USER: <set to the key 'user' of config map 'dbconfig'> Optional: false
  POSTGRES_DB:   <set to the key 'db' of config map 'dbconfig'>   Optional: false
...
```

You should see a block like the above, indicating that the listed environment variables are being populated from the configMap, as expected.

6. Attach to your postgres database using the username and database name you specified in `config.yaml`, to prove to yourself the `configMap` information actually got consumed:

```
[centos@node-0 ~]$ kubectl exec -it -c pg dbdemo -- psql -U moby -d mydb
```

```
mydb=# \du
```

Role name	List of roles Attributes	Member of
moby	Superuser, Create role, Create DB, Replication, Bypass RLS	{ }

```
mydb=# \q
```

The user and database got created with the names we defined.

7. Delete your pod with `kubectl delete -f postgres.yaml`.
8. So far, we've used the environment variables `postgres` looks for for setting user and database names. Some config, however, is expected as a file rather than an environment variable; `configMaps` can provision config files as well as environment config. Create a database initialization script `db-init.sh`:

```
#!/bin/bash
set -e

psql -v ON_ERROR_STOP=1 --username "$POSTGRES_USER" --dbname "$POSTGRES_DB" <<-EOSQL
CREATE TABLE PRODUCTS(PRICE FLOAT, NAME TEXT);
INSERT INTO PRODUCTS VALUES('18.95', 'widget');
INSERT INTO PRODUCTS VALUES('1.45', 'sprocket');
EOSQL
```

9. Turn this entire file into a `configMap`:

```
[centos@node-0 ~]$ kubectl create configmap dbinit --from-file=db-init.sh
```

10. We'll provision this config file to our postgres container by mounting it in as a volume to the correct path; postgres will run all `.sh` files found at the path `/docker-entrypoint-initdb.d` upon initialization. Change your `postgres.yaml` file to look like this:

```
apiVersion: v1
kind: Pod
metadata:
  name: dbdemo
  namespace: default
spec:
  containers:
  - name: pg
    image: postgres:9.6
    volumeMounts:
    - name: dbinit-vol
      mountPath: /docker-entrypoint-initdb.d
  env:
  - name: POSTGRES_USER
    valueFrom:
      configMapKeyRef:
        name: dbconfig
        key: user
  - name: POSTGRES_DB
    valueFrom:
      configMapKeyRef:
        name: dbconfig
```

```

        key: db
volumes:
  - name: dbinit-vol
    configMap:
      name: dbinit

```

Here we've added the `volumeMounts` key describing which volume (`dbinit-vol`) to mount in the container and at what path (`/docker-entrypoint-initdb.d`). We've also added the `volumes` key to define the volumes themselves; we create one volume named `dbinit-vol`, populated from the files contained in the configMap `dbinit` we just created.

11. Deploy postgres with this configuration, and check that the database initialization script actually worked:

```

[centos@node-0 ~]$ kubectl create -f postgres.yaml
[centos@node-0 ~]$ kubectl exec -it -c pg dbdemo -- psql -U moby -d mydb

mydb=# SELECT * FROM products;

 price | name
-----+-----
 18.95 | widget
   1.45 | sprocket
(2 rows)

mydb=# \q

```

Our table was initialized via our config file, as expected. After exiting your pod, delete it with `kubectl delete -f postgres.yaml`.

22.2 Provisioning Secrets

So far, we've provisioned non-sensitive data to our pod, but often we want options for added security when provisioning things like passwords or other access tokens. For this, Kubernetes maintains a separate config provisioning tool, *secrets*.

1. Let's set a custom password for our database using a Kubernetes secret. Create a file `secret.yaml` with the following content:

```

apiVersion: v1
kind: Secret
metadata:
  name: postgres-pwd
  namespace: default
type: Opaque
stringData:
  password: "mypassword"

```

Create the secret via `kubectl create -f secret.yaml`. This will create a secret called `postgres-pwd` that encodes our password.

Note: Of course it's not recommended to leave your secret unencrypted in a file like `secret.yaml`; in practice, we'd delete this file as soon as the secret is created.

2. Update your `postgres.yaml` definition to look like this:

```

apiVersion: v1
kind: Pod
metadata:
  name: dbdemo
  namespace: default
spec:

```

```
containers:
- name: pg
  image: postgres:9.6
  volumeMounts:
  - name: dbinit-vol
    mountPath: /docker-entrypoint-initdb.d
  env:
  - name: POSTGRES_USER
    valueFrom:
      configMapKeyRef:
        name: dbconfig
        key: user
  - name: POSTGRES_DB
    valueFrom:
      configMapKeyRef:
        name: dbconfig
        key: db
  - name: POSTGRES_PASSWORD
    valueFrom:
      secretKeyRef:
        name: postgres-pwd
        key: password
volumes:
- name: dbinit-vol
  configMap:
    name: dbinit
```

This is exactly the same as above, but adds the block at the bottom which populates the POSTGRES_PASSWORD environment variable in the pg container with the value found under the password key in the postgres-pwd secret.

3. Create your pod, and dump its postgres environment variables:

```
[centos@node-0 ~]$ kubectl create -f postgres.yaml
[centos@node-0 ~]$ kubectl exec -it -c pg dbdemo -- env | grep POSTGRES

POSTGRES_USER=moby
POSTGRES_DB=mydb
POSTGRES_PASSWORD=mypassword
```

The POSTGRES_PASSWORD has been provisioned from your Kube secret.

4. Note that anyone with kubectl get secret permissions can recover your secret as follows:

```
[centos@node-0 ~]$ kubectl get secret postgres-pwd -o yaml

apiVersion: v1
data:
  password: bXlwYXNzd29yZA==
kind: Secret
metadata:
  creationTimestamp: "2019-01-07T18:42:01Z"
  name: postgres-pwd
  namespace: default
  resourceVersion: "604412"
  selfLink: /api/v1/namespaces/default/secrets/postgres-pwd
  uid: ebb6f645-12ab-11e9-87ee-0242ac11000a
type: Opaque
```

```
[centos@node-0 ~]$ echo 'bXlwYXNzd29yZA==' | base64 --decode
mypassword
```

Challenge: the secret password can also be recovered in plain text on the node hosting the postgres pod. Can you find it?

5. As usual, clean up by deleting the pods, configMaps, and secret you created in this exercise.

22.3 Conclusion

In this exercise, we looked at configMaps and secrets, two tools for provisioning information to your deployments. When deciding where to place configuration, it can help to prioritize designing for reusability; in the postgres example we saw, we separated out all the environment config from the pod definition, so the same pod yaml could be migrated from environment to environment with no changes; all the environment specific data was captured in the configMaps and secret. If you find yourself having to do heavy reconfiguration of your pods and deployments (or even images) as you migrate from one environment to another, consider if it would be possible to separate this configuration from your pod definition using a configMap or secret.

23 Kubernetes Networking

By the end of this exercise, you should be able to:

- Predict what routing tables rules calico will write to each host in your cluster
- Route and load balance traffic to deployments using clusterIP and nodePort services
- Reconfigure a deployment into a daemonSet (analogous to changing scheduling from 'replicated' to 'global' in a swarm service)

23.1 Routing Traffic with Calico

1. Make sure you're on the master node node-0, and redeploy the nginx deployment defined in `deployment.yaml` from the last exercise.
2. List your pods:

```
[centos@node-0 ~]$ kubectl get pods
```

3. Get some metadata on one of the pods found in the last step:

```
[centos@node-0 ~]$ kubectl describe pods <pod name>
```

which in my case results in:

```
Name:                nginx-deployment-69df458bc5-bb87w
Namespace:           default
Priority:             0
PriorityClassName:    <none>
Node:                node-2/10.10.43.25
Start Time:          Thu, 09 Aug 2018 17:29:52 +0000
Labels:              app=nginx
                    pod-template-hash=2589014671
Annotations:         <none>
Status:              Running
IP:                  192.168.247.10
Controlled By:       ReplicaSet/nginx-deployment-69df458bc5
Containers:
  nginx:
```

```

Container ID:   docker://26e8eac...
Image:         nginx:1.7.9
Image ID:      docker-pullable://nginx@sha256:e3456c...
Port:         <none>
Host Port:    <none>
State:        Running
  Started:    Thu, 09 Aug 2018 17:29:53 +0000
Ready:        True
Restart Count: 0
Environment:   <none>
Mounts:
  /var/run/secrets/kubernetes.io/serviceaccount from default-token-fkf5d (ro)
Conditions:
  Type              Status
  Initialized       True
  Ready             True
  ContainersReady   True
  PodScheduled      True
Volumes:
  default-token-fkf5d:
    Type:          Secret (a volume populated by a Secret)
    SecretName:    default-token-fkf5d
    Optional:      false
QoS Class:       BestEffort
Node-Selectors:  <none>
Tolerations:     node.kubernetes.io/not-ready:NoExecute for 300s
                  node.kubernetes.io/unreachable:NoExecute for 300s
Events:
  Type    Reason      Age   From          Message
  ----    -
  Normal  Scheduled   20s   default-scheduler  Successfully assigned
                        default/nginx-deployment-xxx
                        to node-2
  Normal  Pulled      19s   kubelet, node-2   Container image "nginx:1.7.9"
                        already present on machine
  Normal  Created     19s   kubelet, node-2   Created container
  Normal  Started     19s   kubelet, node-2   Started container
[centos@node-0 ~]$ kubectl describe pods nginx-deployment-69df458bc5-bb87w
Name:          nginx-deployment-69df458bc5-bb87w
Namespace:     default
Priority:       0
PriorityClassName: <none>
Node:          node-2/10.10.43.25
Start Time:    Thu, 09 Aug 2018 17:29:52 +0000
Labels:        app=nginx
                pod-template-hash=2589014671
Annotations:   <none>
Status:        Running
IP:            192.168.247.10
Controlled By: ReplicaSet/nginx-deployment-69df458bc5
Containers:
  nginx:
    Container ID:  docker://26e8eac...
    Image:         nginx:1.7.9
    Image ID:      docker-pullable://nginx@sha256:e3456c8...

```

```

Port:          <none>
Host Port:     <none>
State:        Running
  Started:    Thu, 09 Aug 2018 17:29:53 +0000
Ready:        True
Restart Count: 0
Environment:  <none>
Mounts:
  /var/run/secrets/kubernetes.io/serviceaccount from default-token-fkf5d (ro)
Conditions:
  Type              Status
  Initialized       True
  Ready             True
  ContainersReady   True
  PodScheduled      True
Volumes:
  default-token-fkf5d:
    Type:          Secret (a volume populated by a Secret)
    SecretName:    default-token-fkf5d
    Optional:      false
QoS Class:       BestEffort
Node-Selectors:  <none>
Tolerations:     node.kubernetes.io/not-ready:NoExecute for 300s
                  node.kubernetes.io/unreachable:NoExecute for 300s
Events:
  Type    Reason      Age   From          Message
  ----    -
  Normal  Scheduled   1m    default-scheduler  Successfully assigned
                        default/nginx-deployment-xxx
                        to node-2
  Normal  Pulled      1m    kubelet, node-2    Container image "nginx:1.7.9"
                        already present on machine
  Normal  Created     1m    kubelet, node-2    Created container
  Normal  Started     1m    kubelet, node-2    Started container

```

We can see that in our case the pod has been deployed to node-2 as indicated near the top of the output, and the pod has an IP of 192.168.247.10.

4. Have a look at the routing table on node-0 using `ip route`, which for my example looks like:

```

[centos@node-0 ~]$ ip route

default via 10.10.0.1 dev eth0
10.10.0.0/20 dev eth0 proto kernel scope link src 10.10.7.20
172.17.0.0/16 dev docker0 proto kernel scope link src 172.17.0.1
blackhole 192.168.39.192/26 proto bird
192.168.39.193 dev cali12b0eb5c038 scope link
192.168.39.194 dev calibe752d56965 scope link
192.168.84.128/26 via 10.10.24.89 dev tunl0 proto bird onlink
192.168.247.0/26 via 10.10.43.25 dev tunl0 proto bird onlink

```

Notice the last line; this rule was written by Calico to send any traffic on the 192.168.247.0/26 subnet (which the pod we examined above is on) to the host at IP 10.10.43.25 via IP in IP as indicated by the `dev tunl0` entry. Look at your own routing table and list of VM IPs; what are the corresponding subnets, pod IPs and host IPs in your case? Does that make sense based on the host you found for the nginx pod above?

5. Curl your pod's IP on port 80 from node-0; you should see the HTML for the nginx landing page. By default this pod is reachable at this IP from anywhere in the Kubernetes cluster.

- Head over to the node this pod got scheduled on (node-2 in the example above), and have a look at that host's routing table in the same way:

```
[centos@node-2 ~]$ ip route

default via 10.10.32.1 dev eth0
10.10.32.0/20 dev eth0 proto kernel scope link src 10.10.43.25
172.17.0.0/16 dev docker0 proto kernel scope link src 172.17.0.1
192.168.39.192/26 via 10.10.7.20 dev tunl0 proto bird onlink
192.168.84.128/26 via 10.10.24.89 dev tunl0 proto bird onlink
blackhole 192.168.247.0/26 proto bird
192.168.247.10 dev calia5daa4e7a1d scope link
192.168.247.11 dev cali9ff153fb143 scope link
```

Again notice the second-to-last line; this time, the pod IP is routed to a cali*** device, which is a virtual ethernet endpoint in the host's network namespace, providing a point of ingress into that pod. Once again try `curl <pod IP>:80` - you'll see the nginx landing page html as before.

- Back on node-0, fetch the logs generated by the pod you've been curling:

```
[centos@node-0 ~]$ kubectl logs <pod name>

10.10.52.135 - - [09/May/2018:13:58:42 +0000]
"GET / HTTP/1.1" 200 612 "-" "curl/7.29.0" "-"
192.168.84.128 - - [09/May/2018:14:00:41 +0000]
"GET / HTTP/1.1" 200 612 "-" "curl/7.29.0" "-"
```

We see records of the curls we performed above; like Docker containers, these logs are the STDOUT and STDERR of the containerized processes.

23.2 Routing and Load Balancing with Services

- Above we were able to hit nginx at the pod IP, but there is no guarantee this pod won't get rescheduled to a new IP. If we want a stable IP for this deployment, we need to create a ClusterIP service. In a file `cluster.yaml` on your master node-0:

```
apiVersion: v1
kind: Service
metadata:
  name: cluster-demo
spec:
  selector:
    app: nginx
  ports:
  - port: 8080
    targetPort: 80
```

Create this service with `kubectl create -f cluster.yaml`. This maps the pod internal port 80 to the cluster wide external port 8080; furthermore, this IP and port will only be reachable from *within* the cluster. Also note the `selector: app: nginx` specification; that indicates that this service will route traffic to every pod that has `nginx` as the value of the `app` label in this namespace.

- Let's see what services we have now:

```
[centos@node-0 ~]$ kubectl get services
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
kubernetes	ClusterIP	10.96.0.1	<none>	443/TCP	33m
cluster-demo	ClusterIP	10.104.201.93	<none>	8080/TCP	48s

The second one is the one we just created and we can see that a stable IP address and port 10.104.201.93:8080 has been assigned to our nginx service.

- Let's try to access Nginx now, from any node in our cluster:

```
[centos@node-0 ~]$ curl <nginx CLUSTER-IP>:8080
```

which should return the nginx welcome page. Even if pods get rescheduled to new IPs, this clusterIP service will preserve a stable endpoint for traffic to be load balanced across all pods matching the service's label selector.

- ClusterIP services are reachable only from within the Kubernetes cluster. If you want to route traffic to your pods from an external network, you'll need a NodePort service. On your master node-0, create a file `nodeport.yaml`:

```
apiVersion: v1
kind: Service
metadata:
  name: nodeport-demo
spec:
  type: NodePort
  selector:
    app: nginx
  ports:
    - port: 8080
      targetPort: 80
```

And create this service with `kubectl create -f nodeport.yaml`. Notice this is exactly the same as the ClusterIP service definition, but now we're requesting a type NodePort.

- Inspect this service's metadata:

```
[centos@node-0 ~]$ kubectl describe service nodeport-demo
```

Notice the NodePort field: this is a randomly selected port from the range 30000-32767 where your pods will be reachable externally. Try visiting your nginx deployment at any public IP of your cluster, and the port you found above, and confirming you can see the nginx landing page.

- Clean up the objects you created in this section:

```
[centos@node-0 ~]$ kubectl delete deployment nginx-deployment
[centos@node-0 ~]$ kubectl delete service cluster-demo
[centos@node-0 ~]$ kubectl delete service nodeport-demo
```

23.3 Optional: Deploying DockerCoins onto the Kubernetes Cluster

- First deploy Redis via `kubectl create deployment`:

```
[centos@node-0 ~]$ kubectl create deployment redis --image=redis
```

- And now all the other deployments. To avoid too much typing we do that in a loop:

```
[centos@node-0 ~]$ for DEPLOYMENT in hasher rng webui worker; do
  kubectl create deployment $DEPLOYMENT \
    --image=training/dockercoins-${DEPLOYMENT}:1.0
done
```

- Let's see what we have:

```
[centos@node-0 ~]$ kubectl get pods -o wide -w
```

in my case the result is:

hasher-6c64f78655-rgjk5	1/1	Running	0	53s	10.36.0.1	node-2
redis-75586d7d7c-mmjg7	1/1	Running	0	5m	10.44.0.2	node-1
rng-d94d56d4f-twlwz	1/1	Running	0	53s	10.44.0.1	node-1
webui-6d8668984d-sqtt8	1/1	Running	0	52s	10.36.0.2	node-2
worker-56756ddb8-1bv9r	1/1	Running	0	52s	10.44.0.3	node-1

pods have been distributed across our cluster.

- We can also look at some logs:

```
[centos@node-0 ~]$ kubectl logs deploy/rng
[centos@node-0 ~]$ kubectl logs deploy/worker
```

The rng service (and also the hasher and webui services) seem to work fine but the worker service reports errors. The reason is that unlike on Swarm, Kubernetes does not automatically provide a stable networking endpoint for deployments. We need to create at least a ClusterIP service for each of our deployments so they can communicate.

- List your current services:

```
[centos@node-0 ~]$ kubectl get services
NAME          TYPE          CLUSTER-IP    EXTERNAL-IP    PORT(S)    AGE
kubernetes    ClusterIP     10.96.0.1     <none>         443/TCP    46m
```

- Expose the redis, rng and hasher internally to your cluster, specifying the correct internal port:

```
[centos@node-0 ~]$ kubectl expose deployment redis --port 6379
[centos@node-0 ~]$ kubectl expose deployment rng --port 80
[centos@node-0 ~]$ kubectl expose deployment hasher --port 80
```

- List your services again:

```
[centos@node-0 ~]$ kubectl get services
NAME          TYPE          CLUSTER-IP    EXTERNAL-IP    PORT(S)    AGE
hasher        ClusterIP     10.108.207.22 <none>         80/TCP     20s
kubernetes    ClusterIP     10.96.0.1     <none>         443/TCP    47m
redis         ClusterIP     10.100.14.121 <none>         6379/TCP   31s
rng           ClusterIP     10.111.235.252 <none>         80/TCP     26s
```

Evidently `kubectl expose` creates ClusterIP services allowing stable, internal reachability for your deployments, much like you did via yaml manifests for your nginx deployment in the last section. See the `kubectl` api docs for more command-line alternatives to yaml manifests.

- Get the logs of the worker again:

```
[centos@node-0 ~]$ kubectl logs deploy/worker
```

This time you should see that the worker recovered (give it at least 10 seconds to do so). The worker can now access the other services.

- Now let's expose the webui to the public using a service of type NodePort:

```
[centos@node-0 ~]$ kubectl expose deploy/webui --type=NodePort --port 80
```

- List your services one more time:

```
[centos@node-0 ~]$ kubectl get services
NAME          TYPE          CLUSTER-IP    EXTERNAL-IP    PORT(S)    AGE
hasher        ClusterIP     10.108.207.22 <none>         80/TCP     2m
kubernetes    ClusterIP     10.96.0.1     <none>         443/TCP    49m
redis         ClusterIP     10.100.14.121 <none>         6379/TCP   2m
rng           ClusterIP     10.111.235.252 <none>         80/TCP     2m
webui         NodePort      10.108.88.182 <none>         80:32015/TCP 33s
```

Notice the NodePort service created for webui. This type of service provides similar behavior to the Swarm L4 mesh net: a port (32015 in my case) has been reserved across the cluster; any external traffic hitting any cluster IP on that port will be directed to port 80 inside a webui pod.

11. Visit your Dockercoins web ui at `http://<node IP>:<port>`, where <node IP> is the public IP address any of your cluster members. You should see the dashboard of our DockerCoins application.
12. Let's scale up the worker a bit and see the effect of it:

```
[centos@node-0 ~]$ kubectl scale deploy/worker --replicas=10
```

Observe the result of this scaling in the browser. We do not really get a 10-fold increase in throughput, just as when we deployed DockerCoins on swarm; the rng service is causing a bottleneck.

13. To scale up, we want to run an instance of rng on each node of the cluster. For this we use a DaemonSet. We do this by using a yaml file that captures the desired configuration, rather than through the CLI.

Create a file `deploy-rng.yaml` as follows:

```
[centos@node-0 ~]$ kubectl get deploy/rng -o yaml --export > deploy-rng.yaml
```

Note: `--export` will remove "cluster-specific" information

14. Edit this file to make it describe a DaemonSet instead of a Deployment:

- change kind to DaemonSet
- remove the `progressDeadlineSeconds` field
- remove the `replicas` field
- remove the `strategy` block (which defines the rollout mechanism for a deployment)
- remove the `status: {}` line at the end

15. Now apply this YAML file to create the DaemonSet:

```
[centos@node-0 ~]$ kubectl apply -f deploy-rng.yaml
```

16. We can now look at the DaemonSet that was created:

```
[centos@node-0 ~]$ kubectl get daemonset
```

NAME	DESIRED	CURRENT	READY	UP-TO-DATE	AVAILABLE	NODE SELECTOR	AGE
rng	2	2	2	2	2	<none>	1m

Dockercoins performance should now improve, as illustrated by your web ui.

17. If we do a `kubectl get all` we will see that we now have both a `deployment.apps/rng` AND a `daemonset.apps/rng`. Deployments are not just converted to Daemon sets. Let's delete the rng deployment:

```
[centos@node-0 ~]$ kubectl delete deploy/rng
```

18. Clean up your resources when done:

```
[centos@node-0 ~]$ for D in redis hasher rng webui; \
do kubectl delete svc/$D; done
[centos@node-0 ~]$ for D in redis hasher webui worker; \
do kubectl delete deploy/$D; done
[centos@node-0 ~]$ kubectl delete ds/rng
```

19. Make sure that everything is cleared:

```
[centos@node-0 ~]$ kubectl get all
```

should only show the `svc/kubernetes` resource.

23.4 Conclusion

In this exercise, we looked at some of the key Kubernetes service objects that provide routing and load balancing for collections of pods; clusterIP for internal communication, analogous to Swarm's VIPs, and NodePort, for routing external traffic to an app similarly to Swarm's L4 mesh net. We also briefly touched on the inner workings of Calico, one of many Kubernetes network plugins and the one that ships natively with Docker's Enterprise Edition product. The key networking difference between Swarm and Kubernetes is their approach to default firewalling; while Swarm firewalls software defined networks automatically, all pods can reach all other pods on a Kube cluster, in Calico's case via the BGP-updated control plane and IP-in-IP data plane you explored above.

24 Cleaning up Docker Resources

By the end of this exercise, you should be able to:

- Assess how much disk space docker objects are consuming
- Use `docker prune` commands to clear out unneeded docker objects
- Apply label based filters to prune commands to control what gets deleted in a cleanup operation

1. Find out how much memory Docker is using by executing:

```
[centos@node-3 ~]$ docker system df
```

The output will show us how much space images, containers and local volumes are occupying and how much of this space can be reclaimed.

2. Reclaim all reclaimable space by using the following command:

```
[centos@node-3 ~]$ docker system prune
```

Answer with `y` when asked if we really want to remove all unused networks, containers, images and volumes.

3. Create a couple of containers with labels (these will exit immediately; why?):

```
[centos@node-3 ~]$ docker container run --label apple --name fuji -d alpine
[centos@node-3 ~]$ docker container run --label orange --name clementine -d alpine
```

4. Delete only those stopped containers bearing the apple label:

```
[centos@node-3 ~]$ docker container ls -a
[centos@node-3 ~]$ docker container prune --filter 'label=apple'
[centos@node-3 ~]$ docker container ls -a
```

Only the container named `clementine` should remain after the targeted prune.

5. Finally, prune containers launched before a given timestamp using the `until` filter; start by getting the current RFC 3339 time (<https://tools.ietf.org/html/rfc3339> - note Docker *requires* the otherwise optional `T` separating date and time), then creating a new container:

```
[centos@node-3 ~]$ TIMESTAMP=$(date --rfc-3339=seconds | sed 's/ /T/')
[centos@node-3 ~]$ docker container run --label tomato --name beefsteak -d alpine
```

And use the timestamp returned in a prune:

```
[centos@node-3 ~]$ docker container prune -f --filter "until=$TIMESTAMP"
[centos@node-3 ~]$ docker container ls -a
```

Note the `-f` flag, to suppress the confirmation step. `label` and `until` filters for pruning are also available for networks and images, while data volumes can only be selectively pruned by `label`; finally, images can also be pruned by the boolean `dangling` key, indicating if the image is untagged.

24.1 Conclusion

In this exercise, we saw some very basic `docker prune` usage - most of the top-level docker objects have a `prune` command (`docker container prune`, `docker volume prune` etc). Most docker objects leave something on disk even after being shut down; consider using these cleanup commands as part of your cluster maintenance and garbage collection plan, to avoid accidentally running out of disk on your Docker hosts.

25 Inspection Commands

By the end of this exercise, you should be able to:

- Gather system level info from the docker engine
- Consume and format the docker engine's event stream for monitoring purposes

25.1 Inspecting System Information

1. We can find the `info` command under `system`. Execute:

```
[centos@node-3 ~]$ docker system info
```

This provides some high-level information about the docker deployment on the current node, and the node itself. From this output, identify:

- how many images are cached on your machine?
- how many containers are running or stopped?
- what version of containerd are you running?
- whether Docker is running in swarm mode?

25.2 Monitoring System Events

1. There is another powerful system command that allows us to monitor what's happening on the Docker host. Execute the following command:

```
[centos@node-3 ~]$ docker system events
```

Please note that it looks like the system is hanging, but that is not the case. The system is just waiting for some events to happen.

2. Open a second connection to `node-3` and execute the following command:

```
[centos@node-3 ~]$ docker container run --rm alpine echo 'Hello World!'
```

and observe the generated output in the first terminal. It should look similar to this:

```
2017-01-25T16:57:48.553596179-06:00 container create 30eb63 ...
2017-01-25T16:57:48.556718161-06:00 container attach 30eb63 ...
2017-01-25T16:57:48.698190608-06:00 network connect de1b2b ...
2017-01-25T16:57:49.062631155-06:00 container start 30eb63 ...
2017-01-25T16:57:49.065552570-06:00 container resize 30eb63 ...
2017-01-25T16:57:49.164526268-06:00 container die 30eb63 ...
2017-01-25T16:57:49.613422740-06:00 network disconnect de1b2b ...
2017-01-25T16:57:49.815845051-06:00 container destroy 30eb63 ...
```

Granular information about every action taken by the Docker engine is presented in the events stream.

3. If you don't like the format of the output then we can use the `--format` parameter to define our own format in the form of a Go template. Stop the events watch on your first terminal with `CTRL+C`, and try this:

```
[centos@node-3 ~]$ docker system events --format '--> {{.Type}}-{{.Action}}'
```

now the output looks a little bit less cluttered when we run our alpine container on the second terminal as above.

4. Finally we can find out what the event structure looks like by outputting the events in json format (once again after killing the events watcher on the first terminal and restarting it with):

```
[centos@node-3 ~]$ docker system events --format '{{json .}}' | jq
```

which should give us for the first event in the series after re-running our alpine container on the other connection to node-3 something like this (note, the output has been prettyfied for readability):

```
{
  "status": "create",
  "id": "95ddb6ed4c87d67fa98c3e63397e573a23786046e00c2c68a5bcb9df4c17635c",
  "from": "alpine",
  "Type": "container",
  "Action": "create",
  "Actor": {
    "ID": "95ddb6ed4c87d67fa98c3e63397e573a23786046e00c2c68a5bcb9df4c17635c",
    "Attributes": {
      "image": "alpine",
      "name": "sleepy_roentgen"
    }
  },
  "time": 1485385702,
  "timeNano": 1485385702748011034
}
```

25.3 Conclusion

In this exercise we have learned how to inspect system wide properties of our Docker host by using the `docker system info` command; this is one of the first places to look for general config information to include in a bug report. We also saw a simple example of `docker system events`; the events stream is one of the primary sources of information that should be logged and monitored when running Docker in production. Many commercial as well as open source products (such as Elastic Stack) exist to facilitate aggregating and mining these streams at scale.

26 Plugins

By the end of this exercise, you should be able to:

- Install, configure, and delete any Docker plugin
- Use the `vieux/sshfs` plugin to create ssh-mountable volumes that can be mounted into any container in your cluster

26.1 Installing a Plugin

1. Plugins can be hosted on Docker Hub or any other (private) repository. Let's start with Docker Hub. Browse to <https://hub.docker.com> and enter `vieux/sshfs` in the search box. The result should show you the plugin that we are going to work with.
2. Install the plugin into our Docker Engine:

```
[centos@node-0 ~]$ docker plugin install vieux/sshfs
```

The system should ask us for permission to use privileges. In the case of the `sshfs` plugin there are 4 privileges. Answer with `y`.

3. Once we have successfully installed some plugins we can use the `ls` command to see the status of each of the installed plugins. Execute:

```
[centos@node-0 ~]$ docker plugin ls
```

26.2 Enabling and Disabling a Plugin

1. Once a plugin is installed it is enabled by default. We can disable it using this command:

```
[centos@node-0 ~]$ docker plugin disable vieux/sshfs
```

only when a plugin is disabled can certain operations on it be executed.

2. The plugin can be (re-) enabled by using this command:

```
[centos@node-0 ~]$ docker plugin enable vieux/sshfs
```

Play with the above commands and notice how the status of the plugin changes when displaying it with `docker plugin ls`.

26.3 Inspecting a Plugin

1. We can also use the `inspect` command to further inspect all the attributes of a given plugin. Execute the following command:

```
[centos@node-0 ~]$ docker plugin inspect vieux/sshfs
```

and examine the output. Specifically note that there are two sections in the metadata called `Env`, one is under `Config` and the other under `Settings`. This is where the list of environment variables are listed that the author of the plugin has defined. In this specific situation we can see that there is a single variable called `DEBUG` defined. Its initial value is 0.

2. We can use the `set` command to change values of the environment variables. Execute:

```
[centos@node-0 ~]$ docker plugin set vieux/sshfs DEBUG=1
```

```
Error response from daemon: cannot set on an active plugin,
  disable plugin before setting
```

This is one of those times we have to disable the plugin first; do so, then try the `set` command again:

```
[centos@node-0 ~]$ docker plugin disable vieux/sshfs
[centos@node-0 ~]$ docker plugin set vieux/sshfs DEBUG=1
[centos@node-0 ~]$ docker plugin enable vieux/sshfs
```

and then inspect again the metadata of the plugin. Notice how the value of `DEBUG` has been adjusted. Only the one under the `Settings` node changed but the one under the `Config` node still shows the original (default) value.

26.4 Using the Plugin

1. Make a directory on `node-1` that we will mount as a volume across our cluster:

```
[centos@node-1 ~]$ mkdir ~/demo
```

2. Back on `node-0`, use the plugin to create a volume that can be mounted via `ssh`:

```
[centos@node-0 ~]$ docker volume create -d vieux/sshfs \
  -o sshcmd=centos@<node-1 public IP>:/home/centos/demo \
  -o password=orca \
  sshvolume
```

3. Mount that volume in a new container as per usual:

```
[centos@node-0 ~]$ docker container run --rm -it -v sshvolume:/data alpine sh
```

4. Inside the container navigate to the /data folder and create a new file:

```
/ # cd /data
/ # echo 'Hello from client!' > demo.txt
/ # ls -al
```

5. Head over to node-1, and confirm that demo.txt got written there.

26.5 Removing a Plugin

1. If we don't want or need this plugin anymore we can remove it using the command:

```
[centos@node-0 ~]$ docker volume rm sshvolume
[centos@node-0 ~]$ docker plugin disable vieux/sshfs
[centos@node-0 ~]$ docker plugin rm vieux/sshfs
```

Note how we first have to disable the plugin before we can remove it.

26.6 Conclusion

Docker follows a 'batteries included but swappable' mindset in its product design: everything you need to get started is included, but heavy customization is supported and encouraged. Docker plugins are one aspect of that flexibility, allowing users to define their own volume and networking behavior.

27 Starting a Compose App

By the end of this exercise, you should be able to:

- Read a basic docker compose yaml file and understand what components it is declaring
- Start, stop, and inspect the logs of an application defined by a docker compose file

27.1 Inspecting a Compose App

1. Download the Dockercoins app from github:

```
[centos@node-0 ~]$ git clone -b ee3.0 \
  https://github.com/docker-training/orchestration-workshop.git
[centos@node-0 ~]$ cd orchestration-workshop/dockercoins
```

This app consists of 5 services: a random number generator rng, a hasher, a backend worker, a redis queue, and a web frontend; the code you just downloaded has the source code for each process and a Dockerfile to containerize each of them.

2. Have a brief look at the source for each component of your application. Each folder under ~/orchestration-workshop/dockercoins contains the application logic for the component, and a Dockerfile for building that logic into a Docker image. We've pre-built these images as training/dockercoins-rng:1.0, training/dockercoins-worker:1.0 et cetera, so no need to build them yourself.

3. Have a look in `docker-compose.yml`; especially notice the `services` section. Each block here defines a different Docker service. They each have exactly one image which containers for this service will be started from, as well as other configuration details like network connections and port exposures. Full syntax for Docker Compose files can be found here: <https://docs.docker.com/compose/yml/>.

27.2 Starting the App

1. Stand up the app:

```
[centos@node-0 dockercoins]$ docker-compose up
```

After a moment, your app should be running; visit `<node 0 public IP>:8000` to see the web frontend visualizing your rate of Dockercoin mining.

2. Logs from all the running services are sent to `STDOUT`. Let's send this to the background instead; kill the app with `CTRL+C`, sending a `SIGTERM` to all running processes; some exit immediately, while others wait for a 10s timeout before being killed by a subsequent `SIGKILL`. Start the app again in the background:

```
[centos@node-0 dockercoins]$ docker-compose up -d
```

3. Check out which containers are running thanks to Compose:

```
[centos@node-0 dockercoins]$ docker-compose ps
```

Name	Command	State	Ports
dockercoins_hasher_1	ruby hasher.rb	Up	0.0.0.0:8002->80/tcp
dockercoins_redis_1	docker-entrypoint.sh redis ...	Up	6379/tcp
dockercoins_rng_1	python rng.py	Up	0.0.0.0:8001->80/tcp
dockercoins_webui_1	node webui.js	Up	0.0.0.0:8000->80/tcp
dockercoins_worker_1	python worker.py	Up	

4. Compare this to the usual `docker container ls`; do you notice any differences? If not, start a couple of extra containers using `docker container run...`, and check again.

27.3 Viewing Logs

1. See logs from a Compose-managed app via:

```
[centos@node-0 dockercoins]$ docker-compose logs
```

2. The logging API in Compose follows the main Docker logging API closely. For example, try following the tail of the logs just like you would for regular container logs:

```
[centos@node-0 dockercoins]$ docker-compose logs --tail 10 --follow
```

Note that when following a log, `CTRL+S` and `CTRL+Q` pauses and resumes live following; `CTRL+C` exits follow mode as usual.

27.4 Conclusion

In this exercise, you saw how to start a pre-defined Compose app, and how to inspect its logs. Application logic was defined in each of the five images we used to create containers for the app, but the manner in which those containers were created was defined in the `docker-compose.yml` file; all runtime configuration for each container is captured in this manifest. Finally, the different elements of Dockercoins communicated with each other via service name; the Docker daemon's internal DNS was able to resolve traffic destined for a service, into the IP or MAC address of the corresponding container.

28 Scaling a Compose App

By the end of this exercise, you should be able to:

- Scale a service from Docker Compose up or down

28.1 Scaling a Service

Any service defined in our `docker-compose.yml` can be scaled up from the Compose API; in this context, 'scaling' means launching multiple containers for the same service, which Docker Compose can route requests to and from.

1. Scale up the worker service in our Dockercoins app to have two workers generating coin candidates by re-deploying the app with the `--scale` flag, while checking the list of running containers before and after:

```
[centos@node-0 dockercoins]$ docker-compose ps
[centos@node-0 dockercoins]$ docker-compose up -d --scale worker=2
[centos@node-0 dockercoins]$ docker-compose ps
```

Name	Command	State	Ports
dockercoins_hasher_1	ruby hasher.rb	Up	0.0.0.0:8002->80/tcp
dockercoins_redis_1	docker-entrypoint.sh redis ...	Up	6379/tcp
dockercoins_rng_1	python rng.py	Up	0.0.0.0:8001->80/tcp
dockercoins_webui_1	node webui.js	Up	0.0.0.0:8000->80/tcp
dockercoins_worker_1	python worker.py	Up	
dockercoins_worker_2	python worker.py	Up	

A new worker container has appeared in your list of containers.

2. Look at the performance graph provided by the web frontend; the coin mining rate should have doubled. Also check the logs using the logging API we learned in the last exercise; you should see a second worker instance reporting.

28.2 Investigating Bottlenecks

1. Try running `top` to inspect the system resource usage; it should still be fairly negligible. So, keep scaling up your workers:

```
[centos@node-0 dockercoins]$ docker-compose up -d --scale worker=10
[centos@node-0 dockercoins]$ docker-compose ps
```

2. Check your web frontend again; has going from 2 to 10 workers provided a 5x performance increase? It seems that something else is bottlenecking our application; any distributed application such as Dockercoins needs tooling to understand where the bottlenecks are, so that the application can be scaled intelligently.
3. Look in `docker-compose.yml` at the `rng` and `hasher` services; they're exposed on host ports 8001 and 8002, so we can use `httping` to probe their latency.

```
[centos@node-0 dockercoins]$ httping -c 5 localhost:8001
[centos@node-0 dockercoins]$ httping -c 5 localhost:8002
```

`rng` on port 8001 has the much higher latency, suggesting that it might be our bottleneck. A random number generator based on entropy won't get any better by starting more instances on the same machine; we'll need a way to bring more nodes into our application to scale past this, which we'll explore in the next unit on Docker Swarm.

4. For now, shut your app down:

```
[centos@node-0 dockercoins]$ docker-compose down
```

28.3 Conclusion

In this exercise, we saw how to scale up a service defined in our Compose app using the `--scale` flag. Also, we saw how crucial it is to have detailed monitoring and tooling in a microservices-oriented application, in order to correctly identify bottlenecks and take advantage of the simplicity of scaling with Docker.

29 Appendix: A Brief Introduction to Bash

If you're not familiar with the common Linux commands (`cd`, `ls`, `ps`, `sudo`, etc.), this exercise is designed to lead you through the common linux commands that we use in the Docker training.

By the end of this exercise, you should be able to:

- Manage files and folders using Linux commands
- Inspect processes running on your machine and use some discovery tools.

Pre-requisites:

- A Linux machine.

29.1 Managing Files and Folders

In this first part we will walk through common commands for creating and manipulating files and folders from the command line.

1. Open a terminal.
2. Terminals work by accepting text-based commands, and returning text-based responses. Ask the terminal what directory you are currently in; in all our instructions, `$` represents the terminal prompt, and you should enter what you see *after* the prompt. Yours might look a bit different; in any case, type `pwd`:

```
$ pwd
```

`pwd` stands for Print Working Directory. You should see something like:

```
/Users/billmills
```

This means your terminal is currently in the `billmills` folder, which itself is in the `/Users` folder (yours will have different folder names, but the logic is the same).

3. List the files and folders in your current directory:

```
$ ls
```

```
Applications  Documents  Library    Music
Projects      Desktop
```

Here we show some typical output below the command; again, you only need to enter what comes after the command prompt, `ls` in this case. `ls` lists all the contents of your current directory.

4. Change directory with `cd`:

```
cd Desktop
```

Again, your directories might be named differently - that's ok. Try to navigate to a directory you recognize the names of; `Desktop` or `Downloads` are common.

5. Use `ls` again to see what's in this directory. Compare what you see in your `Desktop` directory to what you actually see on your machine's desktop; the contents are identical. The terminal is just another way to browse the contents of your machine, based on programmatic text rather than visual analogies.
6. Make a new folder on your desktop (or wherever you currently are):

```
$ mkdir demo-dir
```

If you do `ls` again, you'll see a new directory `demo-dir` has been made.

7. Change directory to `demo-dir`, and open a new plain text file called `demo.txt` using the simple text editor `nano`:

```
$ cd demo-dir
$ nano demo.txt
```

You're taken to `nano`'s text-based interface for creating a plain-text file. There are many other plain-text editors available, the most popular being `vim`, but `nano` is probably the simplest if this is your first time using a text editor from the command line.

8. Add any text you want to your file, just by typing something.
9. Save your file by pressing `CMD+O`, and pressing return when asked for a filename. Exit by typing `CMD+X`.
10. Dump the contents of your file to the screen:

```
$ cat demo.txt
```

11. Make a copy of your file with `cp`:

```
$ cp demo.txt demo.copy.txt
```

12. Move and rename your file with `mv`:

```
$ mv demo.txt demo.moved.txt
```

13. Delete a file:

```
$ rm demo.moved.txt
```

14. Check what directory you're in, then back up one level with the special `..` path, and check again:

```
$ pwd

/Users/billmills/Desktop/demo-dir

$ cd ..
$ pwd
/Users/billmills/Desktop
```

`cd ..` took us one level up in our directory tree, from `demo-dir` back out to the directory that contains it, `Desktop` in my example.

15. Delete your `demo-dir`, again with `remove`, but this time using the `-r` flag, which stands for recursive in order to delete the folder and all its contents:

```
rm -r demo-dir
```

29.2 Some Common Tools

In this section we will see the basic usage of different command line tools we'll see again when working with Docker.

29.2.1 Inspecting Processes

1. In the next steps we will use `ps` and `top`. Check whether these tools are existing in our system:

```
$ which pwd
```

`which` is a quick way to identify the location of executables and should return something similar to the following:

```
/bin/pwd
```

The `pwd` command we used above is actually a tiny program, and it lives in our filesystem at the path indicated by `which`. If `which` returns nothing, that means the requested program probably isn't installed on your machine.

2. List all processes running in your system:

```
$ ps -aux
```

A long list of processes is returned.

We will see later how to filter this output to extract only the information we need.

3. There different other tools to inspect processes. Try for example showing the list of processes with usage details (cpu usage, memory, etc.):

```
$ top
```

Exit with CTRL+C.

29.2.2 The Superuser

In some cases you want to run commands that require privileges that you don't have.

1. Try creating a new user:

```
$ adduser myName
```

This should return `Permission denied`.

2. Run the same command with `sudo`:

```
$ sudo adduser myName
```

`Sudo` runs the command with the privileges of the super user.

3. List the existing users:

```
$ cat /etc/passwd
```

You should see the user `myName` in the bottom of the list.

29.2.3 Pinging of an address

`ping` is a tool used to test the reachability of a network address.

1. Send a ping to your `localhost`:

```
$ ping localhost
```

The ping should be successful. Interrupt it with CTRL+C.

2. Try pinging an unreachable address:

```
$ ping -c 3 192.168.1.1
```

The flag `-c` stands for the count of the sent packets. If no device is plugged to your machine and has this address, the request should timeout after 3 packets with 100% packet loss.

29.2.4 Making HTTP Requests

Use the `curl` command to issue HTTP requests across the network.

1. `curl` an example webpage:

```
$ curl example.com
```

You'll get some HTML corresponding to a dummy webpage, downloaded directly to your terminal.

29.2.5 Working with Commands

29.2.5.1 Command piping

So far, every command we've used has accepted text as input, and returned text as output. We can send the text output from one command into the text input of another command using a *pipe*, `|`.

1. Earlier we saw the `ps` command, to write a large table of all the processes running on our machine. We can send that table to `grep`, which is a text search tool that will pick out lines containing something we're interested in:

```
$ ps -aux | grep 'ps'
```

Rather than getting every process on the system, we can just pick out the `ps` process by text-searching for it using `grep`.

2. Another common `grep` usage is with `cat`, to find a string in a file. Search your `/etc/passwd` file for the root user:

```
cat /etc/passwd | grep root
```

Instead of getting every user on the system, only lines with the string `root` are printed out, making it easier to find what you're looking for.

29.2.5.2 Successive Commands

We can run several commands in a one liner using semicolon `;` or double-ampersand `&&`.

1. Create a new directory:

```
$ mkdir newDir
```

2. We know that to remove a directory we need to use `rm -r`. Let's simulate a error by forgetting the `-r` flag, and immediately after removing we'll create a new directory with the same name:

```
$ rm newDir ; mkdir newDir
```

This should return:

```
rm: cannot remove 'newDir': Is a directory
mkdir: cannot create directory 'newDir': File exists
```

The semicolon `;` runs the second command even when the first command wasn't successful.

3. Combining commands with the double-ampersand `&&` insures that the second command will run only if the first command was successful. Try the following:

```
$ rm newDir && mkdir newDir
```

This should return the error for the first command only:

```
rm: cannot remove 'newDir': Is a directory
```

29.2.5.3 Breaking long commands

Some commands get too long, for example because it includes a long file path or it's a one liner that combines several commands. For better readability, you can break long commands into several lines using the backslash:

```
...
$ mkdir aDirectoryWithAVeryLongName ; \
  cd aDirectoryWithAVeryLongName ; \
  echo "this is a test file" > myTestFile ; \
  cat myTestFile ; cd ..
...
```

This one liner will create a directory, `cd` into it, create a file, `cat` the content of the file, and finally change directory a level up.

29.3 Conclusion

We saw most of the Linux commands that we will use in the actual Docker training. Feel free to discover more commands in the following cheat sheet: <https://www.git-tower.com/blog/command-line-cheat-sheet/>

Instructor Demos

1 Instructor Demo: Process Isolation

In this demo, we'll illustrate:

- What containerized process IDs look like inside versus outside of a kernel namespace
- How to impose control group limitations on CPU and memory consumption of a containerized process.

1.1 Exploring the PID Kernel Namespace

1. Start a simple container we can explore:

```
[centos@node-0 ~]$ docker container run -d --name pinger centos:7 ping 8.8.8.8
```

2. Use `docker container exec` to launch a child process inside the container's namespaces:

```
[centos@node-0 ~]$ docker container exec pinger ps -aux
```

USER	PID	%CPU	%MEM	VSZ	RSS	TTY	STAT	START	TIME	COMMAND
root	1	0.1	0.0	24860	1884	?	Ss	02:20	0:00	ping 8.8.8.8
root	5	0.0	0.0	51720	3504	?	Rs	02:20	0:00	ps -aux

3. Run the same `ps` directly on the host, and search for your ping process:

```
[centos@node-0 ~]$ ps -aux | grep ping
```

USER	PID	%CPU	%MEM	VSZ	RSS	TTY	STAT	START	TIME	COMMAND
root	11622	0.0	0.0	24860	1884	?	Ss	02:20	0:00	ping 8.8.8.8
centos	11839	0.0	0.0	112656	2132	pts/0	S+	02:23	0:00	grep --color=auto ping

The ping process appears as PID 1 inside the container, but as some higher PID (11622 in this example) from outside the container.

4. List your containers to show this ping container is still running:

```
[centos@node-0 ~]$ docker container ls
```

CONTAINER ID	IMAGE	COMMAND	...	STATUS	...	NAMES
bb3a3b1cbb78	centos:7	"ping 8.8.8.8"	...	Up 6 minutes		pinger

Kill the ping process by host PID, and show the container has stopped:

```
[centos@node-0 ~]$ sudo kill -9 [host PID of ping]
[centos@node-0 ~]$ docker container ls
```

CONTAINER ID	IMAGE	COMMAND	...	STATUS	...	NAMES
--------------	-------	---------	-----	--------	-----	-------

Killing the ping process on the host also kills the container - all a running container is its PID 1 process, and the kernel tooling that isolates it from the host. Note using `kill -9` is just for demonstration purposes here; never stop containers this way.

1.2 Imposing Resource Limitations With Cgroups

1. Start a container that consumes two full CPUs:

```
[centos@node-0 ~]$ docker container run -d training/stress:3.0 --vm 2
```

Here the `--vm` flag starts 2 dummy processes that allocate and free memory as fast as they can, each consuming as many CPU cycles as possible.

2. Check the CPU consumption of processes in the container:

```
[centos@node-0 ~]$ docker container top <container ID>
```

UID	PID	PPID	C	...	CMD
root	5806	5789	0	...	/usr/bin/stress --verbose --vm 2
root	5828	5806	99	...	/usr/bin/stress --verbose --vm 2
root	5829	5806	99	...	/usr/bin/stress --verbose --vm 2

That C column represents CPU consumption, in percent; this container is hogging two full CPUs! See the same thing by running `ps -aux` both inside and outside this container, like we did above; the same process and its CPU utilization is visible inside and outside the container:

```
[centos@node-0 ~]$ docker container exec <container ID> ps -aux
```

USER	PID	%CPU	%MEM	...	COMMAND
root	1	0.0	0.0	...	/usr/bin/stress --verbose --vm 2
root	5	98.9	6.4	...	/usr/bin/stress --verbose --vm 2
root	6	99.0	0.4	...	/usr/bin/stress --verbose --vm 2
root	7	2.0	0.0	...	ps -aux

And on the host directly, via the PIDs we found from `docker container top` above:

```
[centos@node-0 ~]$ ps -aux | grep <PID>
```

USER	PID	%CPU	%MEM	...	COMMAND
root	5828	99.3	4.9	...	/usr/bin/stress --verbose --vm 2
centos	6327	0.0	0.0	...	grep --color=auto 5828

3. Kill off this container:

```
[centos@node-0 ~]$ docker container rm -f <container ID>
```

This is the right way to kill and remove a running container (not `kill -9`).

4. Run the same container again, but this time with a cgroup limitation on its CPU consumption:

```
[centos@node-0 ~]$ docker container run -d --cpus="1" training/stress:3.0 --vm 2
```

Do `docker container top` and `ps -aux` again, just like above; you'll see the processes taking up half a CPU each, for a total of 1 CPU consumed. The `--cpus="1"` flag has imposed a control group limitation on the processes in this container, constraining them to consume a total of no more than one CPU.

5. Find the host PID of a process running in this container using `docker container top` again, and then see what cgroups that process lives in on the host:

```
[centos@node-0 ~]$ cat /proc/<host PID of containerized process>/cgroup
```

```
12:memory:/docker/31d03...
11:freezer:/docker/31d03...
10:hugetlb:/docker/31d03...
9:perf_event:/docker/31d03...
8:net_cls,net_prio:/docker/31d03...
7:cpuset:/docker/31d03...
```

```
6:pids:/docker/31d03...
5:blkio:/docker/31d03...
4:rdma:/
3:devices:/docker/31d03...
2:cpu,cpuacct:/docker/31d03...
1:name=systemd:/docker/31d03...
```

6. Get a summary of resources consumed by processes in a control group via `systemd-cgtop`:

```
[centos@node-0 ~]$ systemd-cgtop
```

Path	Tasks	%CPU	Memory	Input/s	Output/s
/	68	112.3	1.0G	-	-
/docker	-	99.3	301.0M	-	-
/docker/31d03...	3	99.3	300.9M	-	-
...					

Here again we can see that the processes living in the container's control group (`/docker/31d03...`) are constrained to take up only about 1 CPU.

7. Remove this container, spin up a new one that creates a lot of memory pressure, and check its resource consumption with `docker stats`:

```
[centos@node-0 ~]$ docker container rm -f <container ID>
[centos@node-0 ~]$ docker container run -d training/stress:3.0 --vm 2 --vm-bytes 1024M
[centos@node-0 ~]$ docker stats
```

CONTAINER	CPU %	MEM USAGE / LIMIT	MEM %	...
b29a6d877343	198.94%	937.2MiB / 3.854GiB	23.75%	...

8. Kill this container off, start it again with a memory constraint, and list your containers:

```
[centos@node-0 ~]$ docker container rm -f <container ID>
[centos@node-0 ~]$ docker container run \
    -d -m 256M training/stress:3.0 --vm 2 --vm-bytes 1024M
[centos@node-0 ~]$ docker container ls -a
```

CONTAINER ID	IMAGE	...	STATUS
296c8f76af5c	training/stress:3.0	...	Exited (1) 26 seconds ago

It exited immediately this time.

9. Inspect the metadata for this container, and look for the `OOMKilled` key:

```
[centos@node-0 ~]$ docker container inspect <container ID> | grep 'OOMKilled'

"OOMKilled": true,
```

When the containerized process tried to exceed its memory limitation, it gets killed with an Out Of Memory exception.

1.3 Conclusion

In this demo, we explored some of the most important technologies that make containerization possible: kernel namespaces and control groups. The core message here is that containerized processes are just processes running on their host, isolated and constrained by these technologies. All the tools and management strategies you would use for conventional processes apply just as well for containerized processes.

2 Instructor Demo: Creating Images

In this demo, we'll illustrate:

- How to read each step of the image build output
- How intermediate image layers behave in the cache and as independent images
- What the meanings of 'dangling' and <missing> image layers are

2.1 Understanding Image Build Output

1. Make a folder demo for our image demo:

```
[centos@node-0 ~]$ mkdir demo ; cd demo
```

And create a Dockerfile therein with the following content:

```
FROM centos:7
RUN yum update -y
RUN yum install -y which
RUN yum install -y wget
RUN yum install -y vim
```

2. Build your image from your Dockerfile, just like we did in the last exercise:

```
[centos@node-0 demo]$ docker image build -t demo .
```

3. Examine the output from the build process. The very first line looks like:

```
Sending build context to Docker daemon 2.048kB
```

Here the Docker daemon is archiving everything at the path specified in the `docker image build` command (`.` or the current directory in this example). This is why we made a fresh directory `demo` to build in, so that nothing extra is included in this process.

4. The next lines look like:

```
Step 1/5 : FROM centos:7
---> 49f7960eb7e4
```

Do an image `ls`:

```
[centos@node-0 demo]$ docker image ls
```

REPOSITORY	TAG	IMAGE ID	CREATED	SIZE
demo	latest	59e595750dd5	10 seconds ago	645MB
centos	7	49f7960eb7e4	2 months ago	200MB

Notice the Image ID for `centos:7` matches that second line in the build output. The build starts from the base image defined in the `FROM` command.

5. The next few lines look like:

```
Step 2/5 : RUN yum update -y
---> Running in 8734b14cf011
Loaded plugins: fastestmirror, ovl
...
```

This is the output of the `RUN` command, `yum update -y`. The line `Running in 8734b14cf011` specifies a container that this command is running in, which is spun up based on all previous image layers (just the `centos:7` base at the moment). Scroll down a bit and you should see something like:

```
---> 433e56d735f6
Removing intermediate container 8734b14cf011
```

At the end of this first RUN command, the temporary container 8734b14cf011 is saved as an image layer 433e56d735f6, and the container is removed. This is the exact same process as when you used `docker container commit` to save a container as a new image layer, but now running automatically as part of a Dockerfile build.

- Look at the history of your image:

```
[centos@node-0 demo]$ docker image history demo
```

IMAGE	CREATED	CREATED BY	SIZE
59e595750dd5	2 minutes ago	/bin/sh -c yum install -y vim	142MB
bba17f8df167	2 minutes ago	/bin/sh -c yum install -y wget	87MB
b9f2efa616de	2 minutes ago	/bin/sh -c yum install -y which	86.6MB
433e56d735f6	2 minutes ago	/bin/sh -c yum update -y	129MB
49f7960eb7e4	2 months ago	/bin/sh -c #(nop) CMD ["/bin/bash"]	0B
<missing>	2 months ago	/bin/sh -c #(nop) LABEL org.label-schema....	0B
<missing>	2 months ago	/bin/sh -c #(nop) ADD file:8f4b3be0c1427b1...	200MB

As you can see, the different layers of demo correspond to a separate line in the Dockerfile and the layers have their own ID. You can see the image layer 433e56d735f6 committed in the second build step in the list of layers for this image.

- Look through your build output for where steps 3/5 (installing `which`), 4/5 (installing `wget`), and 5/5 (installing `vim`) occur - the same behavior of starting a temporary container based on the previous image layers, running the RUN command, saving the container as a new image layer visible in your `docker image history` output, and deleting the temporary container is visible.
- Every layer can be used as you would use any image, which means we can inspect a single layer. Let's inspect the `wget` layer, which in my case is `bba17f8df167` (yours will be different, look at your `docker image history` output):

```
[centos@node-0 demo]$ docker image inspect bba17f8df167
```

- Let's look for the command associated with this image layer by using `--format`:

```
[centos@node-0 demo]$ docker image inspect \
  --format='{{.ContainerConfig.Cmd}}' bba17f8df167

[/bin/sh -c yum install -y wget]
```

- We can even start containers based on intermediate image layers; start an interactive container based on the `wget` layer, and look for whether `wget` and `vim` are installed:

```
[centos@node-0 demo]$ docker container run -it bba17f8df167 bash

[root@a766a3d616b7 /]# which wget

/usr/bin/wget

[root@a766a3d616b7 /]# which vim

/usr/bin/which: no vim in
(/usr/local/sbin:/usr/local/bin:/usr/sbin:/usr/bin:/sbin:/bin)
```

`wget` is installed in this layer, but since `vim` didn't arrive until the next layer, it's not available here.

2.2 Managing Image Layers

- Change the last line in the Dockerfile from the last section to install `nano` instead of `vim`:

```
FROM centos:7
RUN yum update -y
RUN yum install -y which
RUN yum install -y wget
RUN yum install -y nano
```

2. Rebuild your image, and list your images again:

```
[centos@node-0 demo]$ docker image build -t demo .
[centos@node-0 demo]$ docker image ls
```

REPOSITORY	TAG	IMAGE ID	CREATED	SIZE
demo	latest	5a6aedc1feab	8 seconds ago	590MB
<none>	<none>	59e595750dd5	23 minutes ago	645MB
centos	7	49f7960eb7e4	2 months ago	200MB

What is that image named <none>? Notice the image ID is the same as the old image ID for demo:latest (see your history output above). The name and tag of an image is just a pointer to the stack of layers that make it up; reuse a name and tag, and you are effectively moving that pointer to a new stack of layers, leaving the old one (the one containing the vim install in this case) as an untagged or 'dangling' image.

3. Rewrite your Dockerfile one more time, to combine some of those install steps:

```
FROM centos:7
RUN yum update -y
RUN yum install -y which wget nano
```

Rebuild using a new tag this time, and list your images one more time:

```
[centos@node-0 demo]$ docker image build -t demo:new .
...
[centos@node-0 demo]$ docker image ls
```

REPOSITORY	TAG	IMAGE ID	CREATED	SIZE
demo	new	568b29a0dce9	20 seconds ago	416MB
demo	latest	5a6aedc1feab	5 minutes ago	590MB
<none>	<none>	59e595750dd5	28 minutes ago	645MB
centos	7	49f7960eb7e4	2 months ago	200MB

Image demo:new is much smaller in size than demo:latest, even though it contains the exact same software - why?

2.3 Conclusion

In this demo, we explored the layered structure of images; each layer is built as a distinct image and can be treated as such, on the host where it was built. This information is preserved on the build host for use in the build cache; build another image based on the same lower layers, and they will be reused to speed up the build process. Notice that the same is not true of downloaded images like centos:7; intermediate image caches are not downloaded, but rather only the final complete image.

3 Instructor Demo: Basic Volume Usage

In this demo, we'll illustrate:

- Creating, updating, destroying, and mounting docker named volumes
- How volumes interact with a container's layered filesystem
- Usecases for mounting host directories into a container

3.1 Using Named Volumes

1. Create a volume, and inspect its metadata:

```
[centos@node-0 ~]$ docker volume create demovol
[centos@node-0 ~]$ docker volume inspect demovol

[
  {
    "CreatedAt": "2018-11-03T19:07:56Z",
    "Driver": "local",
    "Labels": {},
    "Mountpoint": "/var/lib/docker/volumes/demovol/_data",
    "Name": "demovol",
    "Options": {},
    "Scope": "local"
  }
]
```

We can see that by default, named volumes are created under `/var/lib/docker/volumes/<name>/_data`.

2. Run a container that mounts this volume, and list the filesystem therein:

```
[centos@node-0 ~]$ docker container run -it -v demovol:/demo centos:7 bash
[root@f4aca1b60965 /]# ls
anaconda-post.log bin demo dev etc home ...
```

The `demo` directory is created as the mountpoint for our volume, as specified in the flag `-v demovol:/demo`. This should also appear in your container filesystem's list of mountpoints:

```
[root@f4aca1b60965 /]# cat /proc/self/mountinfo | grep demo

1199 1180 202:1 /var/lib/docker/volumes/demovol/_data /demo
      rw,relatime - xfs /dev/xvda1 ...
```

3. Put a file in this volume:

```
[root@f4aca1b60965 /]# echo 'dummy file' > /demo/mydata.dat
```

4. Exit the container, and list the contents of your volume on the host:

```
[centos@node-0 ~]$ sudo ls /var/lib/docker/volumes/demovol/_data
```

You'll see your `mydata.dat` file present at this point in the host's filesystem. Delete the container:

```
[centos@node-0 ~]$ docker container rm -f <container ID>
```

The volume and its contents will still be present on the host.

5. Start a new container mounting the same volume, attach a bash shell to it, and show that the old data is present in your new container:

```
[centos@node-0 ~]$ docker container run -d -v demovol:/demo centos:7 ping 8.8.8.8
[centos@node-0 ~]$ docker container exec -it <container ID> bash
[root@11117d3de672 /]# cat /demo/mydata.dat
```

6. Exit this container, and inspect its mount metadata:

```
[centos@node-0 ~]$ docker container inspect <container ID>

  "Mounts": [
    {
      "Type": "volume",
      "Name": "demovol",
```

```

        "Source": "/var/lib/docker/volumes/demovol/_data",
        "Destination": "/demo",
        "Driver": "local",
        "Mode": "z",
        "RW": true,
        "Propagation": ""
    }
],

```

Here too we can see the volumes and host mountpoints for everything mounted into this container.

7. Build a new image out of this container using `docker container commit`, and start a new container based on that image:

```

[centos@node-0 ~]$ docker container commit <container ID> demo:snapshot
[centos@node-0 ~]$ docker container run -it demo:snapshot bash
[root@ad62f304ba18 /]# cat /demo/mydata.dat
cat: /demo/mydata.dat: No such file or directory

```

The information mounted into the original container is not part of the container's layered filesystem, and therefore is not captured in the image creation process; volume mounts and the layered filesystem are completely separate.

8. Clean up by removing that volume:

```

[centos@node-0 ~]$ docker volume rm demovol

```

You will get an error saying the volume is in use - docker will not delete a volume mounted to any container (even a stopped container) in this way. Remove the offending container first, then remove the volume again.

3.2 Mounting Host Paths

1. Make a directory with some source code in it for your new website:

```

[centos@node-0 ~]$ mkdir /home/centos/myweb
[centos@node-0 ~]$ cd /home/centos/myweb
[centos@node-0 myweb]$ echo "<h1>Hello Wrlld</h1>" > index.html

```

2. Start up an nginx container that mounts this as a static website:

```

[centos@node-0 myweb]$ docker container run -d \
    -v /home/centos/myweb:/usr/share/nginx/html \
    -p 8000:80 nginx

```

Visit your website at the public IP of this node, port 8000.

3. Fix the spelling of 'world' in your HTML, and refresh the webpage; the content served by nginx gets updated without having to restart or replace the nginx container.

3.3 Conclusion

In this demo, we saw two key points about volumes: they exist outside the container's layered filesystem, meaning that not only are they not captured on image creation, they don't participate in the usual copy on write procedure when manipulating files in the writable container layer. Second, we saw that manipulating files on the host that have been mounted into a container immediately propagates those changes to the running container; this is a popular technique for developers who containerize their running environment, and mount in their in-development code so they can edit their code using the tools on their host machine that they are familiar with, and have those changes immediately available inside a running container without having to restart or rebuild anything.

4 Instructor Demo: Single Host Networks

In this demo, we'll illustrate:

- Creating docker bridge networks
- Attaching containers to docker networks
- Inspecting networking metadata from docker networks and containers
- How network interfaces appear in different network namespaces
- What network interfaces are created on the host by docker networking
- What iptables rules are created by docker to isolate docker software-defined networks and forward network traffic to containers

4.1 Following Default Docker Networking

1. Switch to a fresh node you haven't run any containers on yet, list your networks:

```
[centos@node-1 ~]$ docker network ls
```

NETWORK ID	NAME	DRIVER	SCOPE
7c4e63830cbf	bridge	bridge	local
c87d2a849036	host	host	local
902af00d5511	none	null	local

2. Get some metadata about the bridge network, which is the default network containers attach to when doing `docker container run`:

```
[centos@node-1 ~]$ docker network inspect bridge
```

Notice the IPAM section:

```
"IPAM": {
  "Driver": "default",
  "Options": null,
  "Config": [
    {
      "Subnet": "172.17.0.0/16",
      "Gateway": "172.17.0.1"
    }
  ]
}
```

Docker's IP address management driver assigns a subnet (172.17.0.0/16 in this case) to each bridge network, and uses the first IP in that range as the network's gateway.

Also note the containers key:

```
"Containers": {}
```

So far, no containers have been plugged into this network.

3. Have a look at what network interfaces are present on this host:

```
[centos@node-1 ~]$ ip addr
```

```
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN qlen 1000
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
    inet 127.0.0.1/8 scope host lo
        valid_lft forever preferred_lft forever
    inet6 ::1/128 scope host
        valid_lft forever preferred_lft forever
2: eth0: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 9001 qdisc mq state UP qlen 1000
```

```

link/ether 12:eb:dd:4e:07:ec brd ff:ff:ff:ff:ff:ff
inet 10.10.17.74/20 brd 10.10.31.255 scope global dynamic eth0
    valid_lft 2444sec preferred_lft 2444sec
inet6 fe80::10eb:ddff:fe4e:7ec/64 scope link
    valid_lft forever preferred_lft forever
3: docker0: <NO-CARRIER,BROADCAST,MULTICAST,UP> mtu 1500 qdisc noqueue state DOWN
link/ether 02:42:e2:c5:a4:6b brd ff:ff:ff:ff:ff:ff
inet 172.17.0.1/16 scope global docker0
    valid_lft forever preferred_lft forever

```

We see the usual eth0 and loopback interfaces, but also the docker0 linux bridge, which corresponds to the docker software defined network we were inspecting in the previous step; note it has the same gateway IP as we found when doing `docker network inspect`.

4. Create a docker container without specifying any networking parameters, and do the same `docker network inspect` as above:

```

[centos@node-1 ~]$ docker container run -d centos:7 ping 8.8.8.8
[centos@node-1 ~]$ docker network inspect bridge

...
"Containers": {
  "f4e8f3f1b918900dd8c9b8867aa3c81e95cf34aba7e366379f2a9ade9987a40b": {
    "Name": "zealous_kirch",
    "EndpointID": "f9f246a...",
    "MacAddress": "02:42:ac:11:00:02",
    "IPv4Address": "172.17.0.2/16",
    "IPv6Address": ""
  }
}
...

```

The Containers key now contains the metadata for the container you just started; it received the next available IP address from the default network's subnet. Also note that the last four digits of the container's MAC address are the same as its IP on this network - this encoding ensures containers get a locally unique MAC address that linux bridges can route traffic to.

5. Look at your network interfaces again:

```

[centos@node-1 ~]$ ip addr

1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN qlen 1000
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
    inet 127.0.0.1/8 scope host lo
        valid_lft forever preferred_lft forever
    inet6 ::1/128 scope host
        valid_lft forever preferred_lft forever
2: eth0: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 9001 qdisc mq state UP qlen 1000
    link/ether 12:eb:dd:4e:07:ec brd ff:ff:ff:ff:ff:ff
    inet 10.10.17.74/20 brd 10.10.31.255 scope global dynamic eth0
        valid_lft 2188sec preferred_lft 2188sec
    inet6 fe80::10eb:ddff:fe4e:7ec/64 scope link
        valid_lft forever preferred_lft forever
3: docker0: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc noqueue state UP
    link/ether 02:42:e2:c5:a4:6b brd ff:ff:ff:ff:ff:ff
    inet 172.17.0.1/16 scope global docker0
        valid_lft forever preferred_lft forever
    inet6 fe80::42:e2ff:fec5:a46b/64 scope link

```

```

    valid_lft forever preferred_lft forever
5: vethfbd45f0@if4: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc noqueue
    master docker0 state UP
    link/ether 6e:3c:e4:21:7b:e2 brd ff:ff:ff:ff:ff:ff link-netnsid 0
    inet6 fe80::6c3c:e4ff:fe21:7be2/64 scope link
    valid_lft forever preferred_lft forever

```

A new interface has appeared: interface number 5 is the veth connection connecting the container's network namespace to the host's network namespace. But, what happened to interface number 4? It's been skipped in the list.

Look closely at interface number 5:

```
5: vethfbd45f0@if4
```

That @if4 indicates that interface number 5 is connected to interface 4. In fact, these are the two endpoints of the veth connection mentioned above; each end of the connection appears as a distinct interface, and `ip addr` only lists the interfaces in the current network namespace (the host in the above example).

6. Look at the interfaces in your container's network namespace (you'll first need to connect to the container and install `iproute`):

```

[centos@node-1 ~]$ docker container exec -it <container ID> bash
[root@f4e8f3f1b918 /]# yum install -y iproute
...
[root@f4e8f3f1b918 /]# ip addr

1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue
    state UNKNOWN group default qlen 1000
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
    inet 127.0.0.1/8 scope host lo
        valid_lft forever preferred_lft forever
4: eth0@if5: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc noqueue
    state UP group default
    link/ether 02:42:ac:11:00:02 brd ff:ff:ff:ff:ff:ff link-netnsid 0
    inet 172.17.0.2/16 scope global eth0
        valid_lft forever preferred_lft forever

```

Not only does interface number 4 appear inside the container's network namespace connected to interface 5, but we can see that this veth endpoint inside the container is getting treated as the `eth0` interface inside the container.

4.2 Establishing Custom Docker Networks

1. Create a custom bridge network:

```

[centos@node-1 ~]$ docker network create my_bridge
[centos@node-1 ~]$ docker network ls

```

NETWORK ID	NAME	DRIVER	SCOPE
7c4e63830cbf	bridge	bridge	local
c87d2a849036	host	host	local
a04d46bb85b1	my_bridge	bridge	local
902af00d5511	none	null	local

`my_bridge` gets created as another linux bridge-based network by default.

2. Run a couple of containers named `c2` and `c3` attached to this new network:


```
[centos@node-1 ~]$ docker container run \
  --name c2 --network my_bridge -d centos:7 ping 8.8.8.8
[centos@node-1 ~]$ docker container run \
  --name c3 --network my_bridge -d centos:7 ping 8.8.8.8
```

3. Inspect your new bridge:

```
[centos@node-1 ~]$ docker network inspect my_bridge

...
"IPAM": {
  "Driver": "default",
  "Options": {},
  "Config": [
    {
      "Subnet": "172.18.0.0/16",
      "Gateway": "172.18.0.1"
    }
  ]
},
...
"Containers": {
  "084caf415784fb4d58dc6fb4601321114b93dc148793fd66c95fc2c9411b085e": {
    "Name": "c3",
    "EndpointID": "8046005...",
    "MacAddress": "02:42:ac:12:00:03",
    "IPv4Address": "172.18.0.3/16",
    "IPv6Address": ""
  },
  "23d2e307325ec022ce6b08406bfb0f7e307fa533a7a4957a6d476c170d8e8658": {
    "Name": "c2",
    "EndpointID": "730ac71...",
    "MacAddress": "02:42:ac:12:00:02",
    "IPv4Address": "172.18.0.2/16",
    "IPv6Address": ""
  }
},
...
```

The next subnet in sequence (172.18.0.0/16 in my case) has been assigned to my_bridge by the IPAM driver, and containers attached to this network get IPs from this range exactly as they did with the default bridge network.

4. Try to contact container c3 from c2:

```
[centos@node-1 ~]$ docker container exec c2 ping c3
```

It works - containers on the same custom network are able to resolve each other via **DNS lookup of container names**. This means that our application logic (c2 ping c3 in this simple case) doesn't have to do any of its own service discovery; all we need to know are container names, and docker does the rest.

5. Start another container on my_bridge, but don't name it:

```
[centos@node-1 ~]$ docker container run --network my_bridge -d centos:7 ping 8.8.8.8
[centos@node-1 ~]$ docker container ls
```

CONTAINER ID	IMAGE	... STATUS	PORTS	NAMES
625cb95b922d	centos:7	... Up 2 seconds		competent_leavitt
084caf415784	centos:7	... Up 5 minutes		c3

```
23d2e307325e    centos:7    ... Up 5 minutes    c2
f4e8f3f1b918    centos:7    ... Up 21 minutes    zealous_kirch
```

As usual, it got a default name generated for it (`competent_leavitt` in my case). Try resolving this name by DNS as above:

```
[centos@node-1 ~]$ docker container exec c2 ping competent_leavitt

ping: competent_leavitt: Name or service not known
```

DNS resolution fails. Containers must be explicitly named in order to appear in docker's DNS tables.

- Find the IP of your latest container (`competent_leavitt` in my case) via `docker container inspect`, and ping it from `c2` directly by IP:

```
[centos@node-1 ~]$ docker network inspect my_bridge

...
"625cb95b922d2502fd016c6517c51652e84f902f69632d5d399dc38f3f7b2711": {
  "Name": "competent_leavitt",
  "EndpointID": "2fdb093d97b23da43023b07338a329180995fc0564ed0762147c8796380c51e7",
  "MacAddress": "02:42:ac:12:00:04",
  "IPv4Address": "172.18.0.4/16",
  "IPv6Address": ""
}
...

[centos@node-1 ~]$ docker container exec c2 ping 172.18.0.4

PING 172.18.0.4 (172.18.0.4) 56(84) bytes of data.
64 bytes from 172.18.0.4: icmp_seq=1 ttl=64 time=0.083 ms
64 bytes from 172.18.0.4: icmp_seq=2 ttl=64 time=0.060 ms
```

The ping succeeds. While the default-named container isn't resolvable by DNS, it is still reachable on the `my_bridge` network.

- Finally, create container `c1` attached to the default network:

```
[centos@node-1 ~]$ docker container run --name c1 -d centos:7 ping 8.8.8.8
```

Attempt to ping it from `c2` by name:

```
[centos@node-1 ~]$ docker container exec c2 ping c1

ping: c1: Name or service not known
```

DNS resolution is scoped to user-defined docker networks. Find `c1`'s IP manually as above (mine is at `172.17.0.3`), and ping this IP directly from `c2`:

```
[centos@node-1 ~]$ docker container exec c2 ping 172.17.0.3
```

The request hangs until it times out (press `CTRL+C` to give up early if you don't want to wait for the timeout). Different docker networks are firewalled from each other by default; dump your iptables rules and look for lines similar to the following:

```
[centos@node-1 ~]$ sudo iptables-save

...
-A DOCKER-ISOLATION-STAGE-1 -i br-dfda80f70ea5
  ! -o br-dfda80f70ea5 -j DOCKER-ISOLATION-STAGE-2
-A DOCKER-ISOLATION-STAGE-1 -i docker0 ! -o docker0 -j DOCKER-ISOLATION-STAGE-2
-A DOCKER-ISOLATION-STAGE-1 -j RETURN
```

```
-A DOCKER-ISOLATION-STAGE-2 -o br-dfda80f70ea5 -j DROP
-A DOCKER-ISOLATION-STAGE-2 -o docker0 -j DROP
-A DOCKER-ISOLATION-STAGE-2 -j RETURN
...
```

The first line above forwards traffic originating from `br-dfda80f70ea5` (that's your custom bridge) but destined somewhere else to the stage 2 isolation chain, where if it is destined for the `docker0` bridge, it gets dropped, preventing traffic from going from one bridge to another.

4.3 Forwarding a Host Port to a Container

1. Start an `nginx` container with a port exposure:

```
[centos@node-1 ~]$ docker container run -d -p 8000:80 nginx
```

This syntax asks docker to forward all traffic arriving on port 8000 of the host's network namespace to port 80 of the container's network namespace. Visit the `nginx` landing page at `<node-1 public IP>:8000`.

2. Inspect your iptables rules again to see how docker forwarded this traffic:

```
[centos@node-1 ~]$ sudo iptables-save | grep 8000

-A DOCKER ! -i docker0 -p tcp -m tcp --dport 8000
  -j DNAT --to-destination 172.17.0.4:80
```

Inspect your default bridge network to find the IP of your `nginx` container; you should find that it matches the IP in the network address translation rule above, which states that any traffic arriving on port `tcp/8000` on the host should be network address translated to `172.17.0.4:80` - the IP of our `nginx` container and the port we exposed with the `-p 8000:80` flag when we created this container.

3. Clean up your containers and networks:

```
[centos@node-1 ~]$ docker container rm -f $(docker container ls -aq)
[centos@node-1 ~]$ docker network rm my_bridge
```

4.4 Conclusion

In this demo, we stepped through the basic behavior of docker software defined bridge networks, and looked at the technology underpinning them such as linux bridges, veth connections, and iptables rules. From a practical standpoint, in order for containers to communicate they must be attached to the same docker software defined network (otherwise they'll be firewalled from each other by the cross-network iptables rules we saw), and in order for containers to resolve each other's name by DNS, they must also be explicitly named upon creation.

5 Instructor Demo: Self-Healing Swarm

In this demo, we'll illustrate:

- Setting up a swarm
- How swarm makes basic scheduling decisions
- Actions swarm takes to self-heal a docker service

5.1 Setting Up a Swarm

1. Start by making sure no containers are running on any of your nodes:

```
[centos@node-0 ~]$ docker container rm -f $(docker container ls -aq)
[centos@node-1 ~]$ docker container rm -f $(docker container ls -aq)
[centos@node-2 ~]$ docker container rm -f $(docker container ls -aq)
[centos@node-3 ~]$ docker container rm -f $(docker container ls -aq)
```

2. Initialize a swarm on one node:

```
[centos@node-0 ~]$ docker swarm init
```

Swarm initialized: current node (xyz) is now a manager.

To add a worker to this swarm, run the following command:

```
docker swarm join --token SWMTKN-1-0s96... 10.10.1.40:2377
```

To add a manager to this swarm, run

```
'docker swarm join-token manager' and follow the instructions.
```

3. List the nodes in your swarm:

```
[centos@node-0 ~]$ docker node ls
```

ID	HOSTNAME	STATUS	AVAILABILITY	MANAGER STATUS
xyz *	node-0	Ready	Active	Leader

4. Add some workers to your swarm by cutting and pasting the `docker swarm join... token` Docker provided in step 2 above:

```
[centos@node-1 ~]$ docker swarm join --token SWMTKN-1-0s96... 10.10.1.40:2377
[centos@node-2 ~]$ docker swarm join --token SWMTKN-1-0s96... 10.10.1.40:2377
[centos@node-3 ~]$ docker swarm join --token SWMTKN-1-0s96... 10.10.1.40:2377
```

Each node should report `This node joined a swarm as a worker.` after joining.

5. Back on your first node, list your swarm members again:

```
[centos@node-0 ~]$ docker node ls
```

ID	HOSTNAME	STATUS	AVAILABILITY	MANAGER STATUS
ghi	node-3	Ready	Active	
def	node-2	Ready	Active	
abc	node-1	Ready	Active	
xyz *	node-0	Ready	Active	Leader

You have a four-member swarm, ready to accept workloads.

5.2 Scheduling Workload

1. Create a service on your swarm:

```
[centos@node-0 ~]$ docker service create \
  --replicas 4 \
  --name service-demo \
  centos:7 ping 8.8.8.8
```

2. List what processes have been started for your service:

```
[centos@node-0 ~]$ docker service ps service-demo
```

ID	NAME	IMAGE	NODE	DESIRED STATE	CURRENT STATE
----	------	-------	------	---------------	---------------

g3dimc0nkoha	service-demo.1	centos:7	node-3	Running	Running 18 seconds ago
e7d7sy5saqqo	service-demo.2	centos:7	node-0	Running	Running 18 seconds ago
wv0culf6w8m6	service-demo.3	centos:7	node-1	Running	Running 18 seconds ago
ty35gss71mpf	service-demo.4	centos:7	node-2	Running	Running 18 seconds ago

Our service has scheduled four tasks, one on each node in our cluster; by default, swarm tries to spread tasks out evenly across hosts, but much more sophisticated scheduling controls are also available.

5.3 Maintaining Desired State

1. Switch to node-1, and list the containers running there:

```
[centos@node-1 ~]$ docker container ls
```

ID	IMAGE	COMMAND	CREATED	STATUS	NAMES
5b5...	centos:7	"ping 8.8.8.8"	4 minutes ago	Up 4 minutes	service-demo.3.wv0...

Note the container's name indicates the service it belongs to.

2. Let's simulate a container crash, by killing off this container:

```
[centos@node-1 ~]$ docker container rm -f <container ID>
```

Back on our swarm manager, list the processes running for our service-demo service again:

```
[centos@node-0 ~]$ docker service ps service-demo
```

ID	NAME	IMAGE	NODE	DESIRED STATE	CURRENT STATE
g3d...	service-demo.1	centos:7	node-3	Running	Running 6 minutes ago
e7d...	service-demo.2	centos:7	node-0	Running	Running 6 minutes ago
u7l...	service-demo.3	centos:7	node-1	Running	Running 3 seconds ago
wv0...	_ service-demo.3	centos:7	node-1	Shutdown	Failed 3 seconds ago
ty3...	service-demo.4	centos:7	node-2	Running	Running 6 minutes ago

Swarm has automatically started a replacement container for the one you killed on node-1. Go back over to node-1, and do `docker container ls` again; you'll see a new container for this service up and running.

3. Next, let's simulate a complete node failure by rebooting one of our nodes:

```
[centos@node-3 ~]$ sudo reboot now
```

4. Back on your swarm manager, check your service containers again:

```
[centos@node-0 ~]$ docker service ps service-demo
```

ID	NAME	IMAGE	NODE	DESIRED STATE	CURRENT STATE
ral...	service-demo.1	centos:7	node-0	Running	Running 19 seconds ago
g3d...	_ service-demo.1	centos:7	node-3	Shutdown	Running 38 seconds ago
e7d...	service-demo.2	centos:7	node-0	Running	Running 12 minutes ago
u7l...	service-demo.3	centos:7	node-1	Running	Running 5 minutes ago
wv0...	_ service-demo.3	centos:7	node-1	Shutdown	Failed 5 minutes ago
ty3...	service-demo.4	centos:7	node-2	Running	Running 12 minutes ago

The process on node-3 has been scheduled for SHUTDOWN when the swarm manager lost connection to that node, and meanwhile the workload has been rescheduled onto node-0 in this case. When node-3 comes back up and rejoins the swarm, its container will be confirmed to be in the SHUTDOWN state, and reconciliation is complete.

5. Remove your service-demo:

```
[centos@node-0 ~]$ docker service rm service-demo
```

All tasks and containers will be removed.

5.4 Conclusion

One of the great advantages of the portability of containers is that we can imagine orchestrators like Swarm which can schedule and re-schedule workloads across an entire datacenter, such that if a given node fails, all its workload can be automatically moved to another host with available resources. In the above example, we saw the most basic examples of this 'reconciliation loop' that swarm provides: the swarm manager is constantly monitoring all the containers it has scheduled, and replaces them if they fail or their hosts become unreachable, completely automatically.

6 Instructor Demo: Kubernetes Basics

In this demo, we'll illustrate:

- Setting up a Kubernetes cluster with one master and two nodes
- Scheduling a pod, including the effect of taints on scheduling
- Namespaces shared by containers in a pod

6.1 Initializing Kubernetes

1. On node-0, initialize the cluster with kubeadm:

```
[centos@node-0 ~]$ sudo kubeadm init --pod-network-cidr=192.168.0.0/16 \
--ignore-preflight-errors=SystemVerification
```

If successful, the output will end with a join command:

```
...
You can now join any number of machines by running the following on each node
as root:

kubeadm join 10.10.29.54:6443 --token xxx --discovery-token-ca-cert-hash sha256:yyy
```

2. To start using you cluster, you need to run:

```
[centos@node-0 ~]$ mkdir -p $HOME/.kube
[centos@node-0 ~]$ sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
[centos@node-0 ~]$ sudo chown $(id -u):$(id -g) $HOME/.kube/config
```

3. List all your nodes in the cluster:

```
[centos@node-0 ~]$ kubectl get nodes
```

Which should output something like:

NAME	STATUS	ROLES	AGE	VERSION
node-0	NotReady	master	2h	v1.16.3

The NotReady status indicates that we must install a network for our cluster.

4. Let's install the Calico network driver:

```
[centos@node-0 ~]$ kubectl apply -f https://bit.ly/2KHCEg0
```

5. After a moment, if we list nodes again, ours should be ready:

```
[centos@node-0 ~]$ kubectl get nodes -w
NAME          STATUS    ROLES    AGE    VERSION
node-0        NotReady  master   1m     v1.16.3
node-0        NotReady  master   1m     v1.16.3
node-0        NotReady  master   1m     v1.16.3
node-0        Ready     master   2m     v1.16.3
node-0        Ready     master   2m     v1.16.3
```

6.2 Exploring Kubernetes Scheduling

1. Let's create a demo-pod.yaml file on node-0 after enabling Kubernetes on this single node:

```
apiVersion: v1
kind: Pod
metadata:
  name: demo-pod
spec:
  volumes:
    - name: shared-data
      emptyDir: {}
  containers:
    - name: nginx
      image: nginx
    - name: mydemo
      image: centos:7
      command: ["ping", "8.8.8.8"]
```

2. Deploy the pod:

```
[centos@node-0 ~]$ kubectl create -f demo-pod.yaml
```

3. Check to see if the pod is running:

```
[centos@node-0 ~]$ kubectl get pod demo-pod
```

NAME	READY	STATUS	RESTARTS	AGE
demo-pod	0/2	Pending	0	7s

The status should be stuck in pending. Why is that?

4. Let's attempt to troubleshoot by obtaining some information about the pod:

```
[centos@node-0 ~]$ kubectl describe pod demo-pod
```

In the bottom section titled Events:, we should see something like this:

```
...
Events:
  Type     Reason             ... Message
  ----     -
  Warning   FailedScheduling   ... 0/1 nodes are available: 1 node(s)
                                     had taints that the pod didn't tolerate.
```

Note how it states that the one node in your cluster has a taint, which is Kubernetes's way of saying there's a reason you might not want to schedule pods there.

5. Get some state and config information about your single kubernetes node:

```
[centos@node-0 ~]$ kubectl describe nodes
```

If we scroll a little, we should see a field titled Taints, and it should say something like:

```
Taints: node-role.kubernetes.io/master:NoSchedule
```

By default, Kubernetes masters carry a taint that disallows scheduling pods on them. While this can be overridden, it is best practice to not allow pods to get scheduled on master nodes, in order to ensure the stability of your cluster.

6. Execute the join command you found above when initializing Kubernetes on node-1 and node-2 (you'll need to add sudo to the start, and --ignore-preflight-errors=SystemVerification to the end), and then check the status back on node-0:

```
[centos@node-1 ~]$ sudo kubeadm join...--ignore-preflight-errors=SystemVerification
[centos@node-2 ~]$ sudo kubeadm join...--ignore-preflight-errors=SystemVerification
[centos@node-0 ~]$ kubectl get nodes
```

After a few moments, there should be three nodes listed - all with the Ready status.

- Let's see what system pods are running on our cluster:

```
[centos@node-0 ~]$ kubectl get pods -n kube-system
```

which results in something similar to this:

NAME	READY	STATUS	RESTARTS	AGE
calico-kube-controllers-6b64bcd855-lnrw7	1/1	Running	0	2m13s
calico-node-btgn6	1/1	Running	0	54s
calico-node-rdhrb	0/1	Running	0	20s
calico-node-s7dnc	1/1	Running	0	2m13s
coredns-5644d7b6d9-c5j5c	1/1	Running	0	7m29s
coredns-5644d7b6d9-gm4rl	1/1	Running	0	7m29s
etcd-node-0	1/1	Running	0	6m39s
kube-apiserver-node-0	1/1	Running	0	6m41s
kube-controller-manager-node-0	1/1	Running	0	6m26s
kube-proxy-jc6h7	1/1	Running	0	20s
kube-proxy-nczr9	1/1	Running	0	7m29s
kube-proxy-sf9n6	1/1	Running	0	54s
kube-scheduler-node-0	1/1	Running	0	6m51s

We can see the pods running on the master: etcd, api-server, controller manager and scheduler, as well as calico and DNS infrastructure pods deployed when we installed calico.

- Finally, let's check the status of our demo pod now:

```
[centos@node-0 ~]$ kubectl get pod demo-pod
```

Everything should be working correctly with 2/2 containers in the pod running, now that there are un-tainted nodes for the pod to get scheduled on.

6.3 Exploring Containers in a Pod

- Let's interact with the centos container running in demo-pod by getting a shell in it:

```
[centos@node-0 ~]$ kubectl exec -it -c mydemo demo-pod -- /bin/bash
```

Try listing the processes in this container:

```
[root@demo-pod /]# ps -aux
USER      PID %CPU %MEM    VSZ   RSS TTY      STAT START   TIME COMMAND
root         1  0.0  0.0  24860  1992 ?        Ss   14:48   0:00 ping 8.8.8.8
root         5  0.0  0.0  11832  3036 pts/0    Ss   14:48   0:00 /bin/bash
root        20  0.0  0.0  51720  3508 pts/0    R+   14:48   0:00 ps -aux
```

We can see the ping process we containerized in our yaml file running as PID 1 inside this container, just like we saw for plain containers.

- Try reaching Nginx:

```
[root@demo-pod /]# curl localhost:80
```

You should see the HTML for the default nginx landing page. Notice the difference here from a regular container; we were able to reach our nginx deployment from our centos container on a port on localhost. The nginx and centos containers share a network namespace and therefore all their ports, since they are part of the same pod.

6.4 Conclusion

In this demo, we saw two scheduling innovations Kubernetes offers: taints, which provide ‘anti-affinity’, or reasons not to schedule a pod on a given node; and pods, which are groups of containers that are always scheduled on the same node, and share network, IPC and hostname namespaces. These are both examples of Kubernetes’s highly expressive scheduling, and are both difficult to reproduce with the simpler scheduling offered by Swarm.

7 Instructor Demo: Docker Compose

In this demo, we’ll illustrate:

- Starting an app defined in a docker compose file
- Inter-service communication using DNS resolution of service names

7.1 Exploring the Compose File

1. Please download the DockerCoins app from Github and change directory to ~/orchestration-workshop/dockercoins.

```
[centos@node-0 ~]$ git clone -b ee3.0 \
  https://github.com/docker-training/orchestration-workshop.git
[centos@node-0 ~]$ cd ~/orchestration-workshop/dockercoins
```

2. Let’s take a quick look at our Compose file for Dockercoins:

```
version: "3.1"

services:
  rng:
    image: training/dockercoins-rng:1.0
    networks:
      - dockercoins
    ports:
      - "8001:80"

  hasher:
    image: training/dockercoins-hashier:1.0
    networks:
      - dockercoins
    ports:
      - "8002:80"

  webui:
    image: training/dockercoins-webui:1.0
    networks:
      - dockercoins
    ports:
      - "8000:80"

  redis:
    image: redis
    networks:
      - dockercoins

  worker:
    image: training/dockercoins-worker:1.0
    networks:
```

```
- dockercoins

networks:
  dockercoins:
```

This Compose file contains 5 services, along with a bridge network.

- When we start the app, we will see the service images getting downloaded one at a time:

```
[centos@node-0 dockercoins]$ docker-compose up -d
```

- After starting, the images required for this app have been downloaded:

```
[centos@node-0 dockercoins]$ docker image ls | grep "dockercoins"
```

- Make sure the services are up and running, as is the dedicated network:

```
[centos@node-0 dockercoins]$ docker-compose ps
[centos@node-0 dockercoins]$ docker network ls
```

- If everything is up, visit your app at <node-0 public IP>:8000 to see Dockercoins in action.

7.2 Communicating Between Containers

- In this section, we'll demonstrate that containers created as part of a service in a Compose file are able to communicate with containers belonging to other services using just their service names. Let's start by listing our DockerCoins containers:

```
[centos@node-0 dockercoins]$ docker container ls | grep 'dockercoins'
```

- Now, connect into one container; let's pick webui:

```
[centos@node-0 dockercoins]$ docker container exec -it <Container ID> bash
```

- From within the container, ping rng by name:

```
[root@<Container ID>]# ping rng
```

Logs should be outputted resembling this:

```
PING rng (172.18.0.5) 56(84) bytes of data.
64 bytes from dockercoins_rng_1... (172.18.0.5): icmp_seq=1 ttl=64 time=0.108 ms
64 bytes from dockercoins_rng_1... (172.18.0.5): icmp_seq=2 ttl=64 time=0.049 ms
64 bytes from dockercoins_rng_1... (172.18.0.5): icmp_seq=3 ttl=64 time=0.073 ms
64 bytes from dockercoins_rng_1... (172.18.0.5): icmp_seq=4 ttl=64 time=0.067 ms
64 bytes from dockercoins_rng_1... (172.18.0.5): icmp_seq=5 ttl=64 time=0.057 ms
64 bytes from dockercoins_rng_1... (172.18.0.5): icmp_seq=6 ttl=64 time=0.074 ms
64 bytes from dockercoins_rng_1... (172.18.0.5): icmp_seq=7 ttl=64 time=0.052 ms
64 bytes from dockercoins_rng_1... (172.18.0.5): icmp_seq=8 ttl=64 time=0.057 ms
64 bytes from dockercoins_rng_1... (172.18.0.5): icmp_seq=9 ttl=64 time=0.080 ms
```

Use CTRL+C to terminate the ping. DNS lookup for the services in DockerCoins works because they are all attached to the user-defined dockercoins network.

- After exiting this container, let's navigate to the worker folder and take a look at a section of worker.py:

```
[centos@node-0 dockercoins]$ cd worker
[centos@node-0 worker]$ cat worker.py

import logging
import os
from redis import Redis
import requests
```

```
import time

DEBUG = os.environ.get("DEBUG", "").lower().startswith("y")

log = logging.getLogger(__name__)
if DEBUG:
    logging.basicConfig(level=logging.DEBUG)
else:
    logging.basicConfig(level=logging.INFO)
    logging.getLogger("requests").setLevel(logging.WARNING)

redis = Redis("redis")

def get_random_bytes():
    r = requests.get("http://rng/32")
    return r.content

def hash_bytes(data):
    r = requests.post("http://hasher/",
                      data=data,
                      headers={"Content-Type": "application/octet-stream"})
    hex_hash = r.text
    return hex_hash
```

As we can see in the last two stanzas, we can direct traffic to a service via a DNS name that exactly matches the service name defined in the docker compose file.

5. Shut down Dockercoins and clean up its resources:

```
[centos@node-0 dockercoins]$ docker-compose down
```

7.3 Conclusion

In this exercise, we stood up an application using Docker Compose. The most important new idea here is the notion of Docker Services, which are collections of identically configured containers. Docker Service names are resolvable by DNS, so that we can write application logic designed to communicate service to service; all service discovery and load balancing between your application's services is abstracted away and handled by Docker.